

NEURAL NETWORKS AND DEEP LEARNING

MICHELLE KUCHERA
DAVIDSON COLLEGE

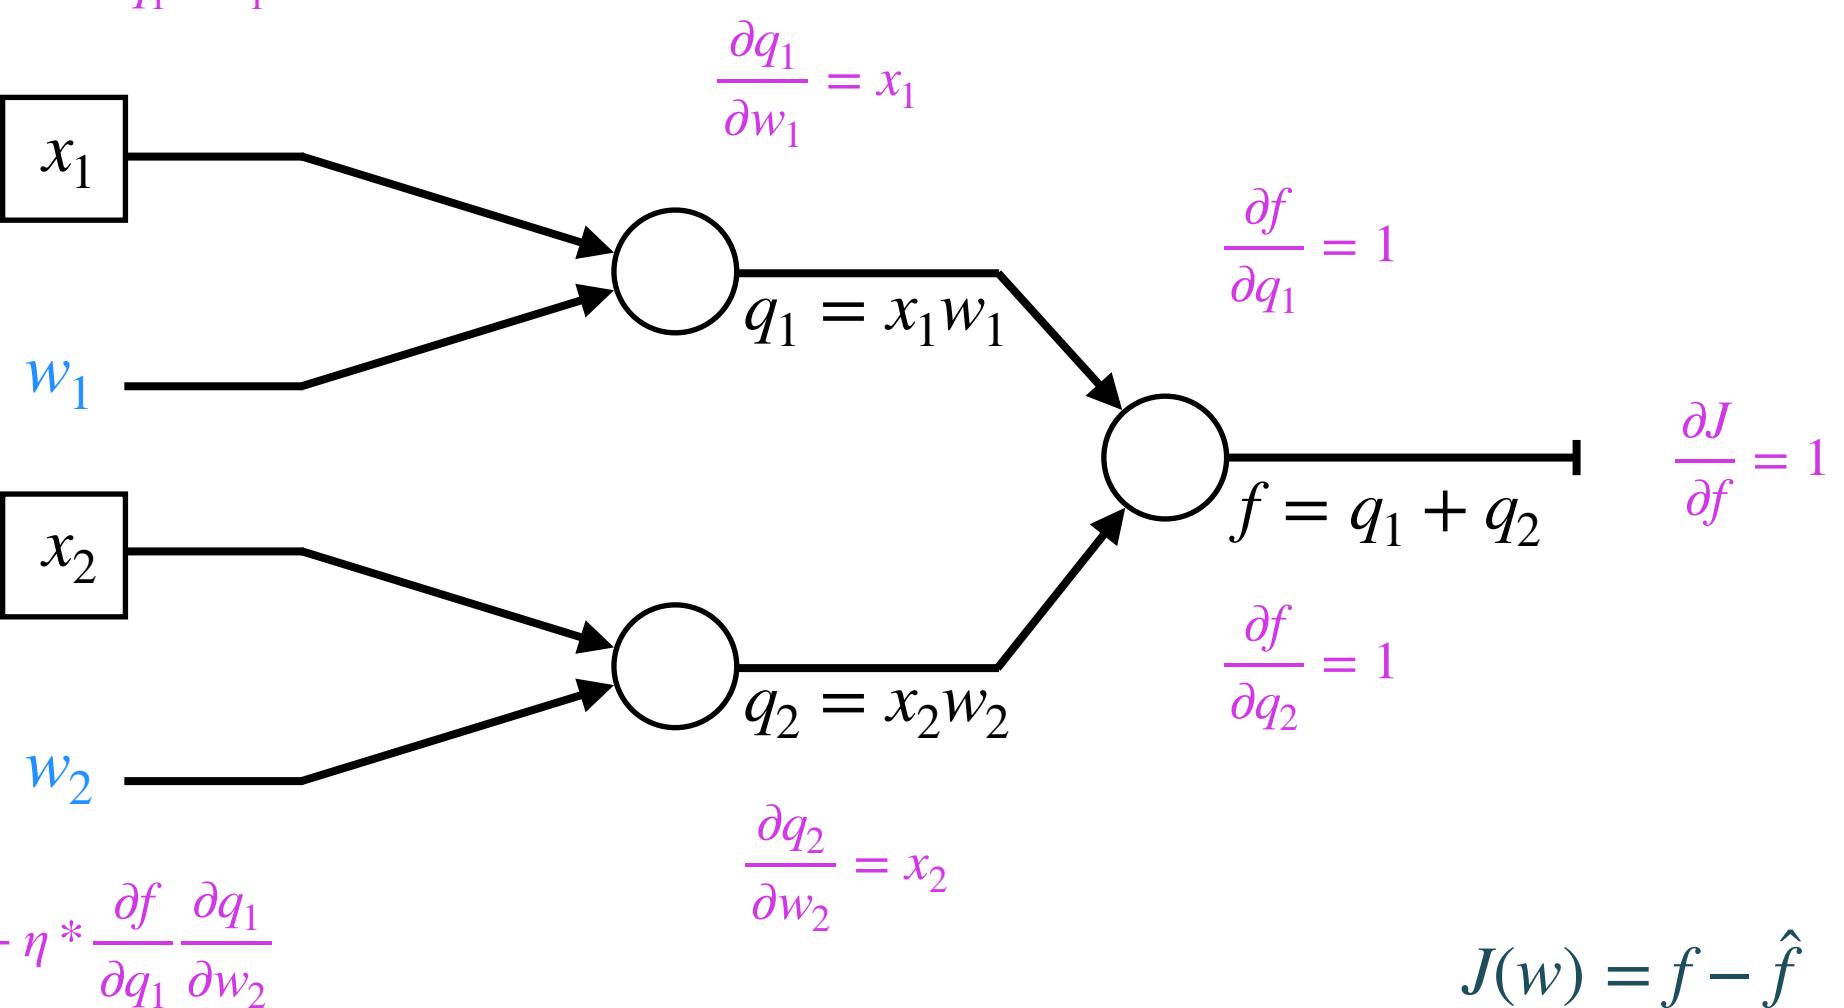
CPS-FR
MIT
25 AUGUST 2023

MICHELLE KUCHERA

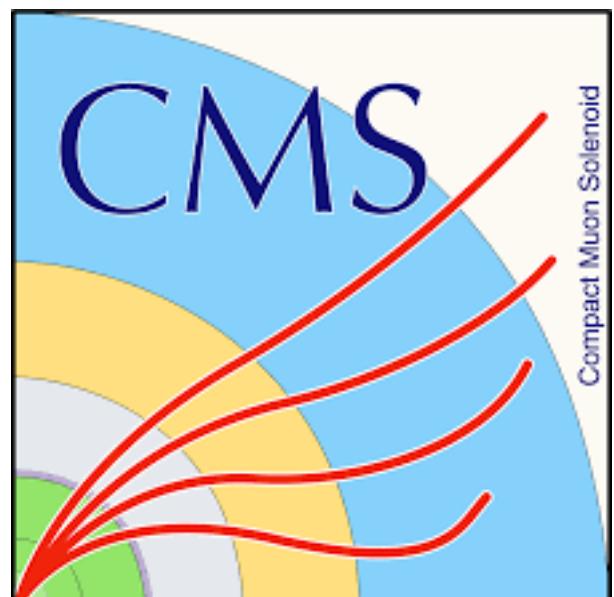
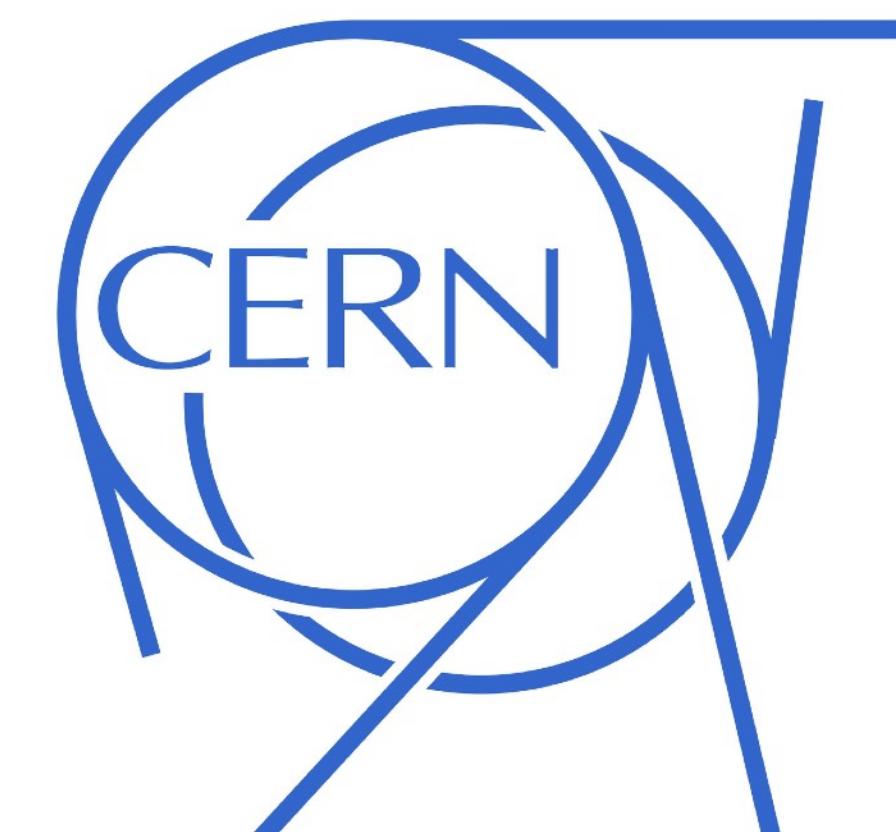
B.S., M.S. PHYSICS
M.S., PH.D. COMPUTATIONAL SCIENCE



$$w_1 = w_1 + \eta * \frac{\partial f}{\partial q_1} \frac{\partial q_1}{\partial w_1}$$



ALPhA
DAVIDSON COLLEGE



Jefferson Lab

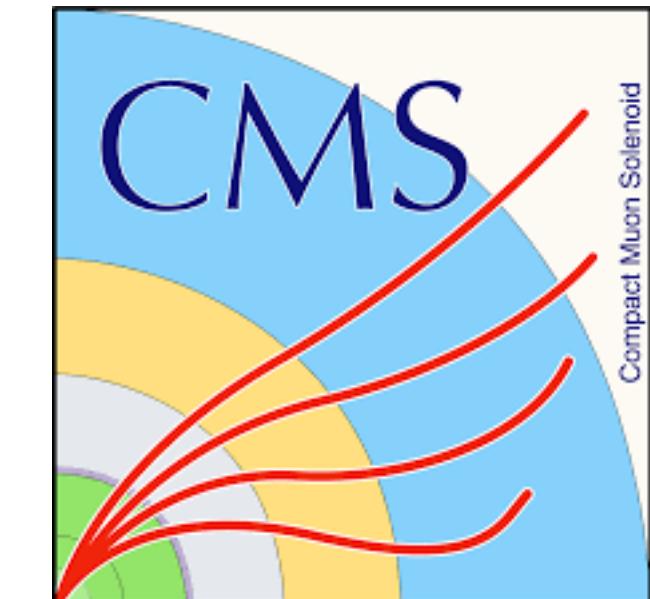
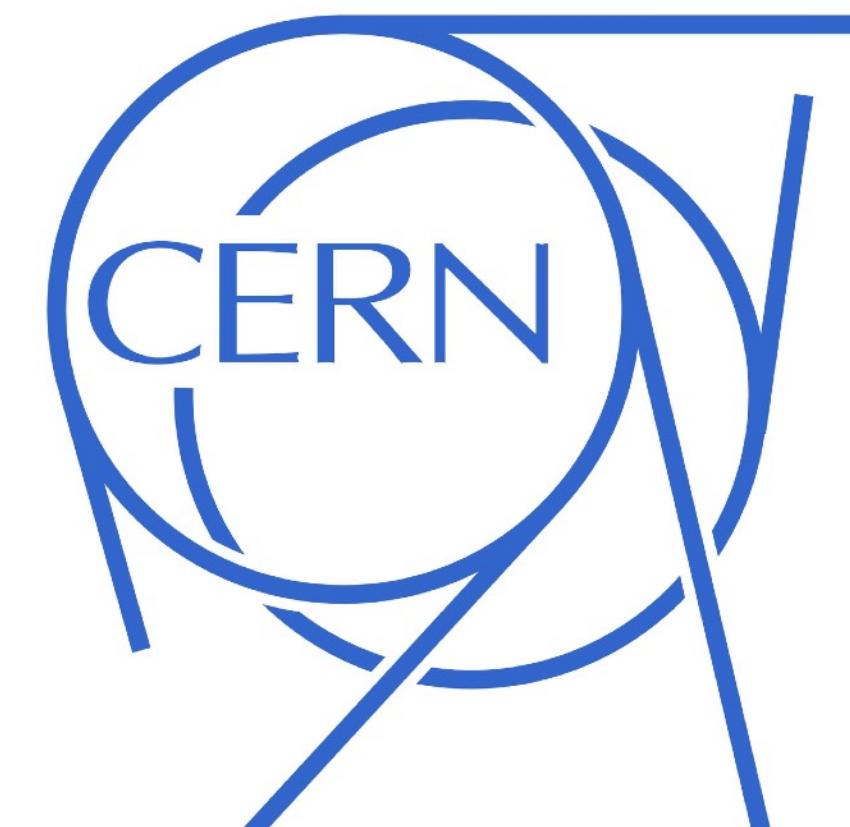
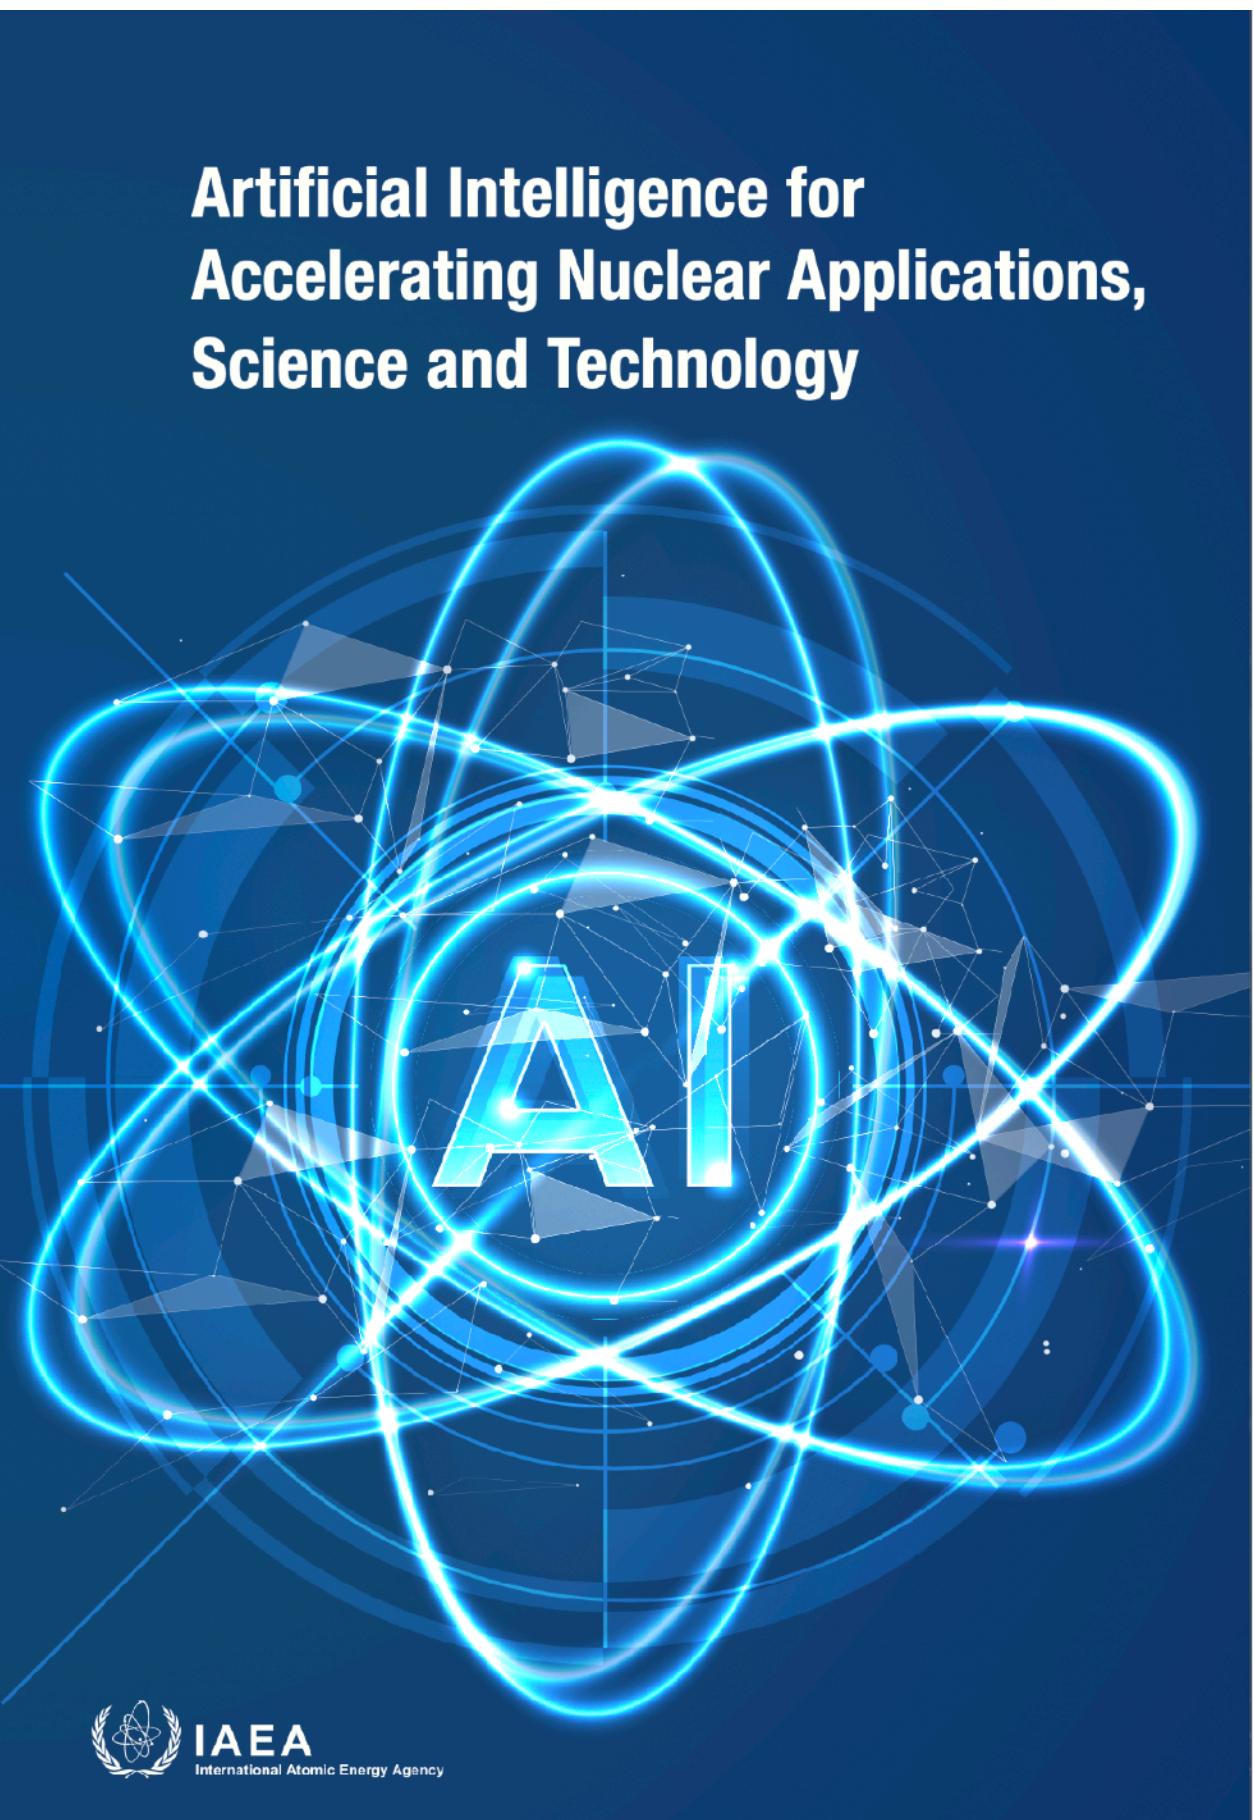
NSCL

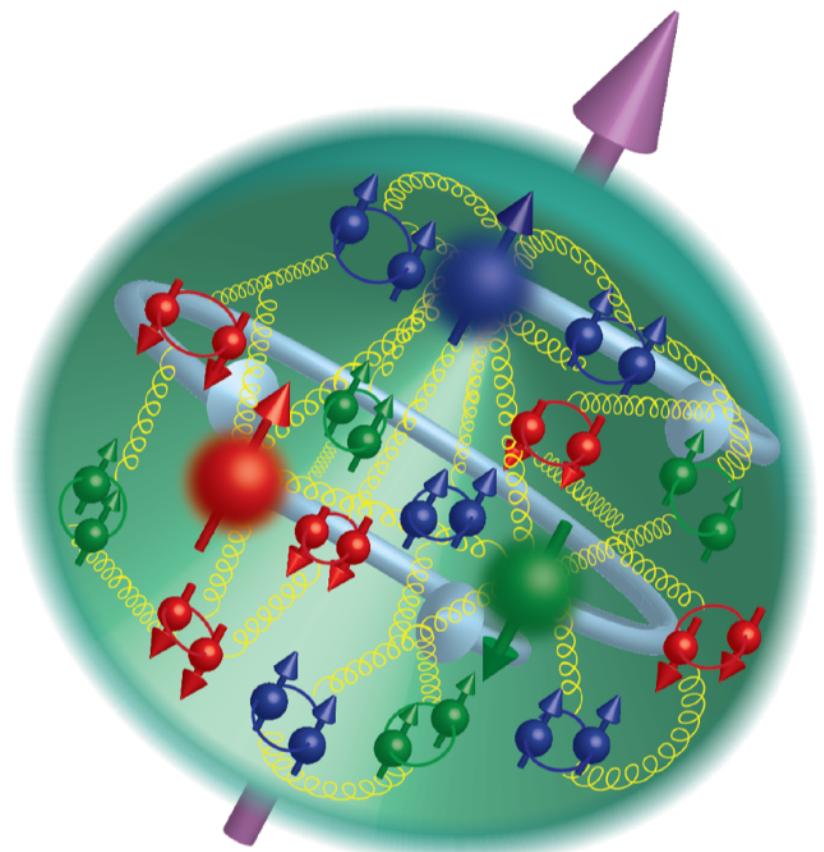
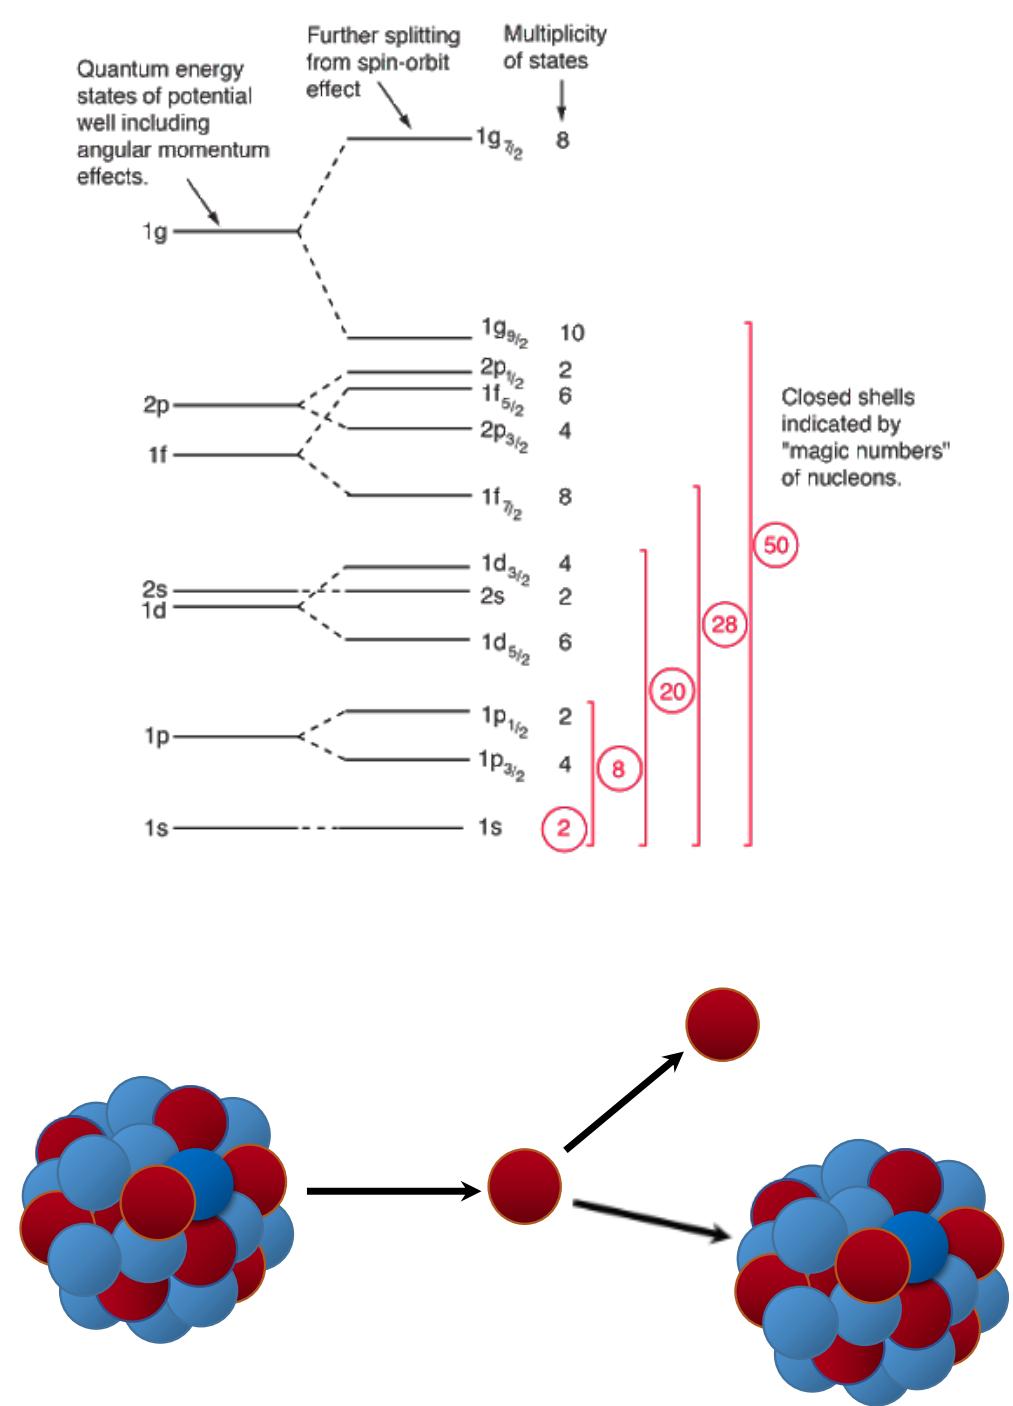
PhD: GPUs for Bayesian Neural Networks (😢)

MICHELLE KUCHERA

B.S., M.S. PHYSICS

M.S., PH.D. COMPUTATIONAL SCIENCE

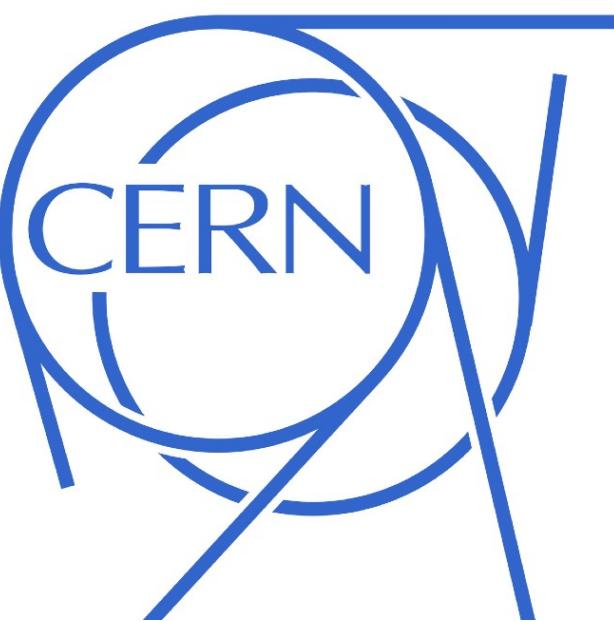




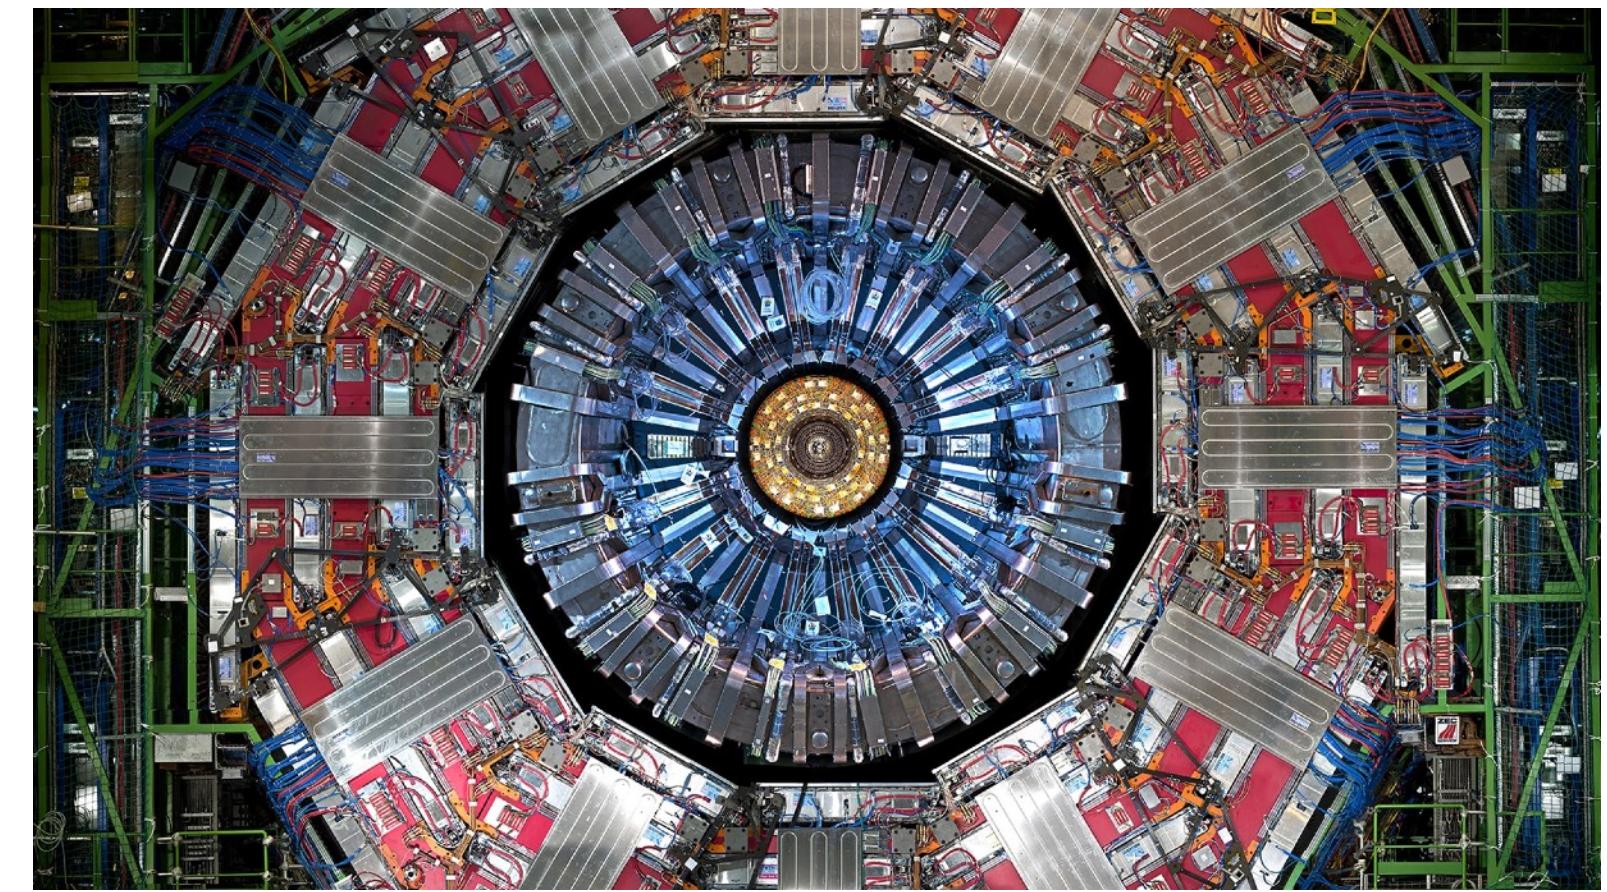
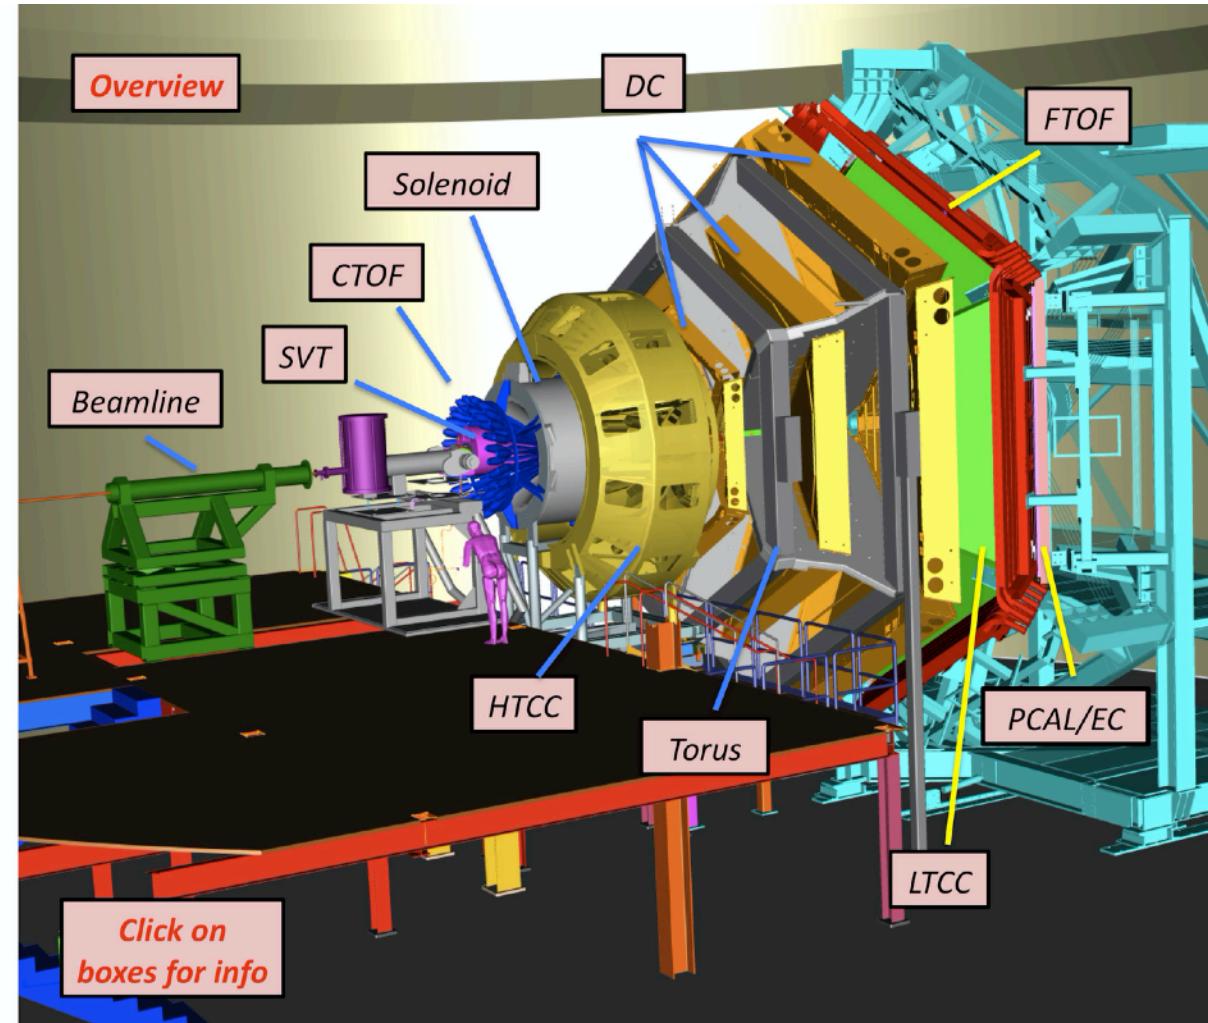
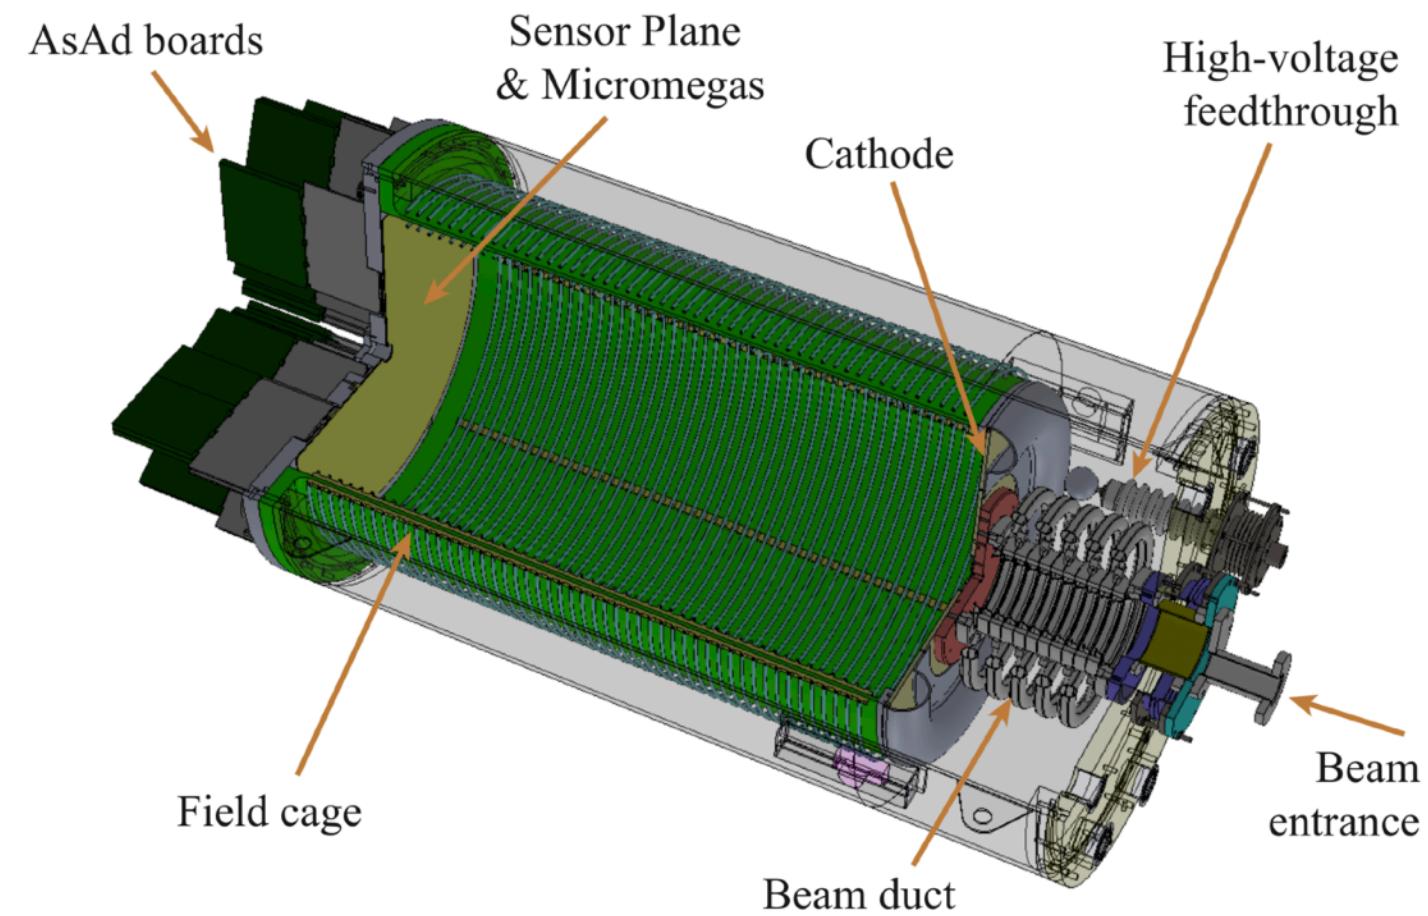
QUARKS		GAUGE BOSONS	
mass $\rightarrow +0.3 \text{ MeV/c}^2$	charge $\rightarrow +2/3$	u	H
spin $\rightarrow 1/2$		up	Higgs boson
$+4.8 \text{ MeV/c}^2$	$+1/3 \text{ e}$	c	
$+95 \text{ MeV/c}^2$	$+2/3 \text{ e}$	t	
$+4.10 \text{ GeV/c}^2$	$+1/2 \text{ e}$	charm	
$+13 \text{ GeV/c}^2$	$+1/2 \text{ e}$	top	
$+0.5 \text{ MeV/c}^2$	$-1/3 \text{ e}$	g	
-13 GeV/c^2	$-1/2 \text{ e}$	gluon	
-13 GeV/c^2	$-1/2 \text{ e}$	γ	
-91.2 GeV/c^2	0 e	b	
-13 GeV/c^2	0 e	bottom	
-91.2 GeV/c^2	0 e	Z	
-91.2 GeV/c^2	0 e	Z boson	
-0.2 eV/c^2	0 e	e	
-0.2 eV/c^2	0 e	μ	
-0.2 eV/c^2	0 e	τ	
$< 10^{-3} \text{ MeV/c}^2$	0 e	ν_e	
$< 10^{-3} \text{ MeV/c}^2$	0 e	ν_μ	
$< 10^{-3} \text{ MeV/c}^2$	0 e	ν_τ	
$< 10^{-3} \text{ MeV/c}^2$	0 e	$\bar{\nu}_e$	
$< 10^{-3} \text{ MeV/c}^2$	0 e	$\bar{\nu}_\mu$	
$< 10^{-3} \text{ MeV/c}^2$	0 e	$\bar{\nu}_\tau$	
$< 10^{-3} \text{ MeV/c}^2$	0 e	W	
$< 10^{-3} \text{ MeV/c}^2$	0 e	W boson	



Jefferson Lab



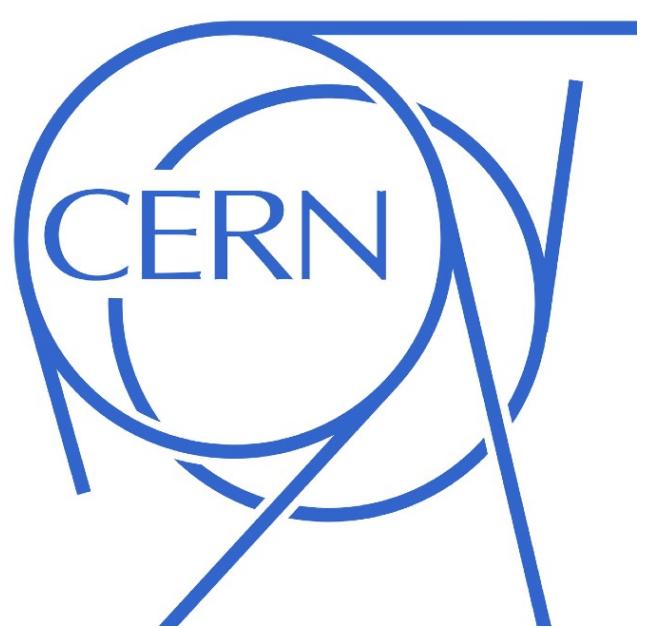
EXPERIMENTAL DATA



J. BRADT ET. AL., NUCLEAR INSTRUMENTS AND METHODS, 2017.



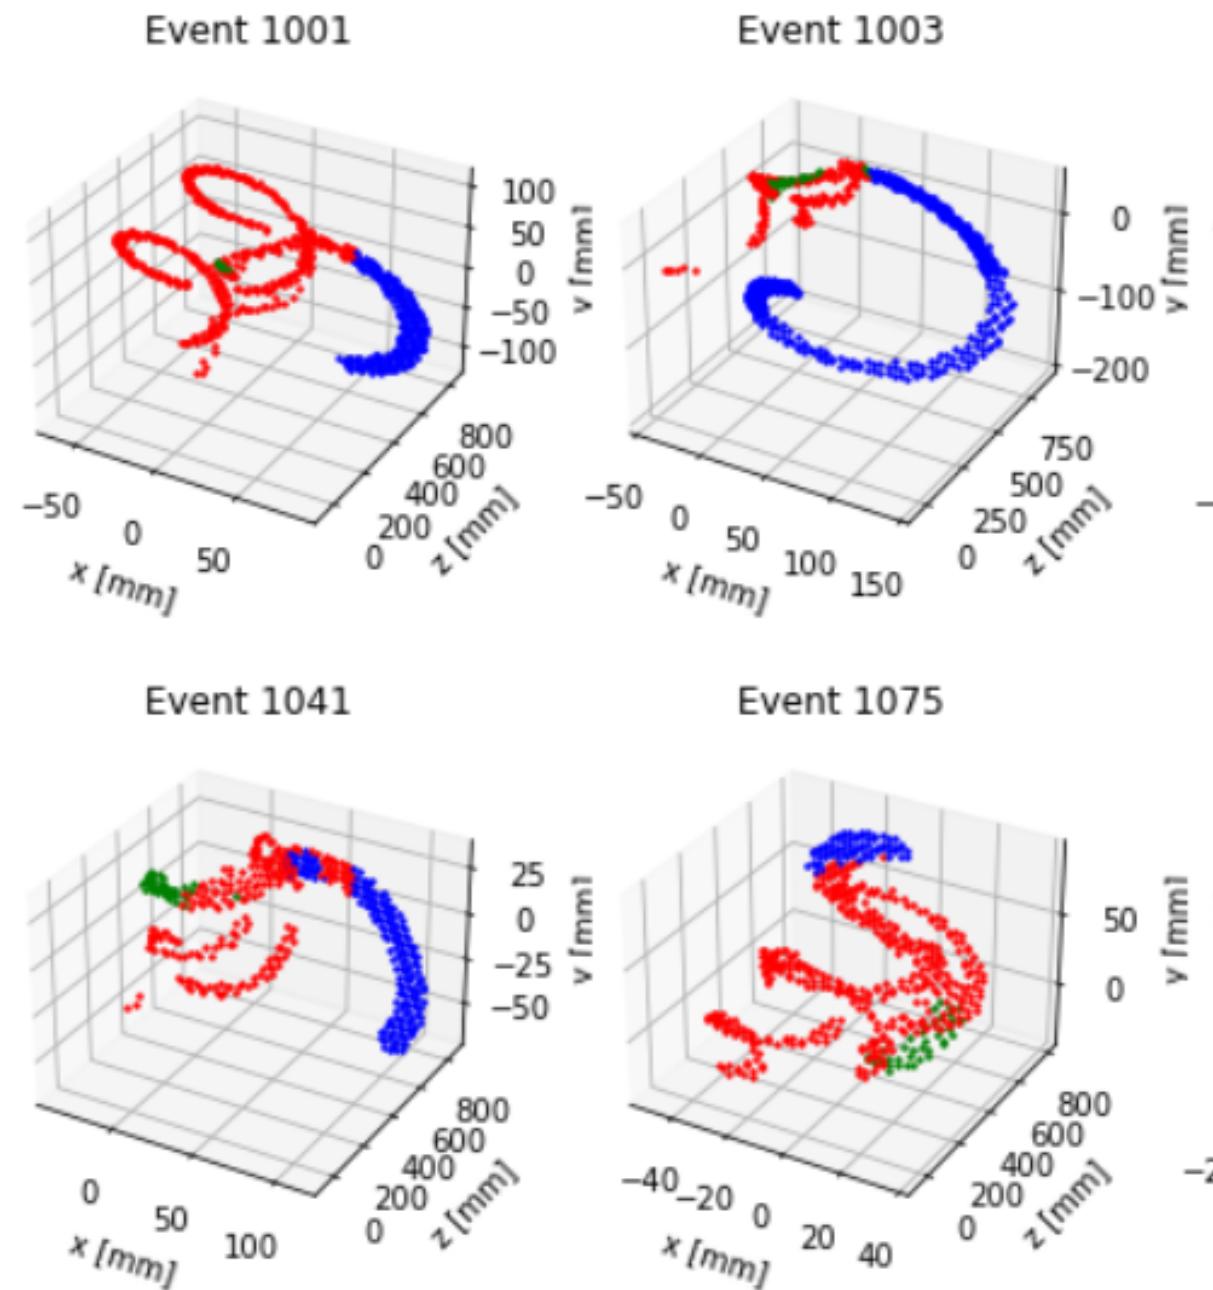
AT-TPC



CLAS 12

CMS

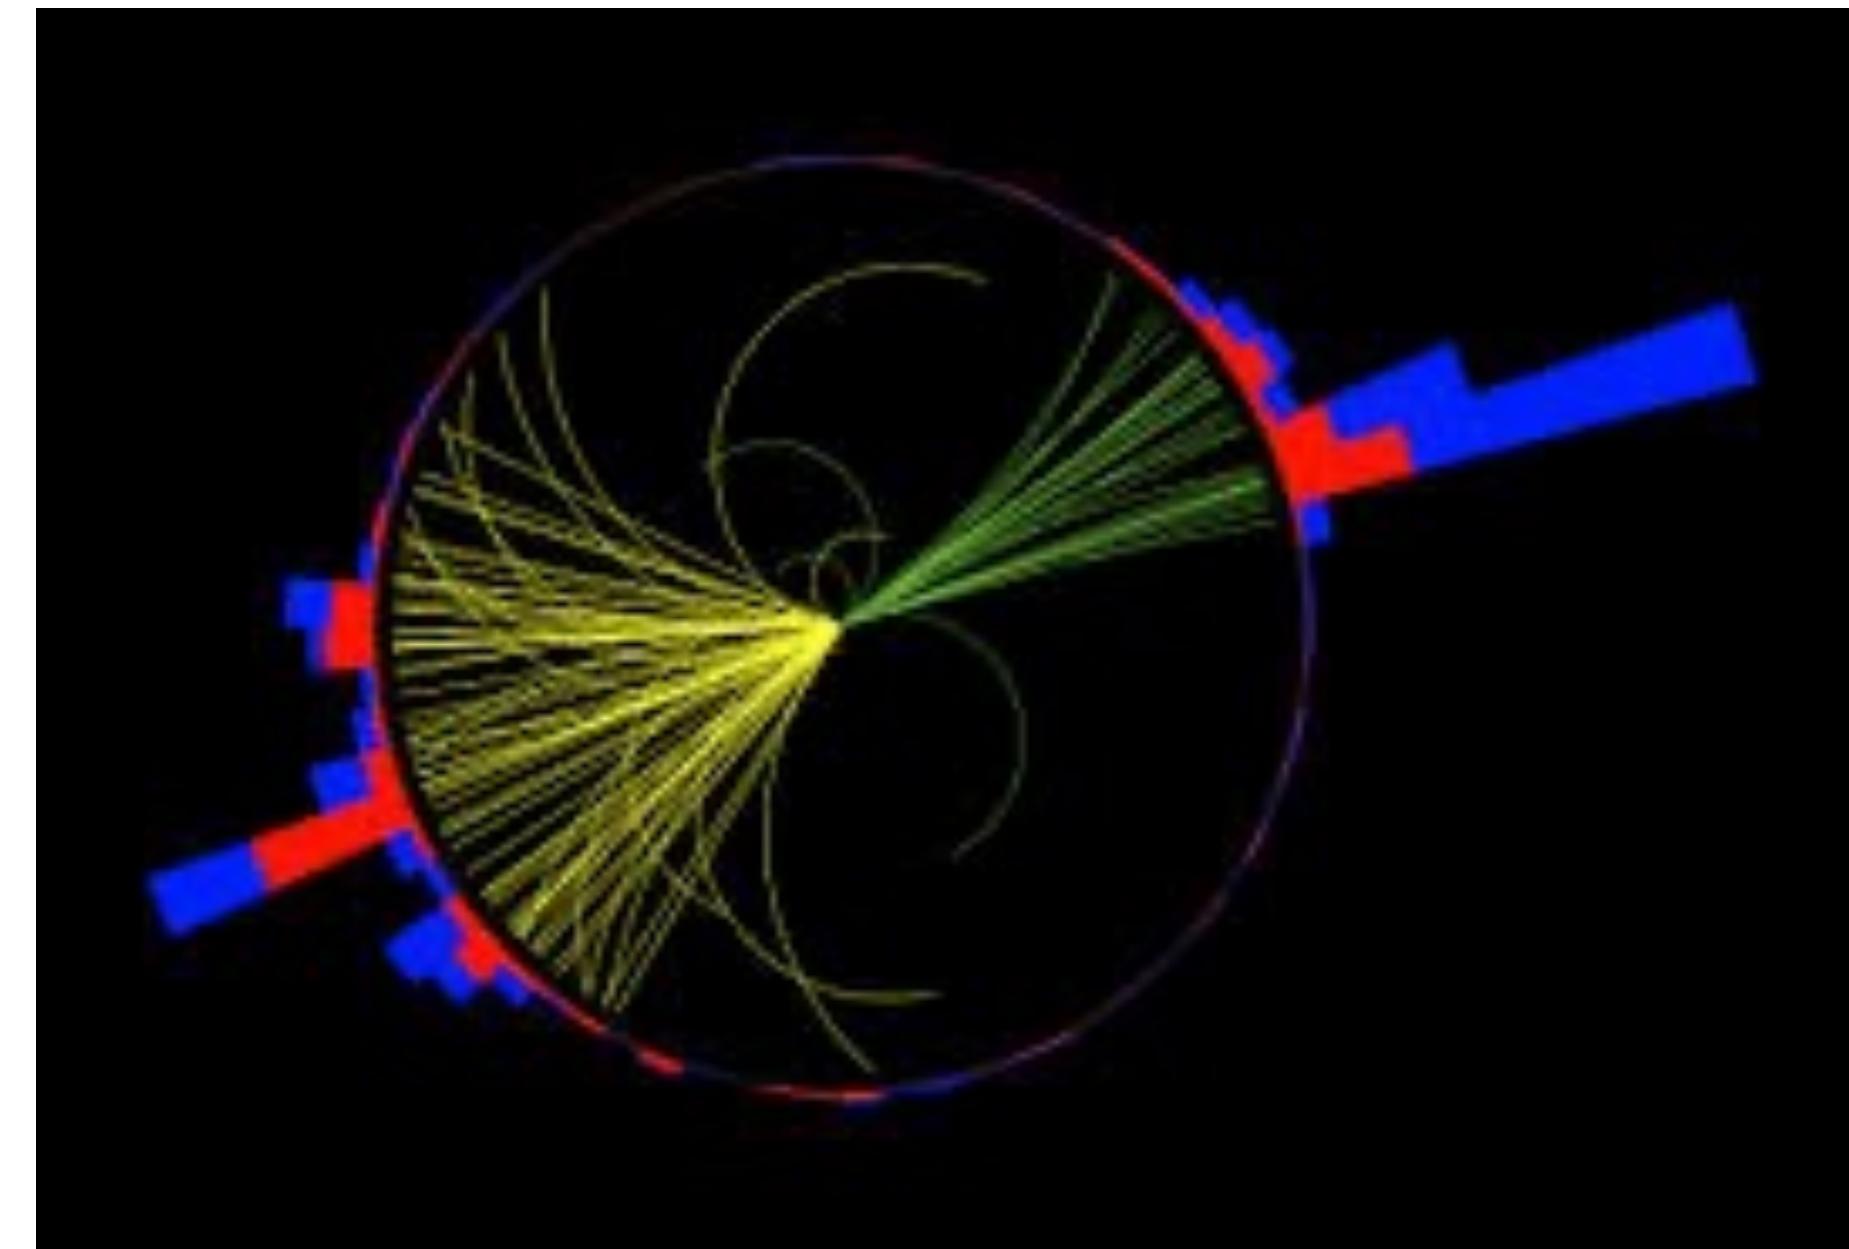
EXPERIMENTAL DATA



AT-TPC



CLAS 12



CMS

LECTURE 1 TOPICS

- Computational graphs
- Gradient-descent optimization
- Logistic regression
- Regression neural networks

MICHELLE KUCHERA
DAVIDSON COLLEGE

CPS-FR
MIT
25 AUGUST 2023

GOALS

- Each of us learns something today
- Stop me with any questions

MICHELLE KUCHERA
DAVIDSON COLLEGE

CPS-FR
MIT
25 AUGUST 2023

COMMUNITY

- Each of you arrived here with your own backgrounds, specialty, and path in life
- Your experience and expertise are valuable here, no matter what it is
- If the activity is within your background, help others!
- If you are totally (or a little) lost, ask for help!
- It is our shared goal to have **each** of us leave with some new skill/knowledge/understanding

Without Machine Learning



With Machine Learning

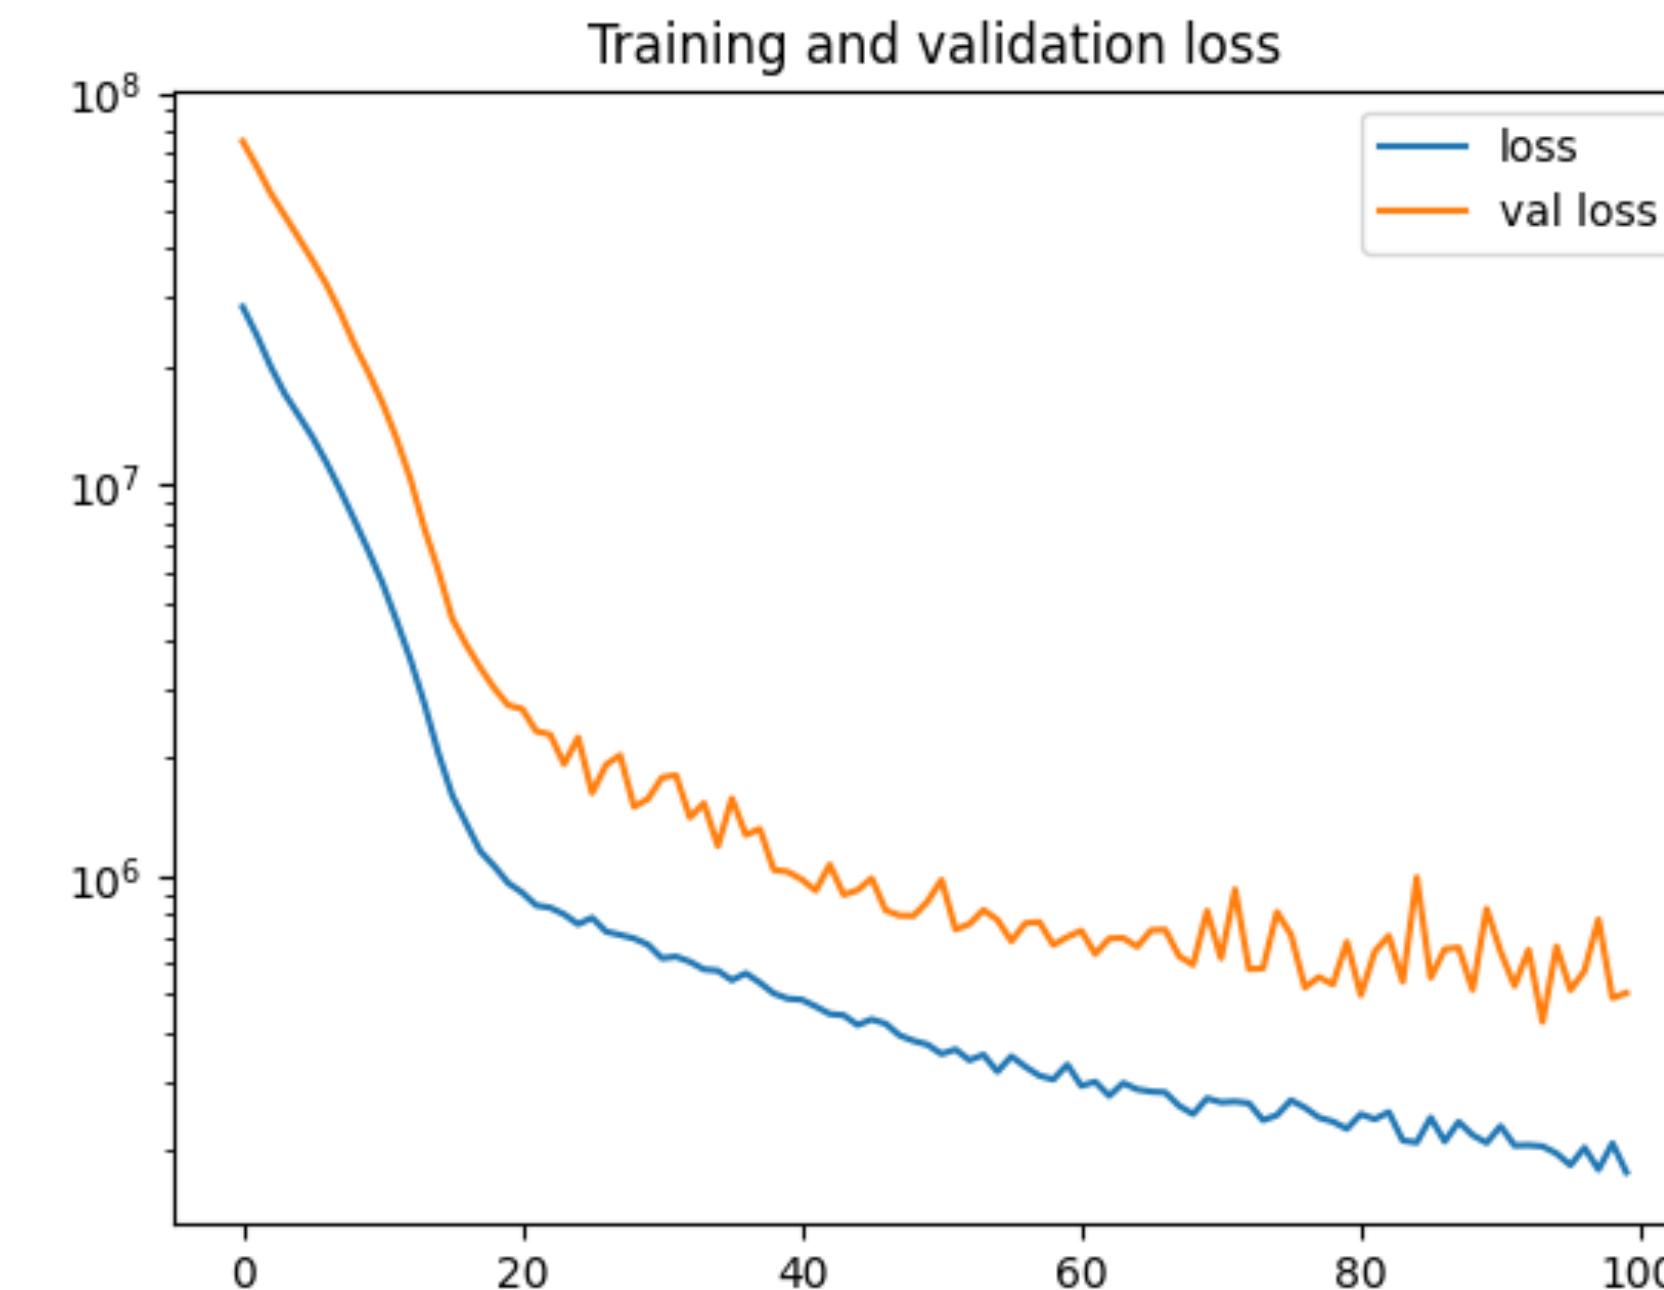
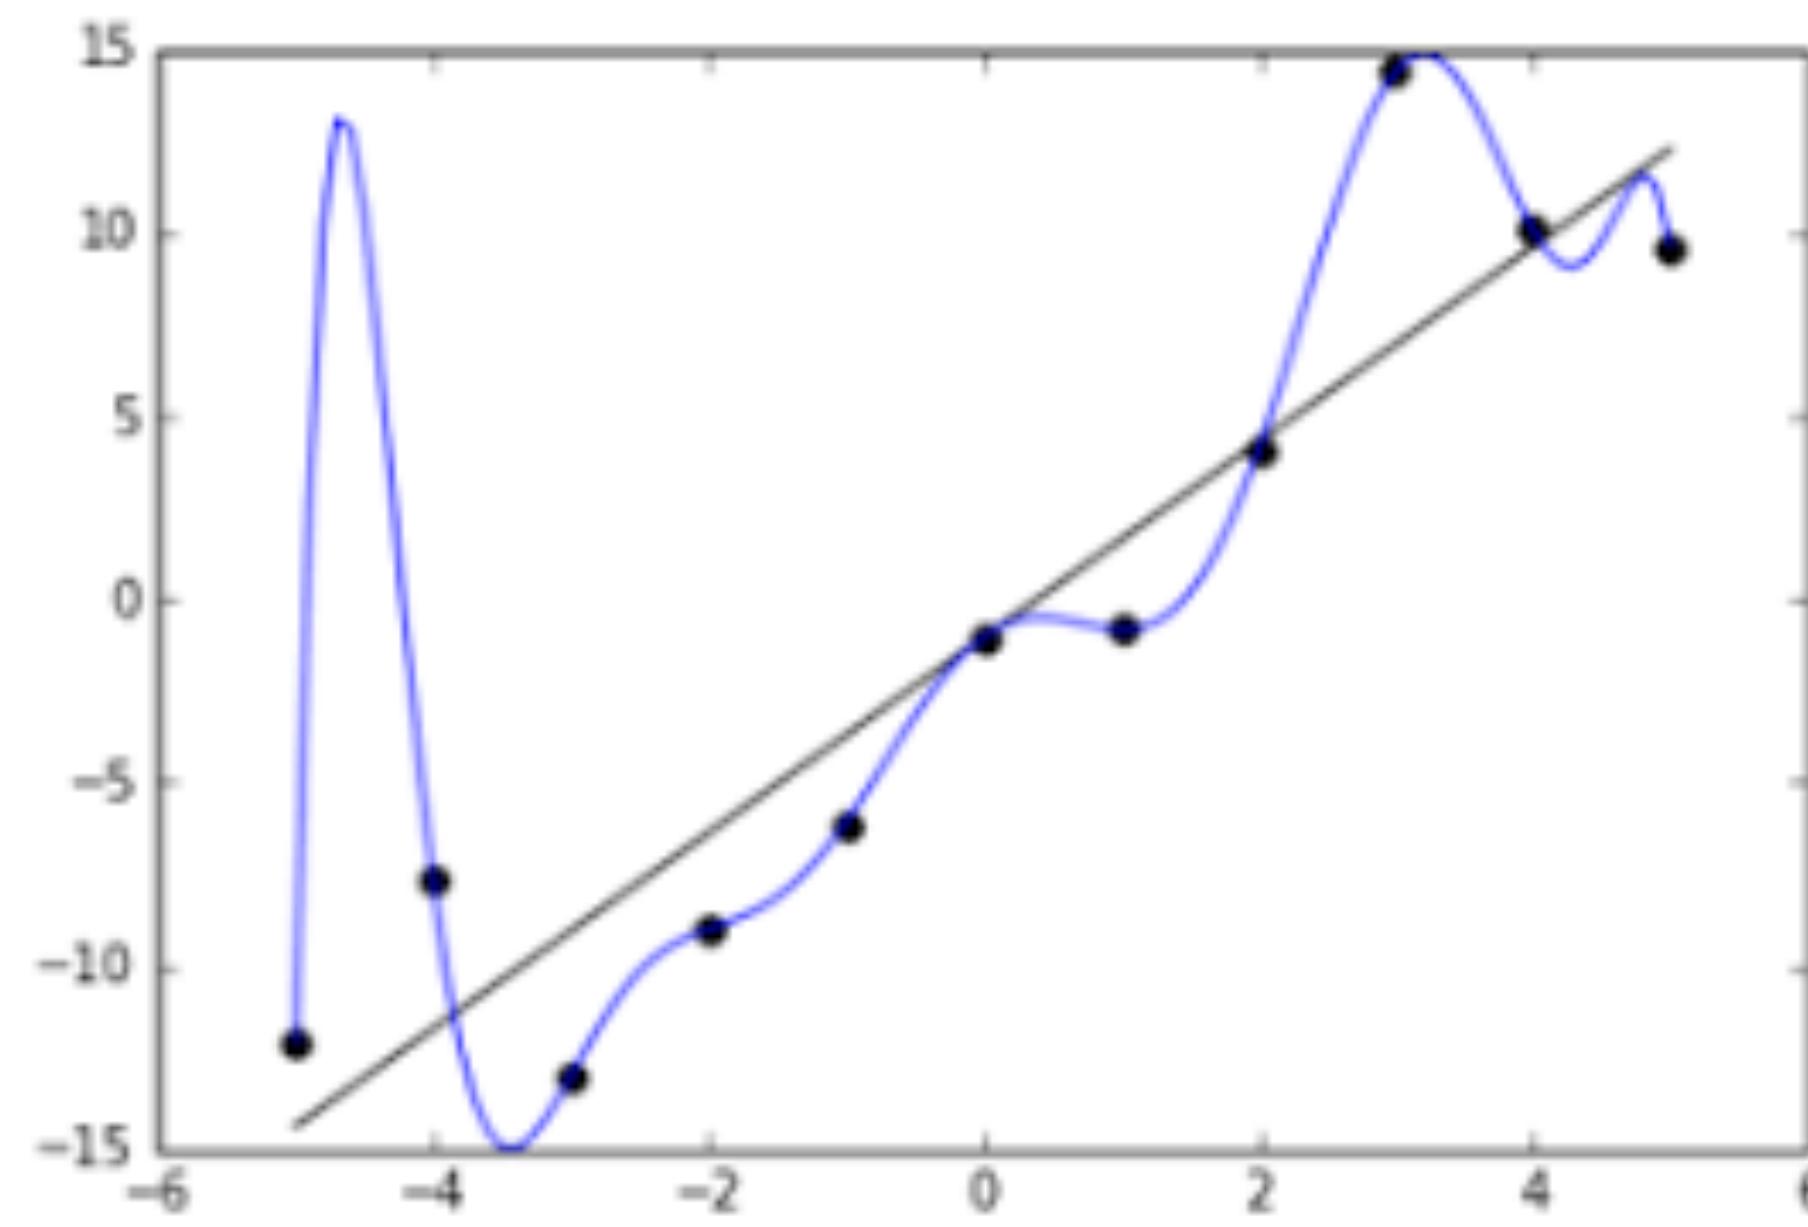


Without Machine Learning

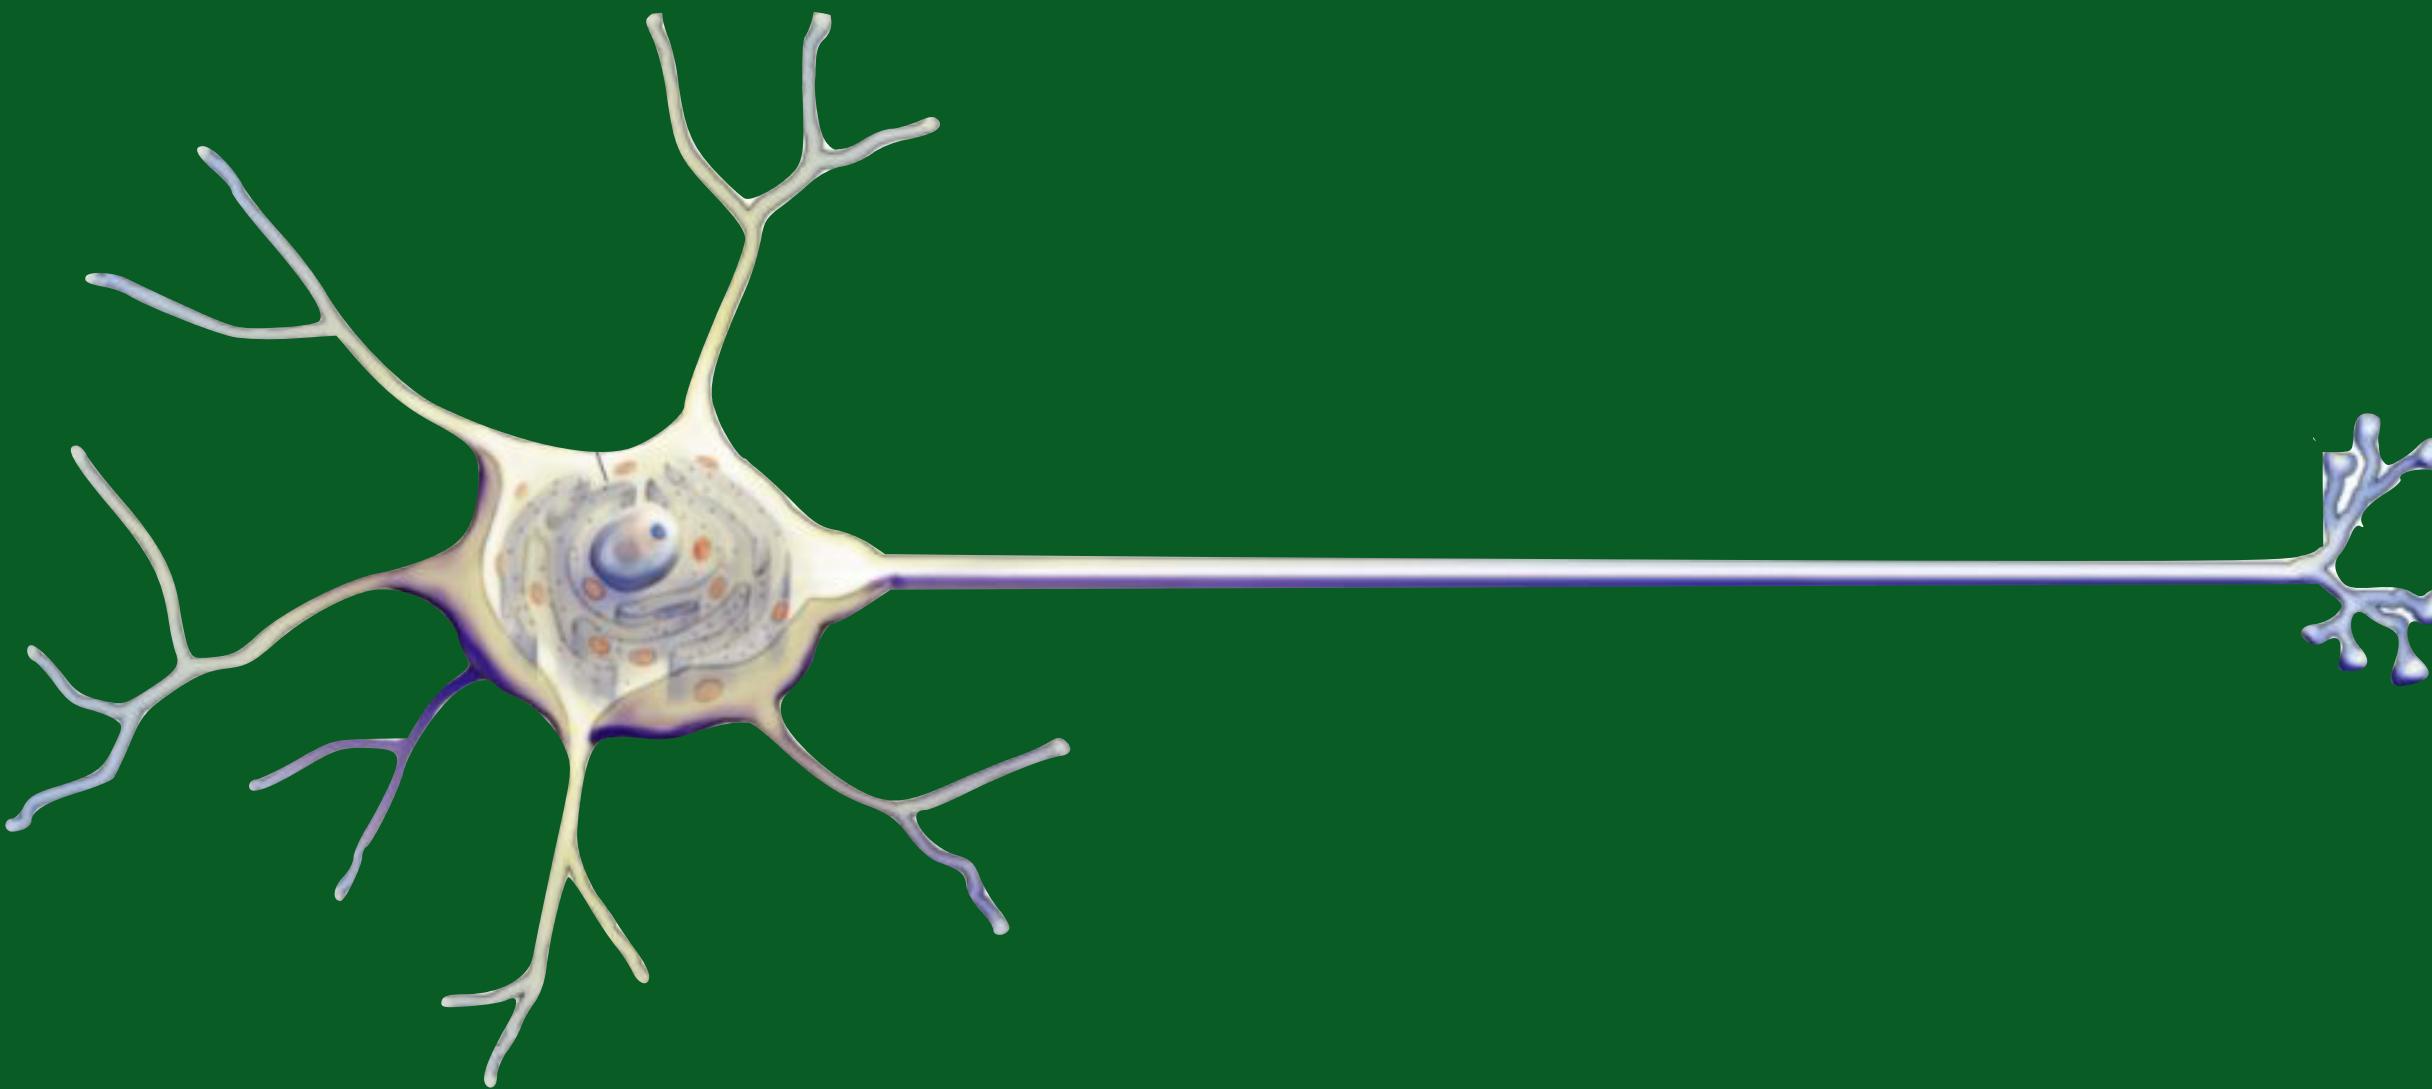
With Machine Learning

Learning from data is a **paradigm shift** in thinking about predictive models

/ \ * VERY SPECIFIC
INSTRUCTIONS



NEURON



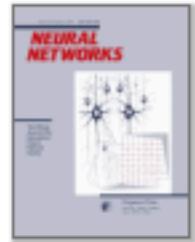
MATHEMATICS



ELSEVIER

Neural Networks

Volume 4, Issue 2, 1991, Pages 251-257



Approximation capabilities of multilayer feedforward networks

Kurt Hornik

Show more

Share Cite

[https://doi.org/10.1016/0893-6080\(91\)90009-T](https://doi.org/10.1016/0893-6080(91)90009-T)

[Get rights and content](#)

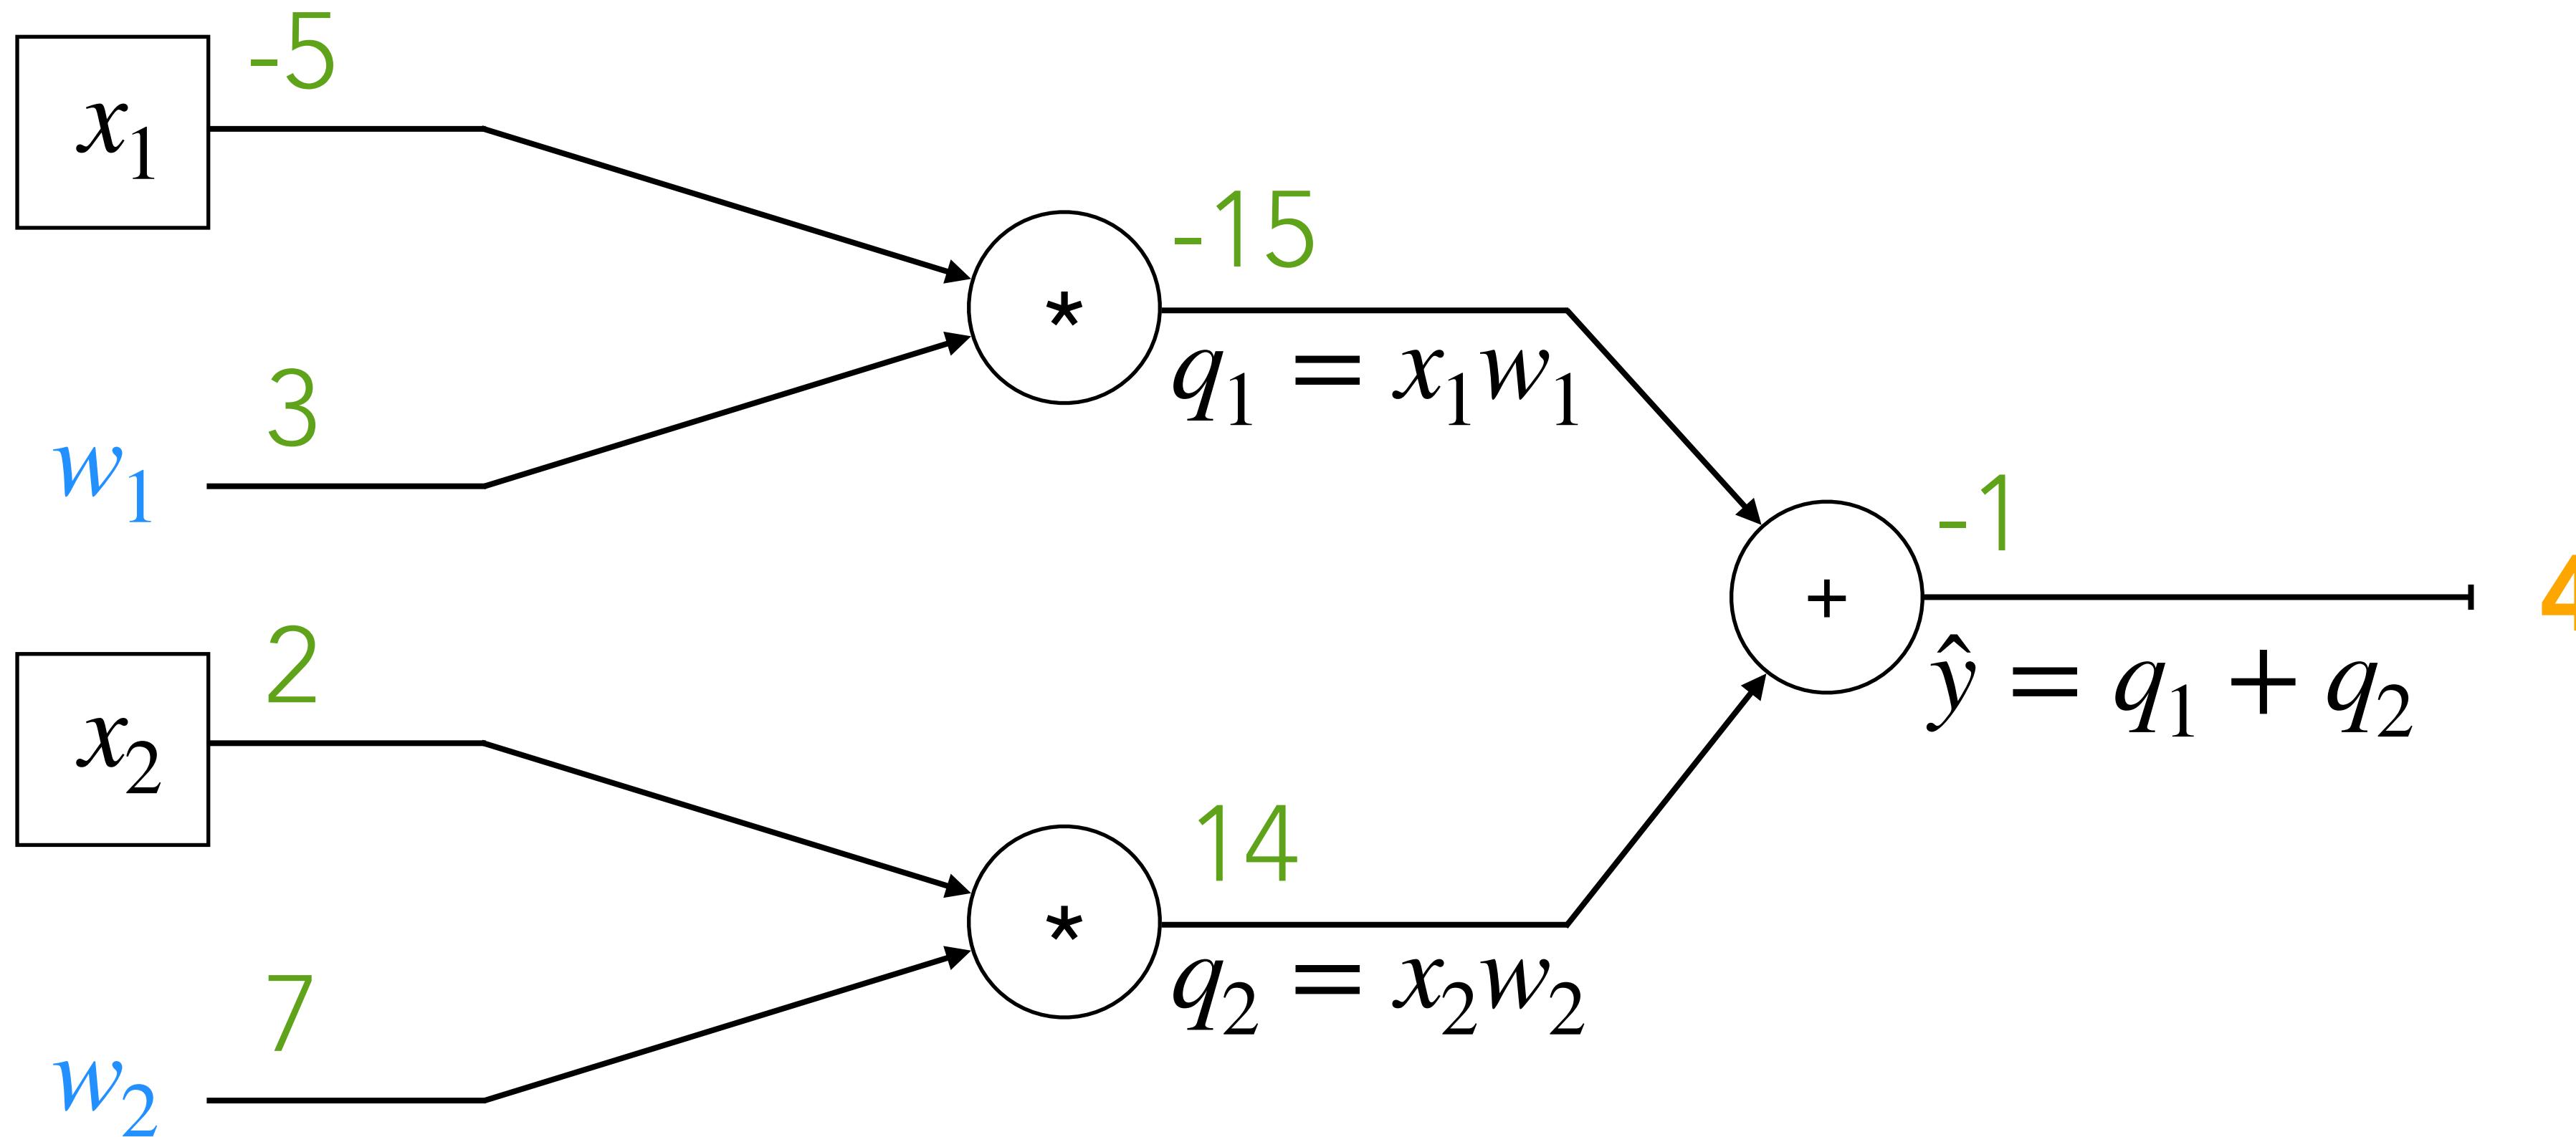
Abstract

We show that standard multilayer feedforward networks with as few as a single hidden layer and arbitrary bounded and nonconstant activation function are universal approximators with respect to $L^p(\mu)$ performance criteria, for arbitrary finite input environment measures μ , provided only that sufficiently many hidden units are available. If the activation function is continuous, bounded and nonconstant, then continuous mappings can be learned uniformly over compact input sets. We also give very general conditions ensuring that networks with sufficiently smooth activation functions are capable of arbitrarily accurate approximation to a function and its derivatives.

MATHEMATICS

COMPUTATIONAL GRAPH

$$\hat{y} = x_1 w_1 + x_2 w_2$$

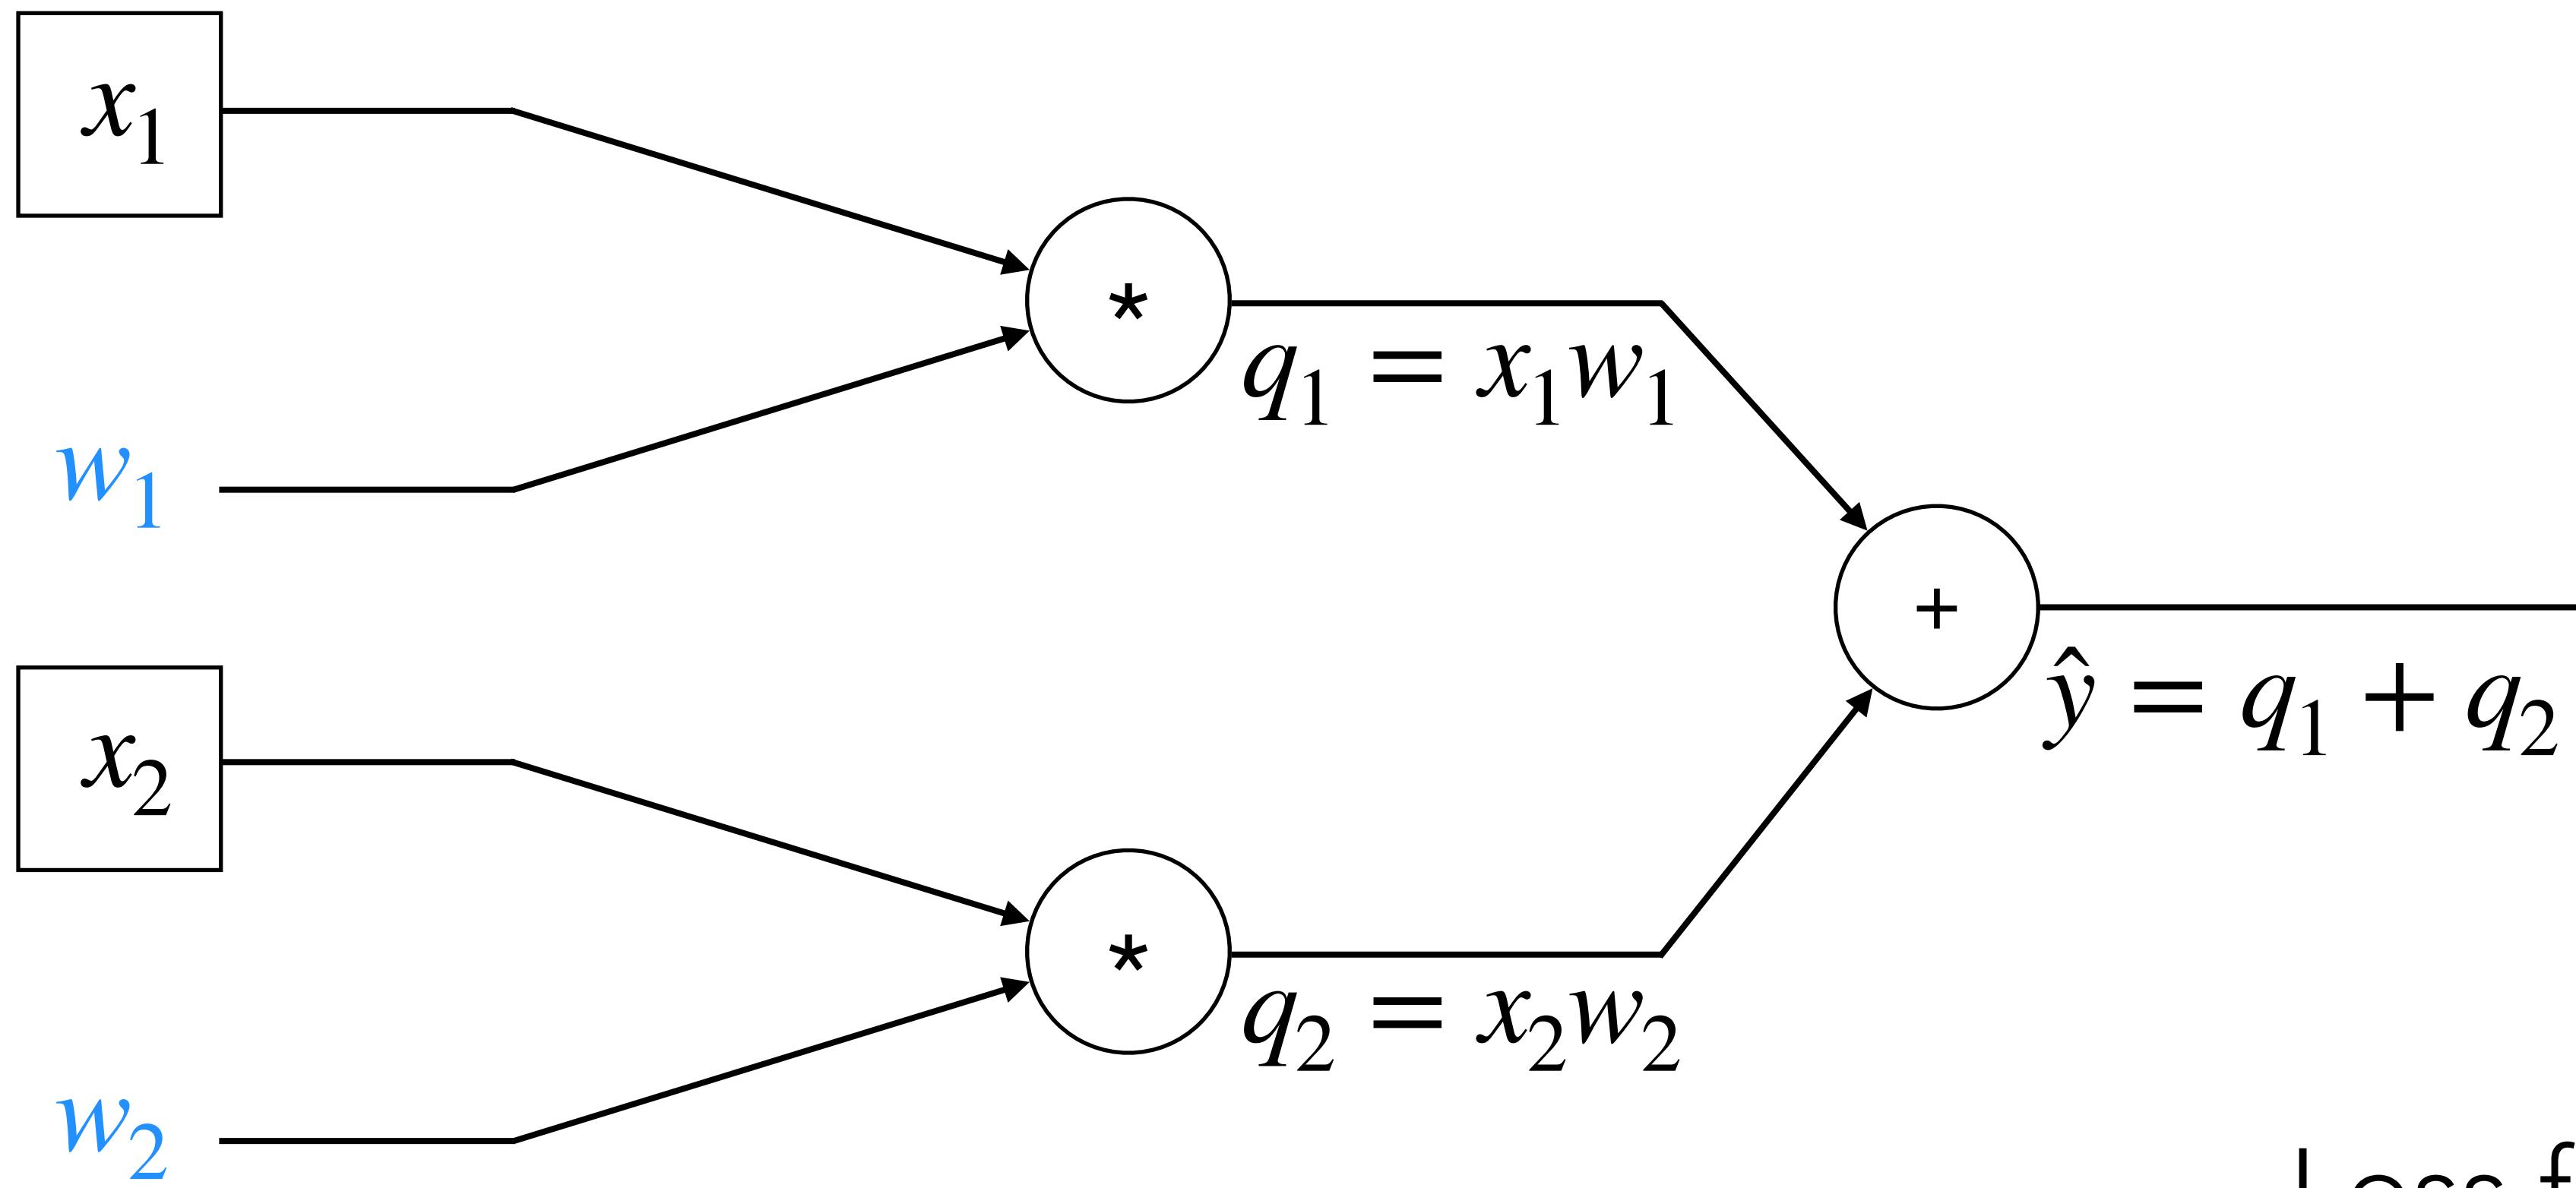


MACHINE LEARNING

SUPERVISED LEARNING

REGRESSION

$$\hat{y} = x_1 w_1 + x_2 w_2$$

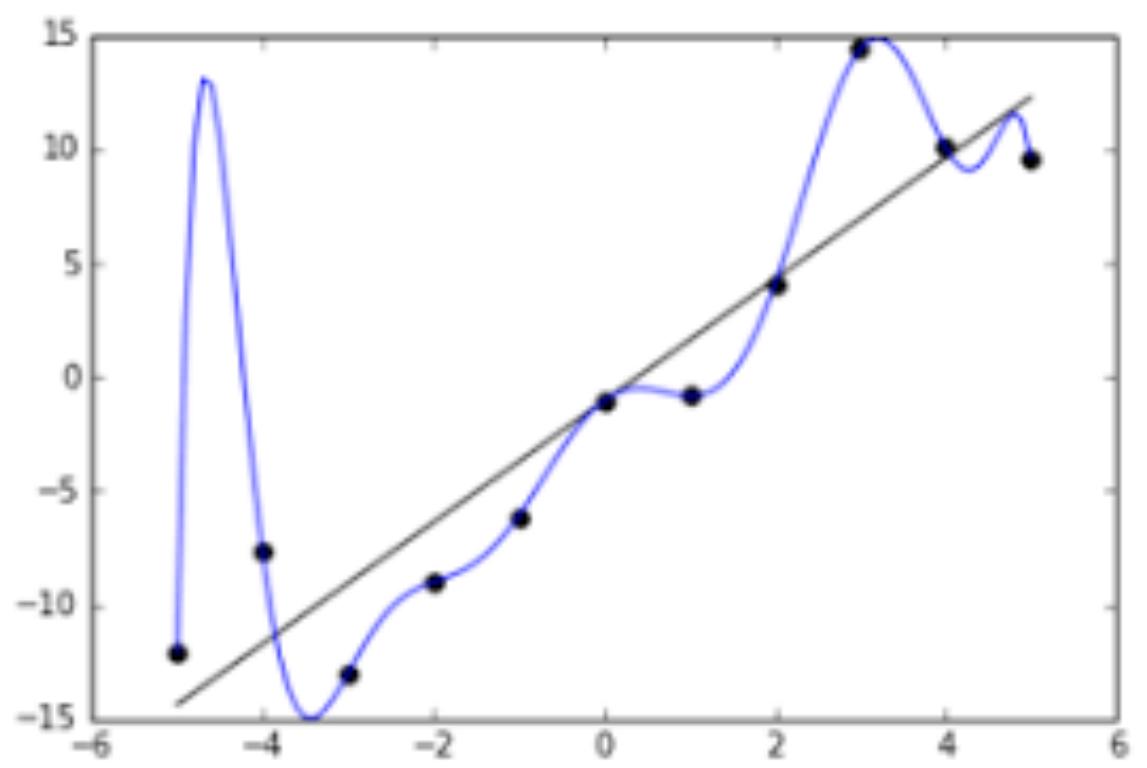
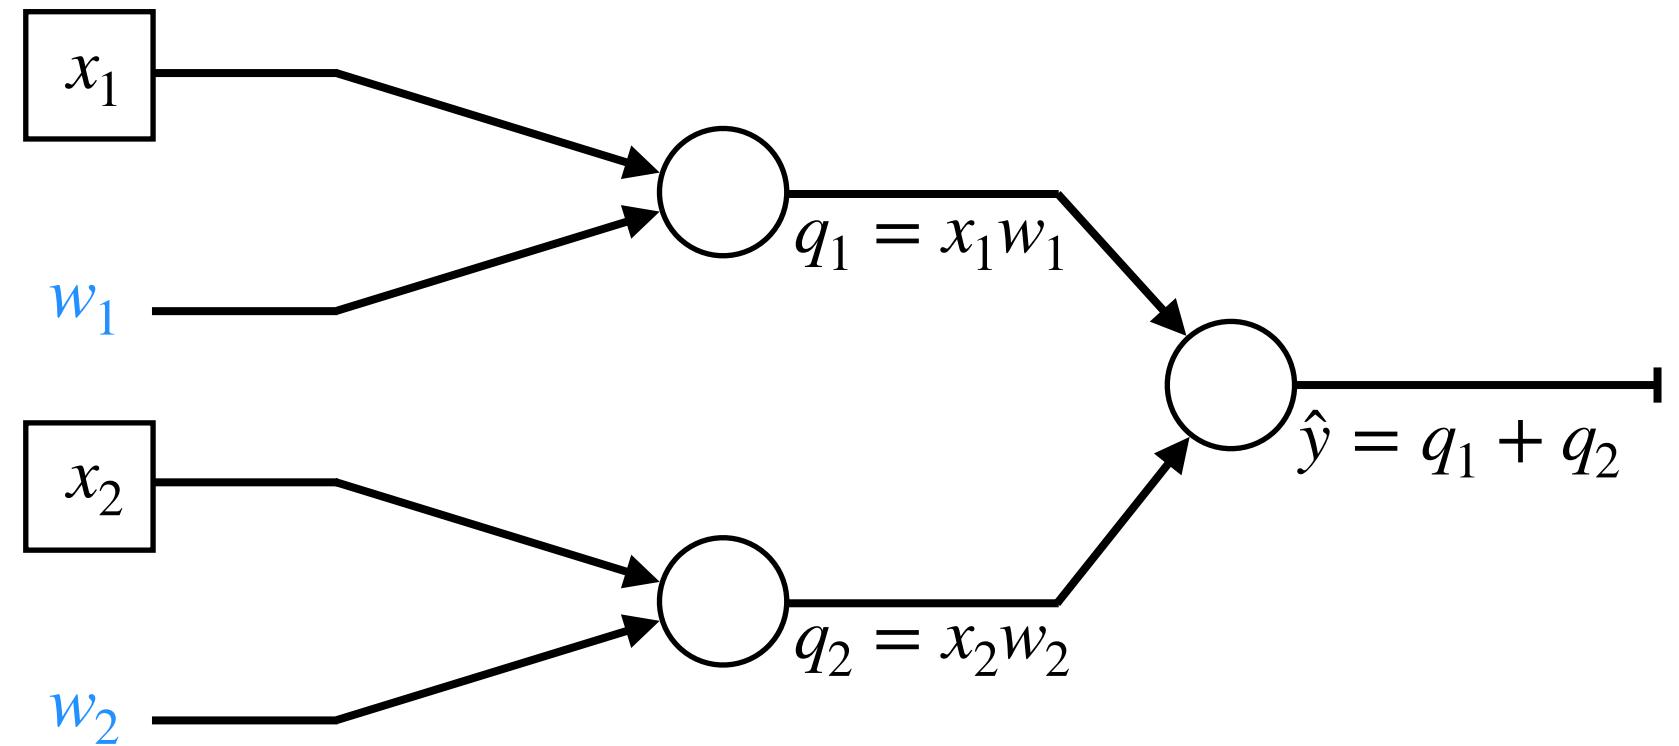


Loss function

$$J(w) = \hat{y} - y$$

SUPERVISED LEARNING

$$\hat{y} = x_1 w_1 + x_2 w_2$$



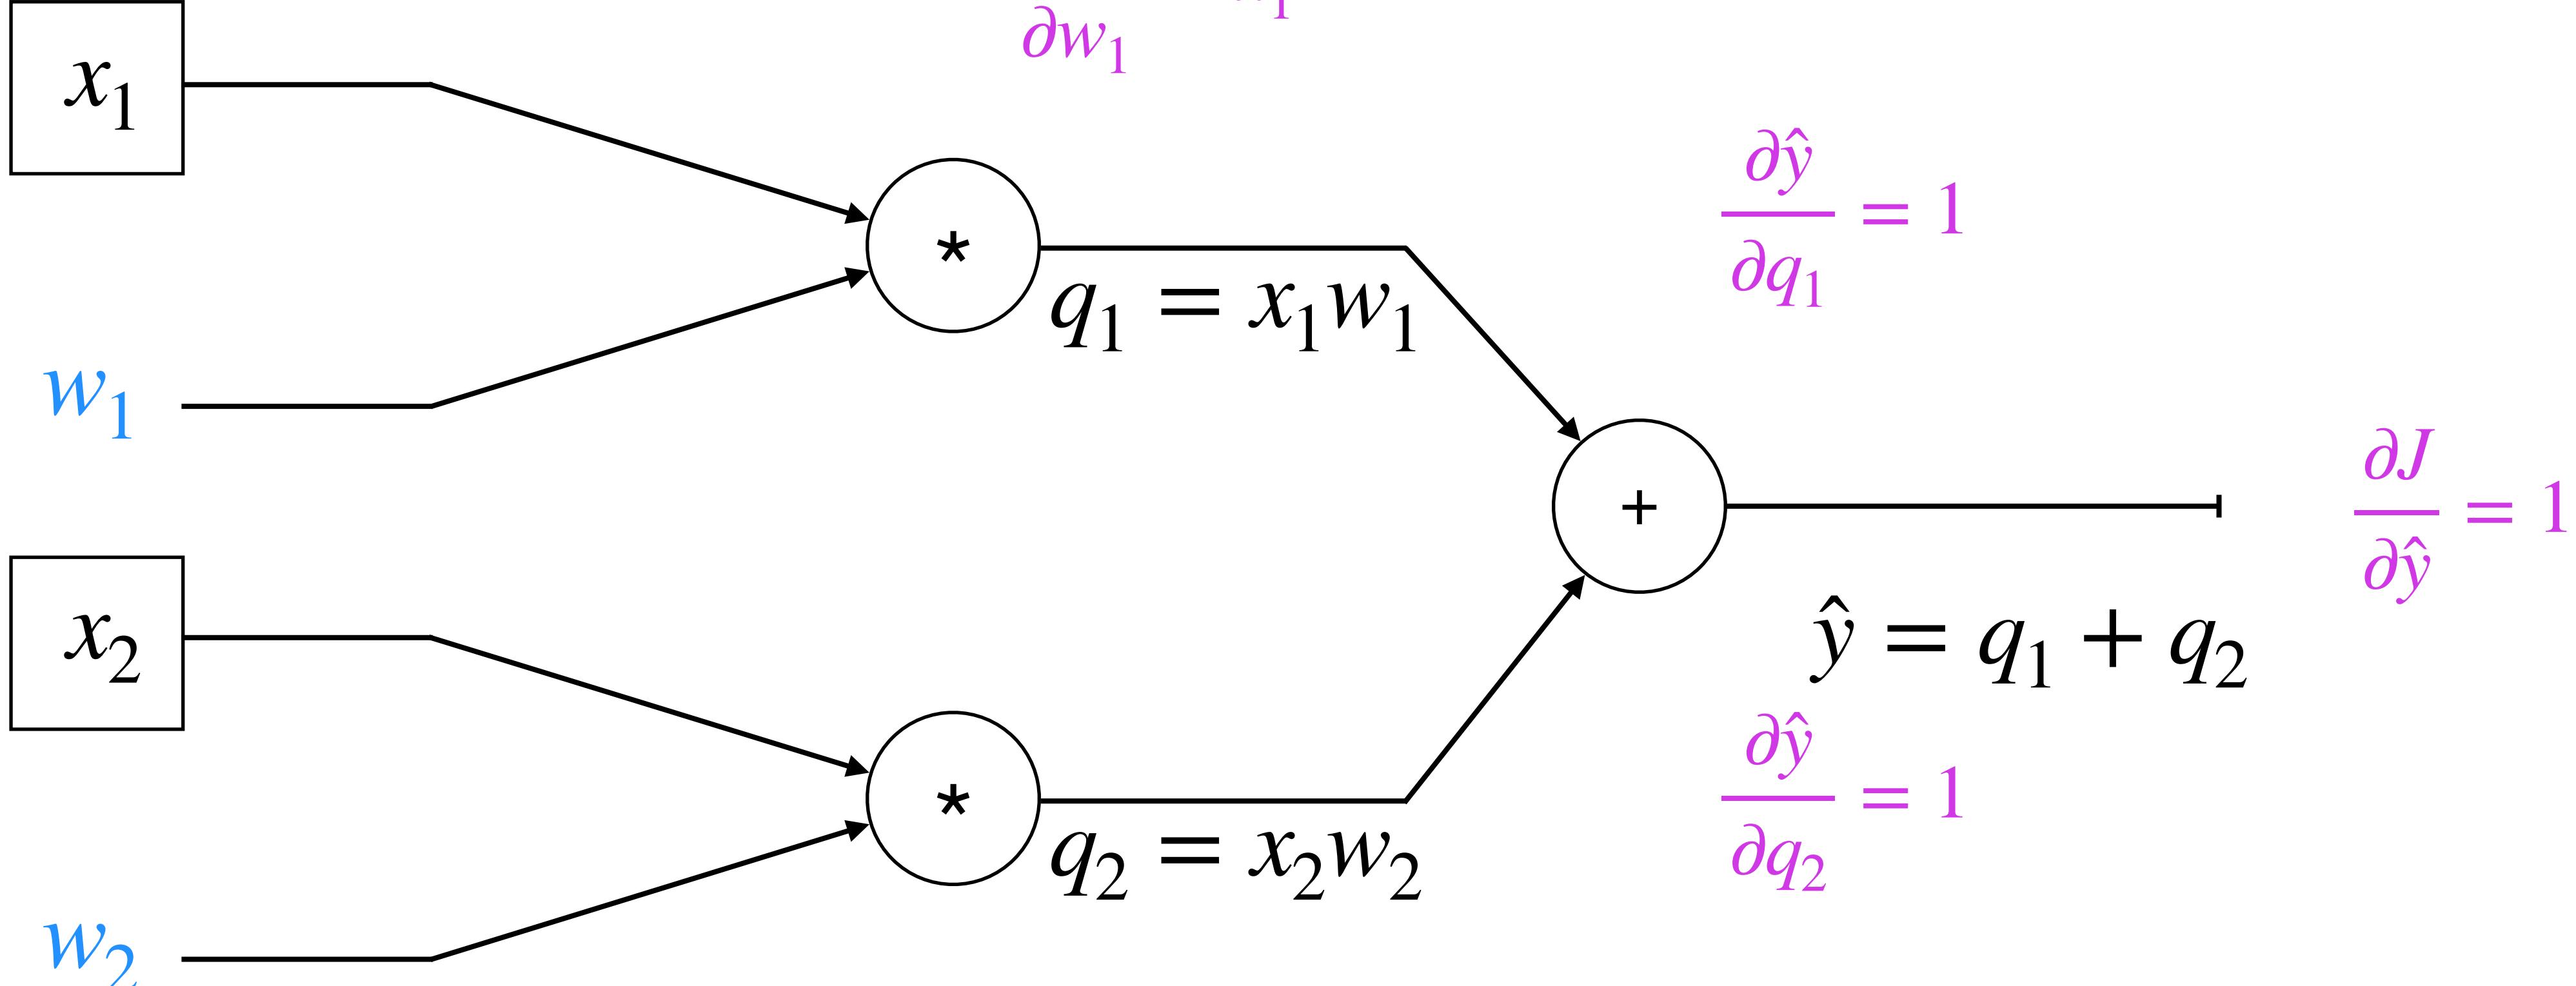
	Feature 1	Feature 2	Feature 3	Target
Example 1				
Example 2				
Example 3				
Example 4				

Loss function
MSE across N examples

$$J(w) = \frac{1}{N} \sum_{i=0}^N (\hat{y}_i - y_i)^2$$

BACKPROPAGATION

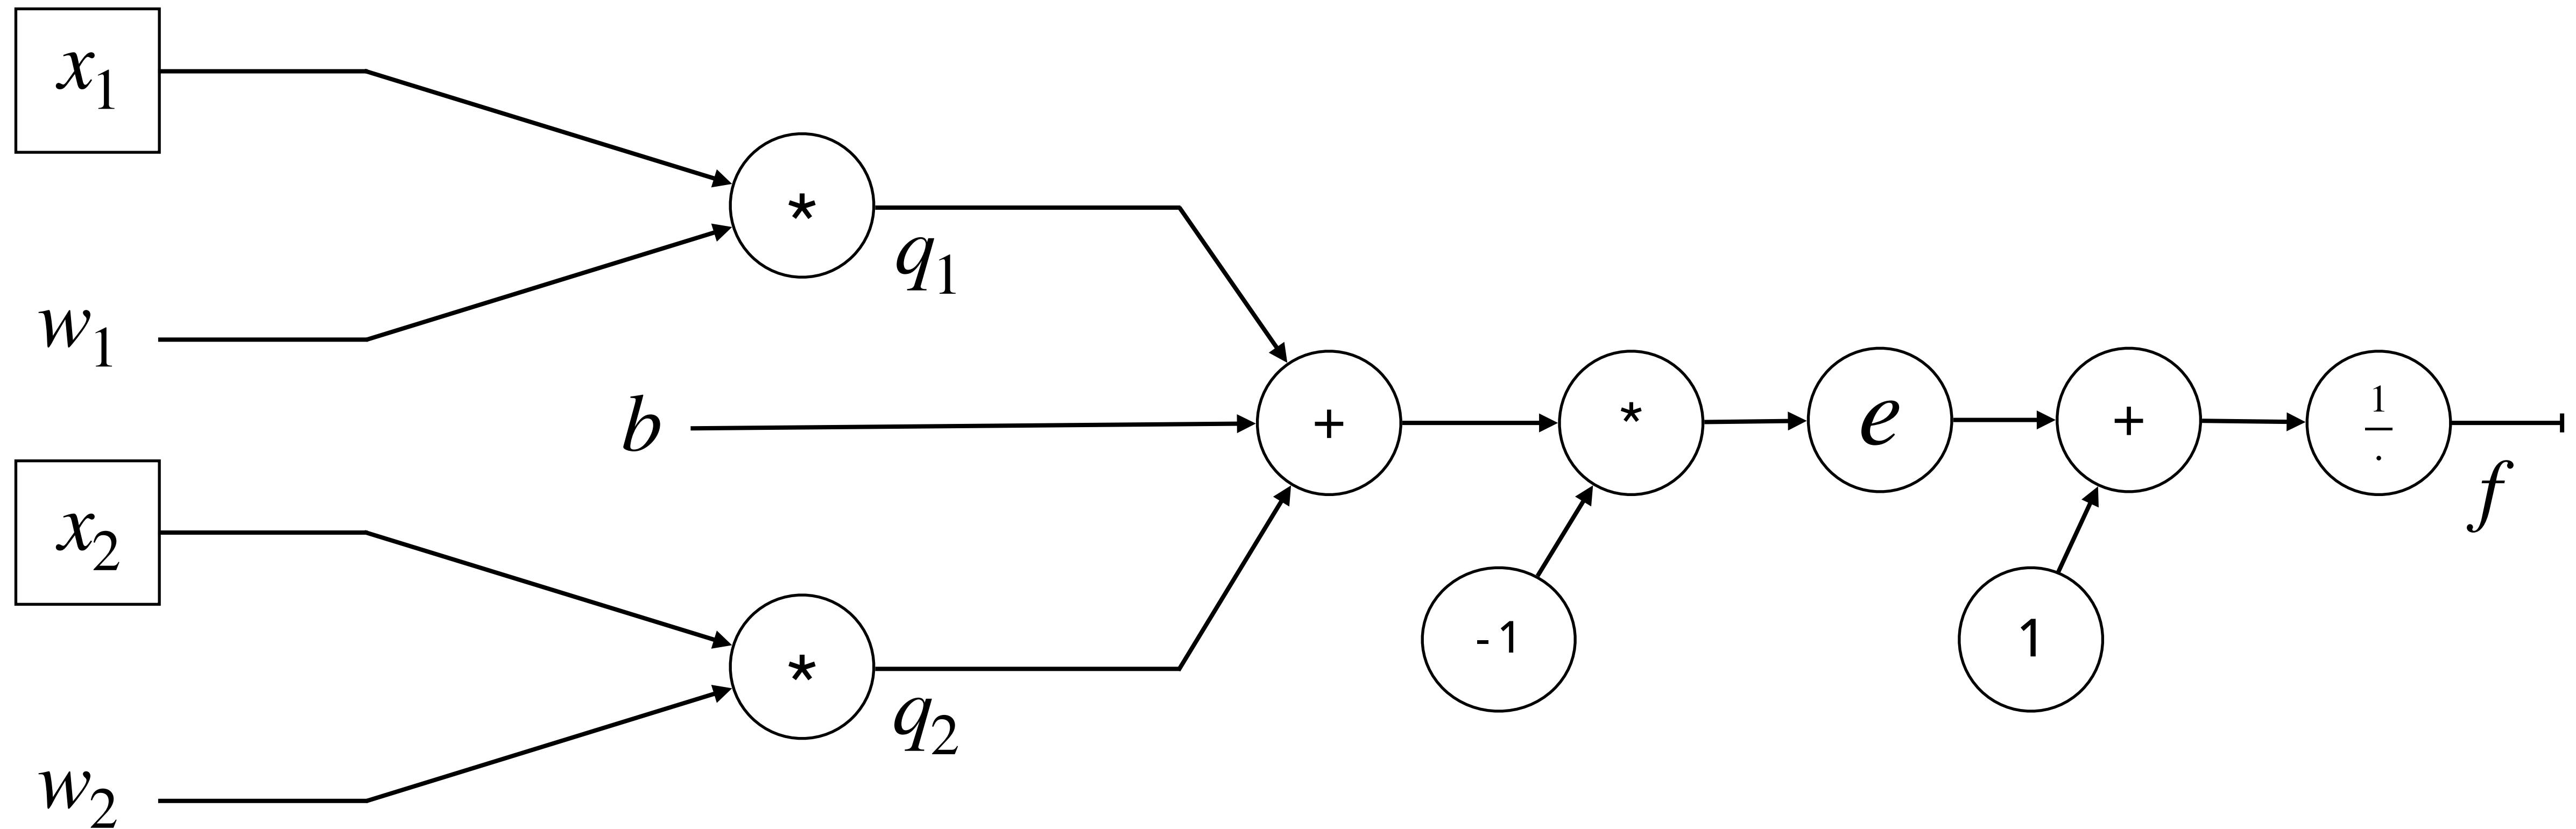
$$w_1 = w_1 - \eta * \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial q_1} \frac{\partial q_1}{\partial w_1}$$



$$w_2 = w_2 - \eta * \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial q_2} \frac{\partial q_2}{\partial w_2}$$

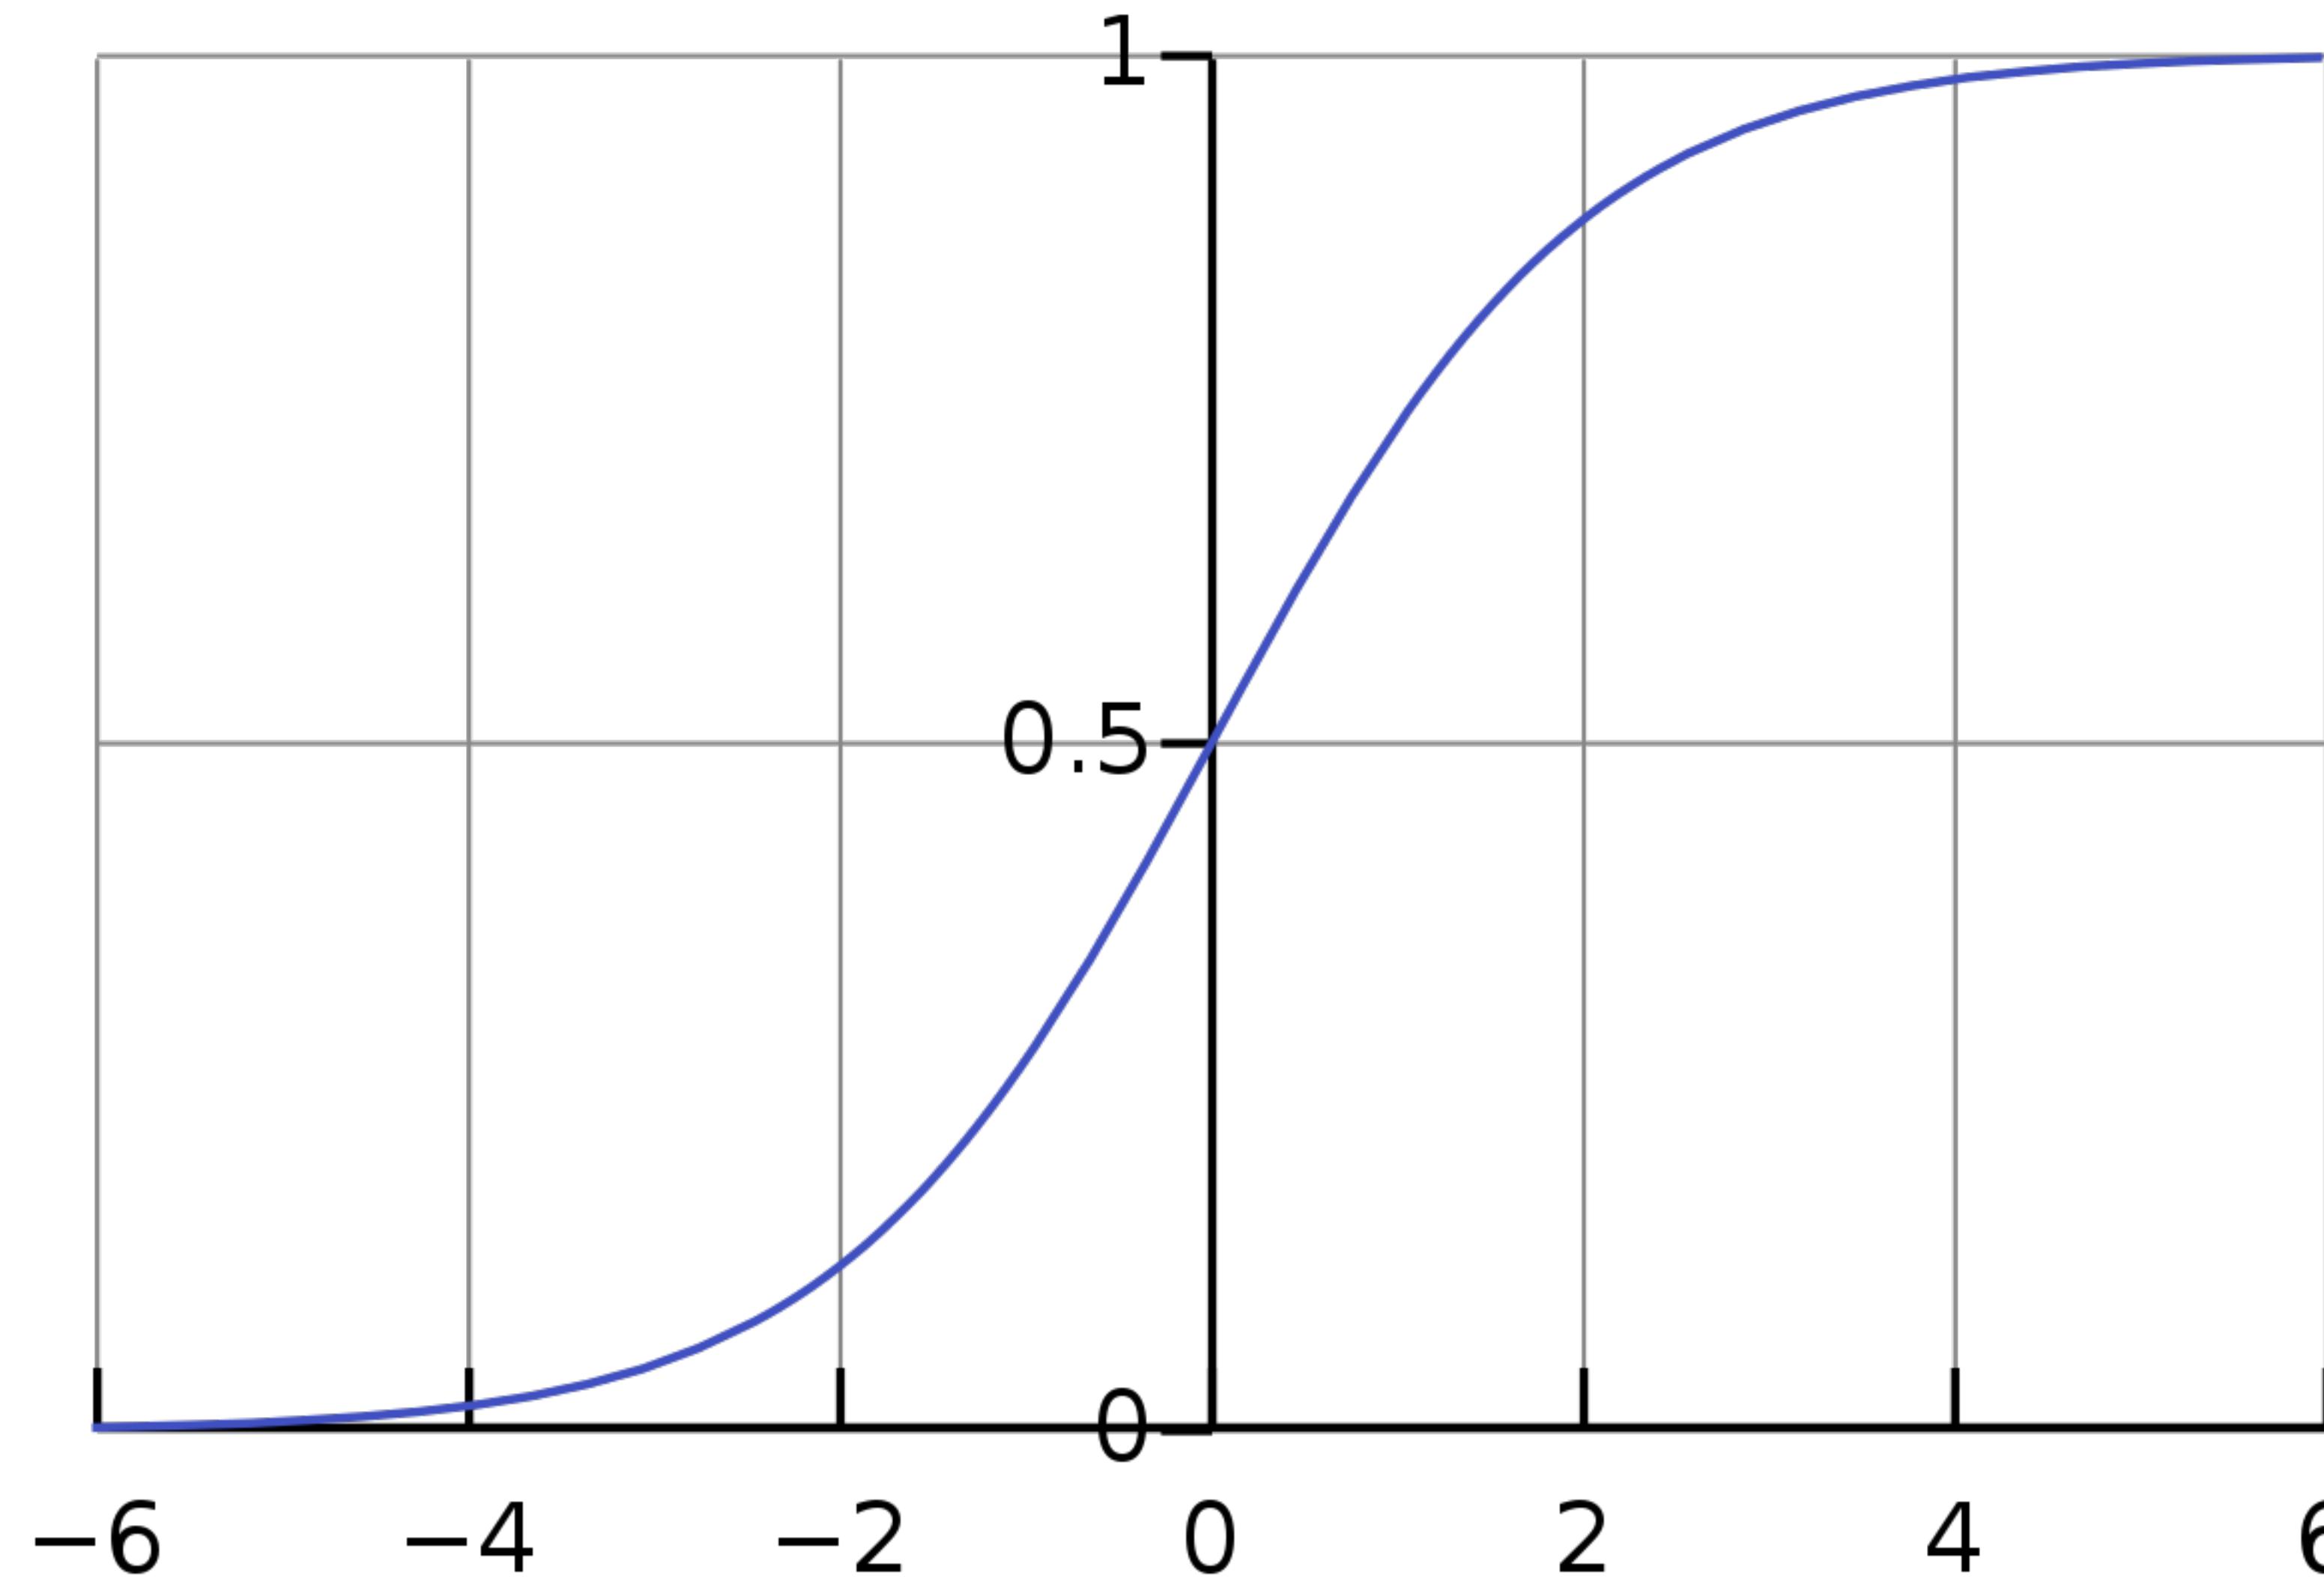
Loss function
 $J(w) = \hat{y} - y$

LOGISTIC REGRESSION

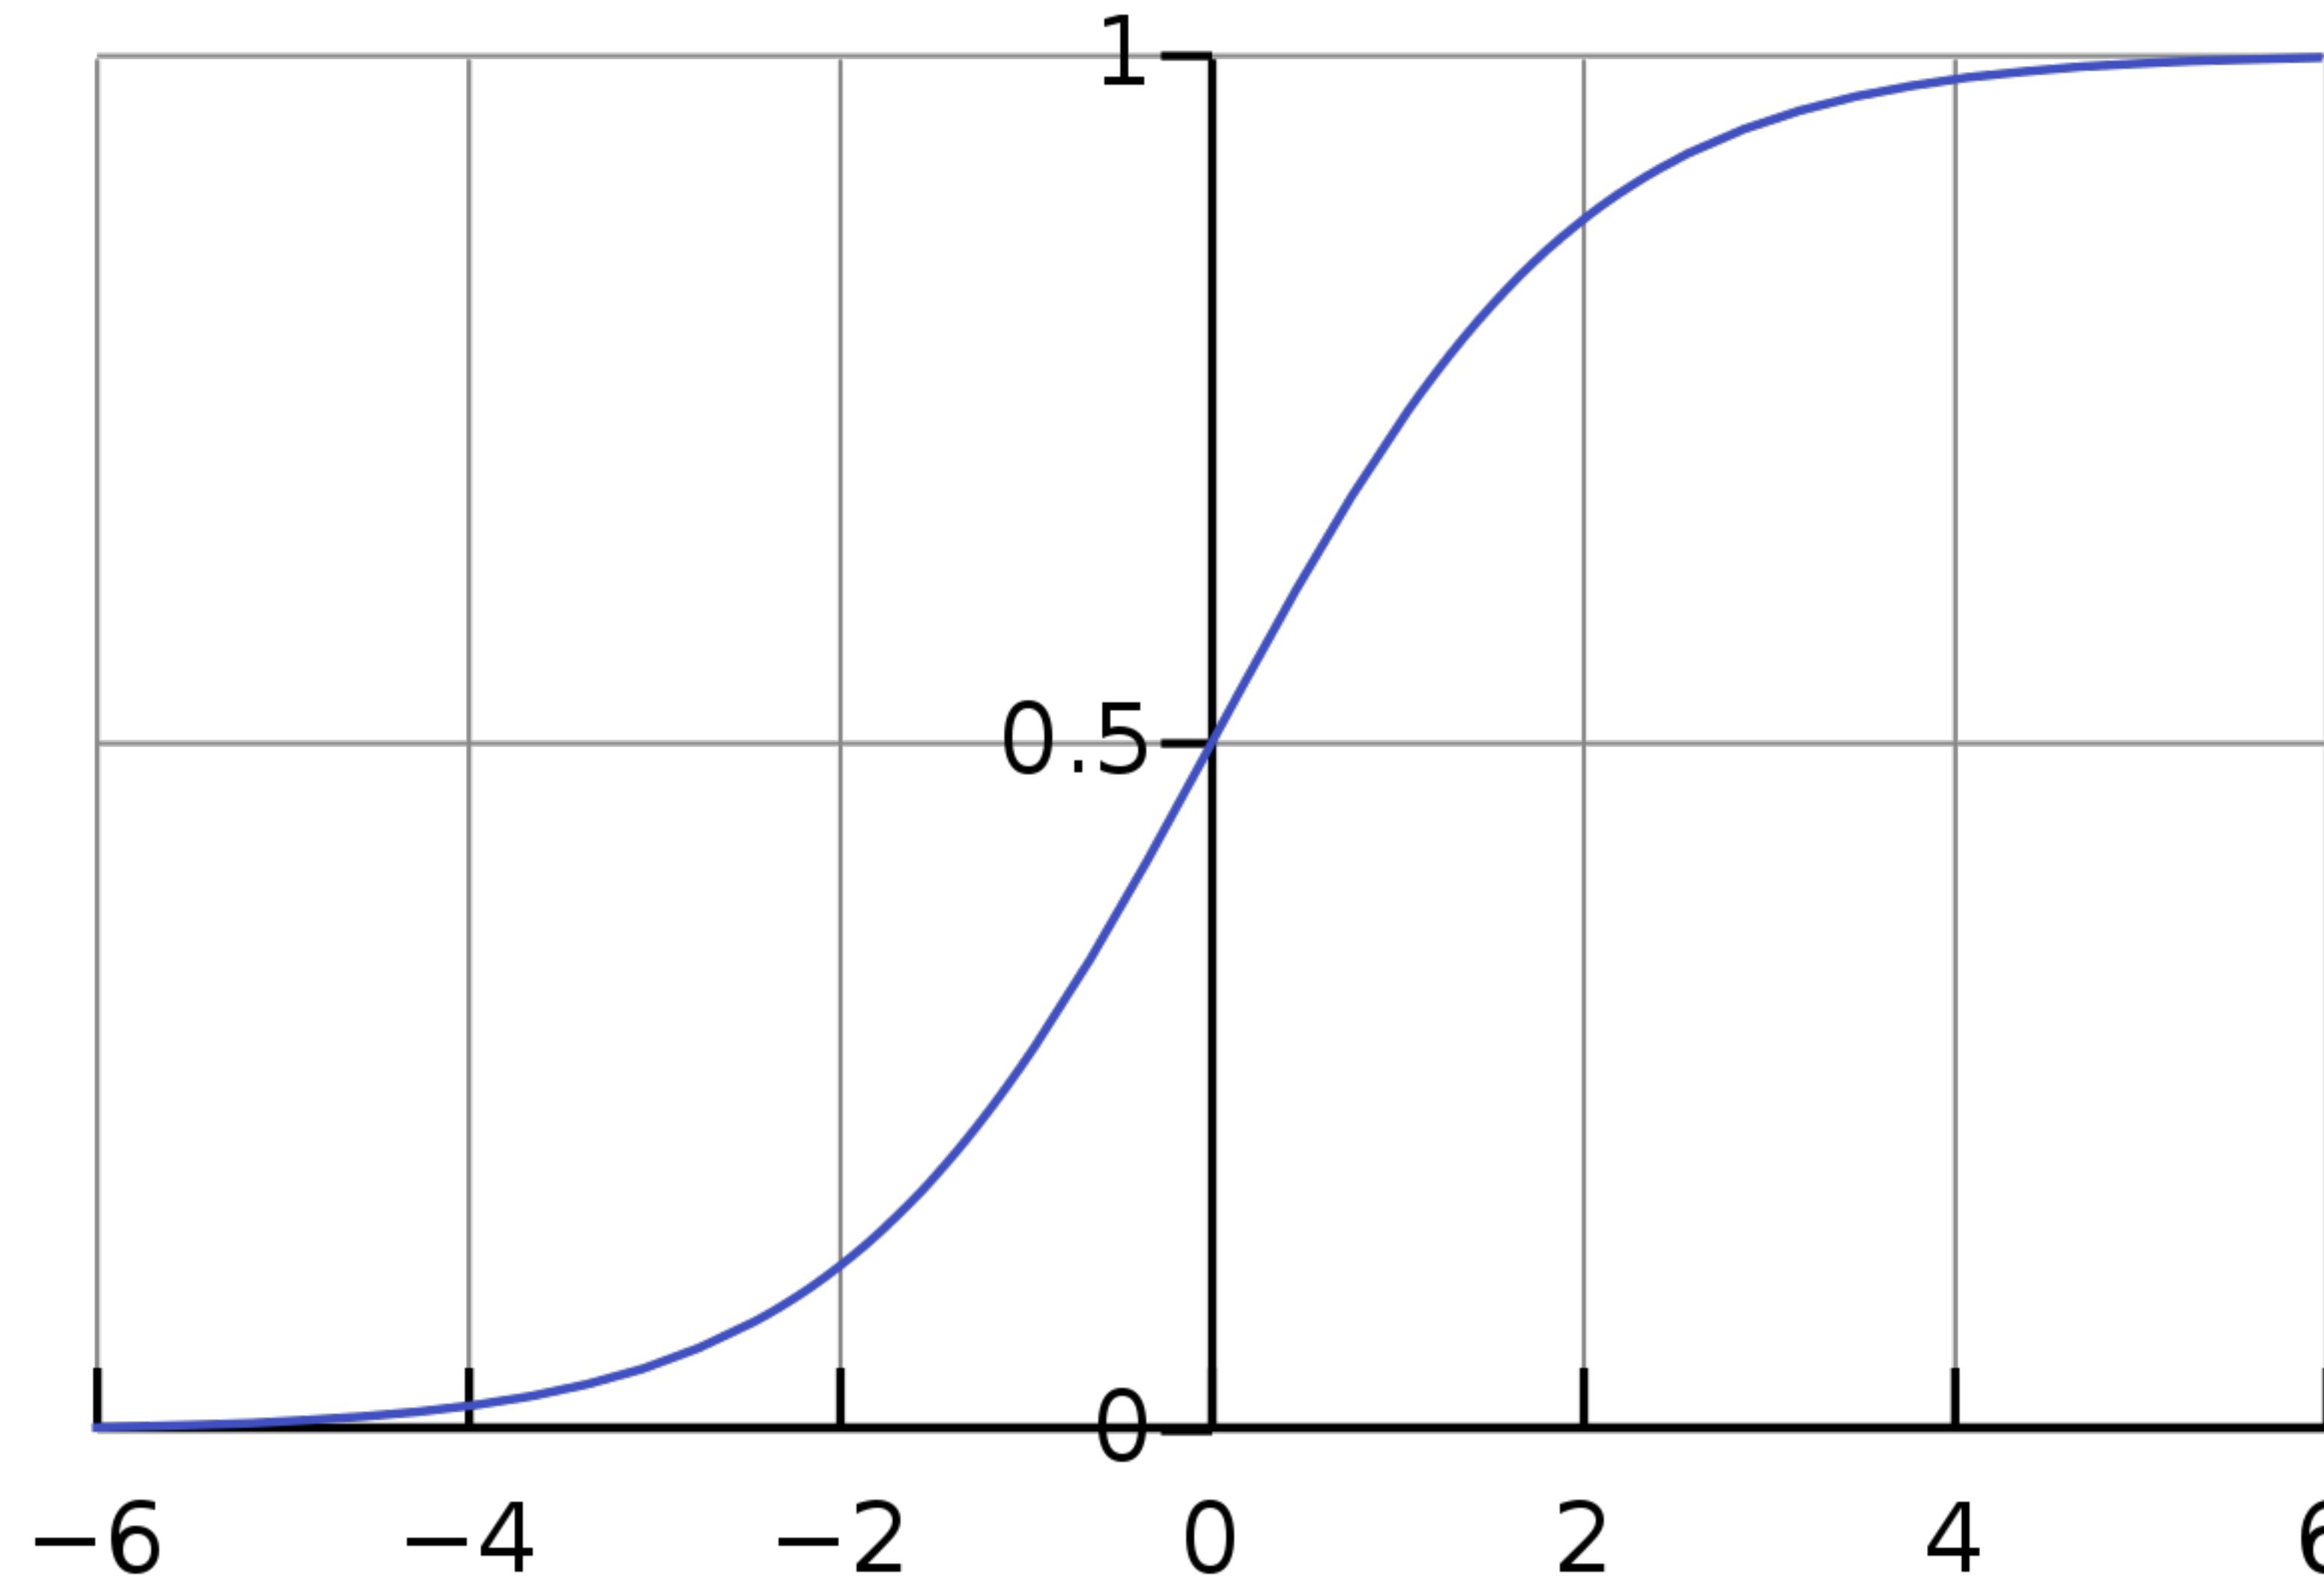


$$f = \frac{1}{1 + e^{-(x_1 w_1 + x_2 w_2 + b)}}$$

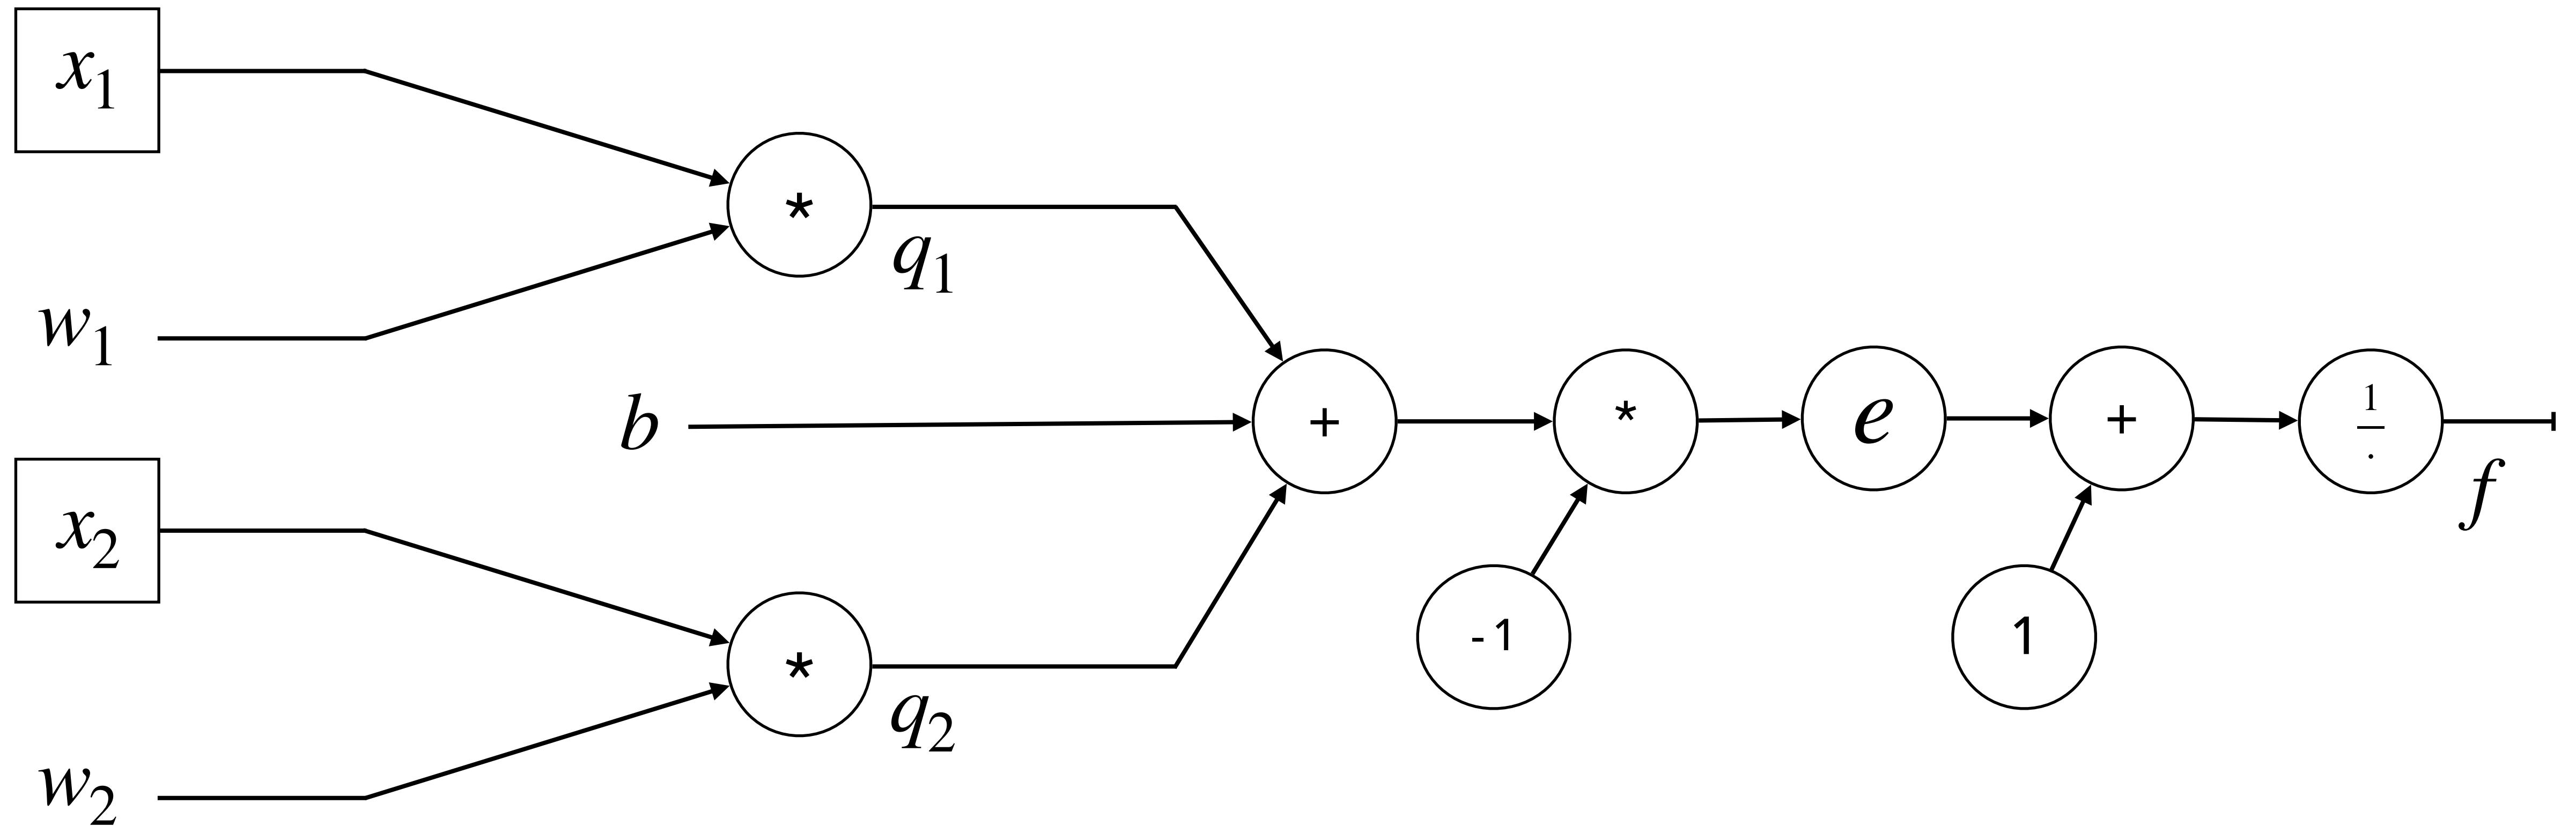
LOGISTIC REGRESSION



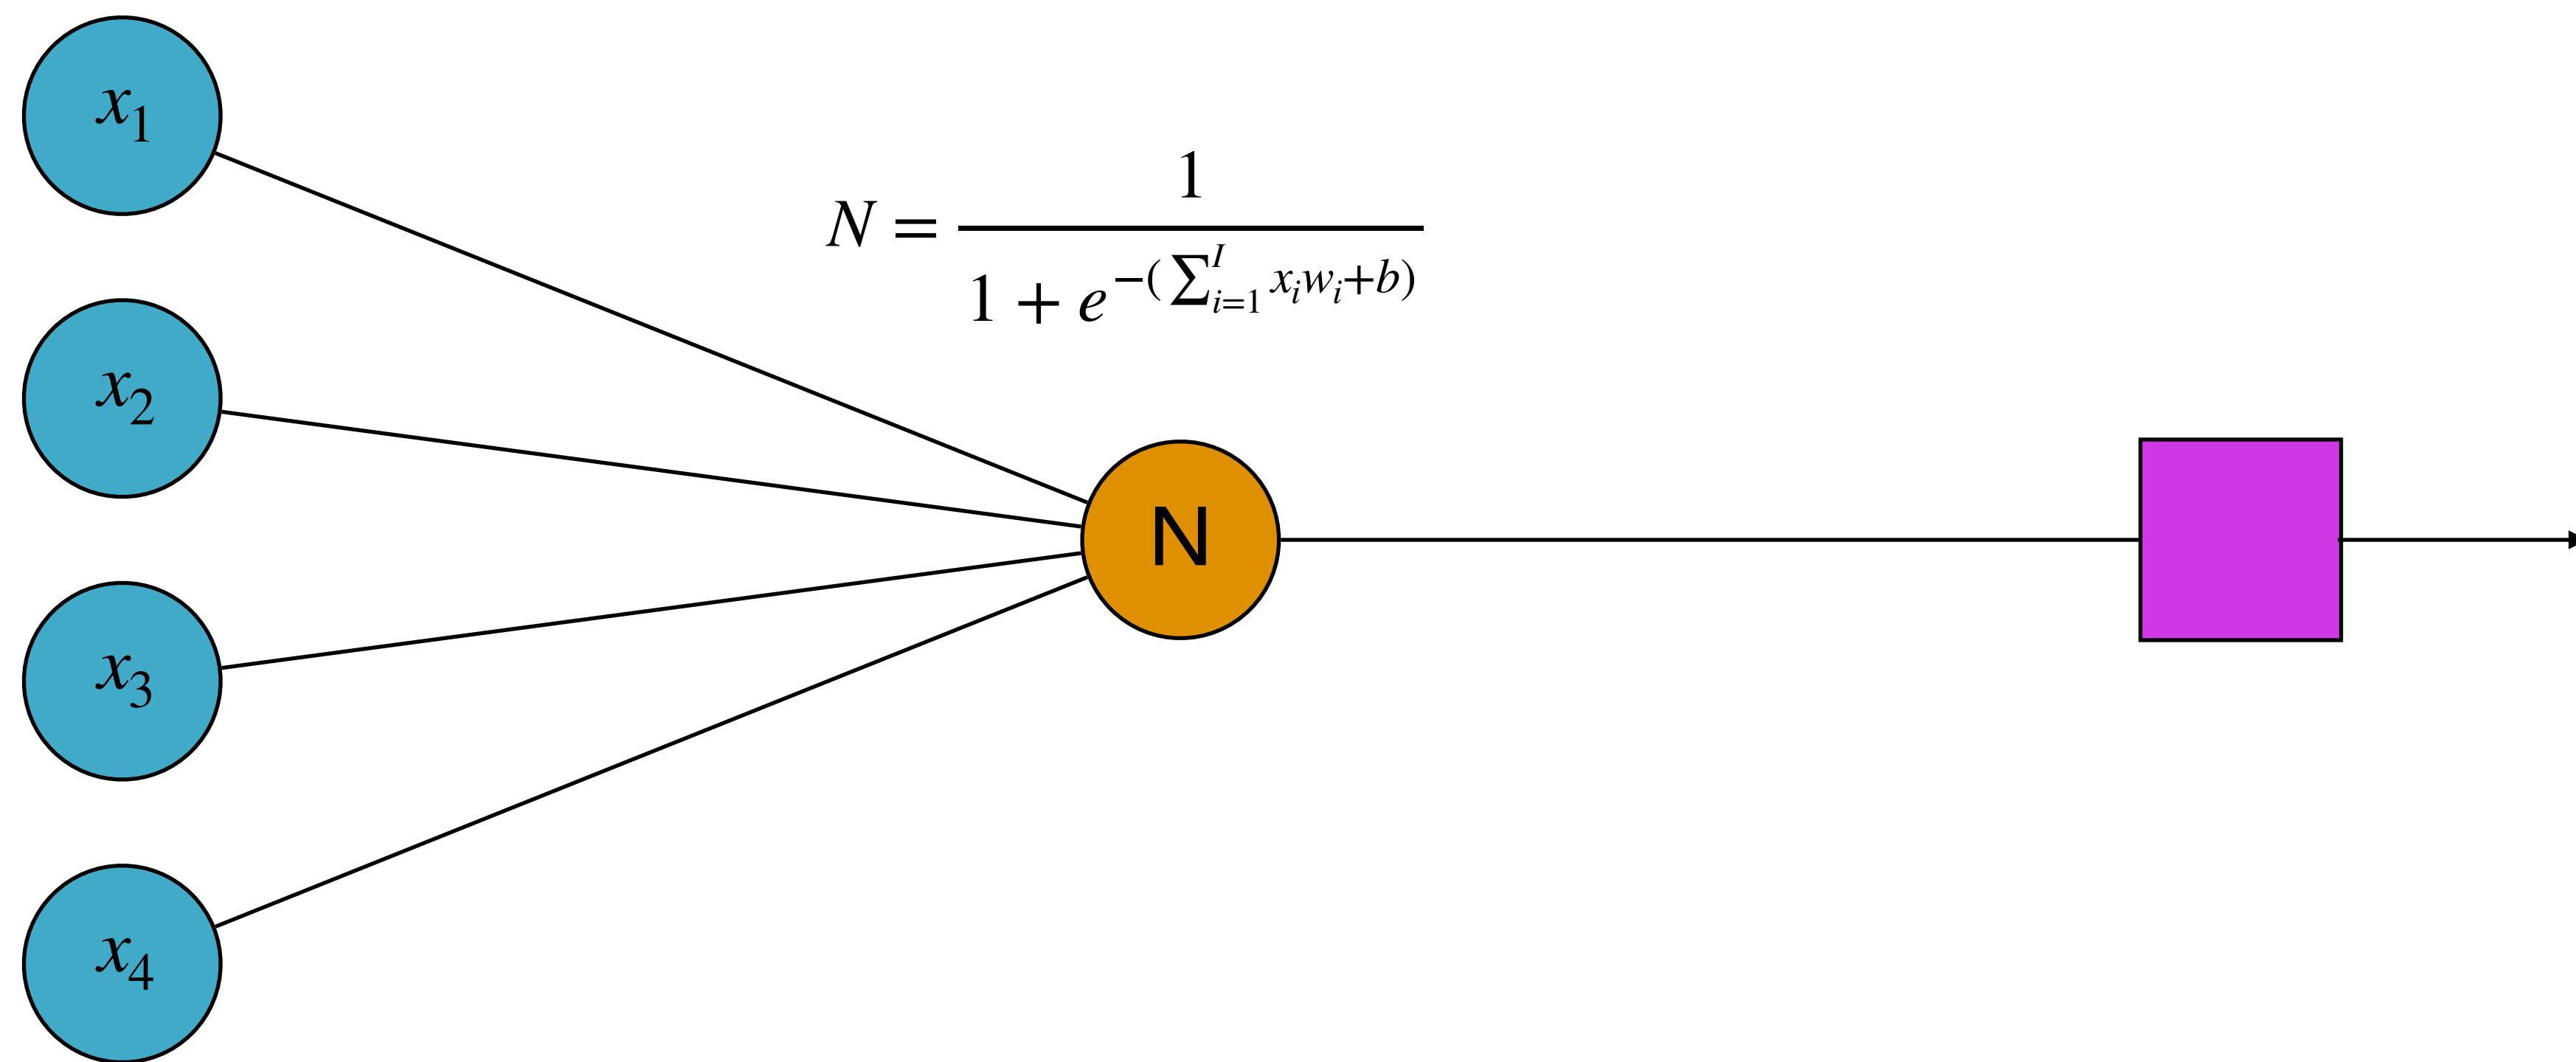
CLASSIFICATION



LOGISTIC REGRESSION



$$f = \frac{1}{1 + e^{-(x_1 w_1 + x_2 w_2 + b)}}$$

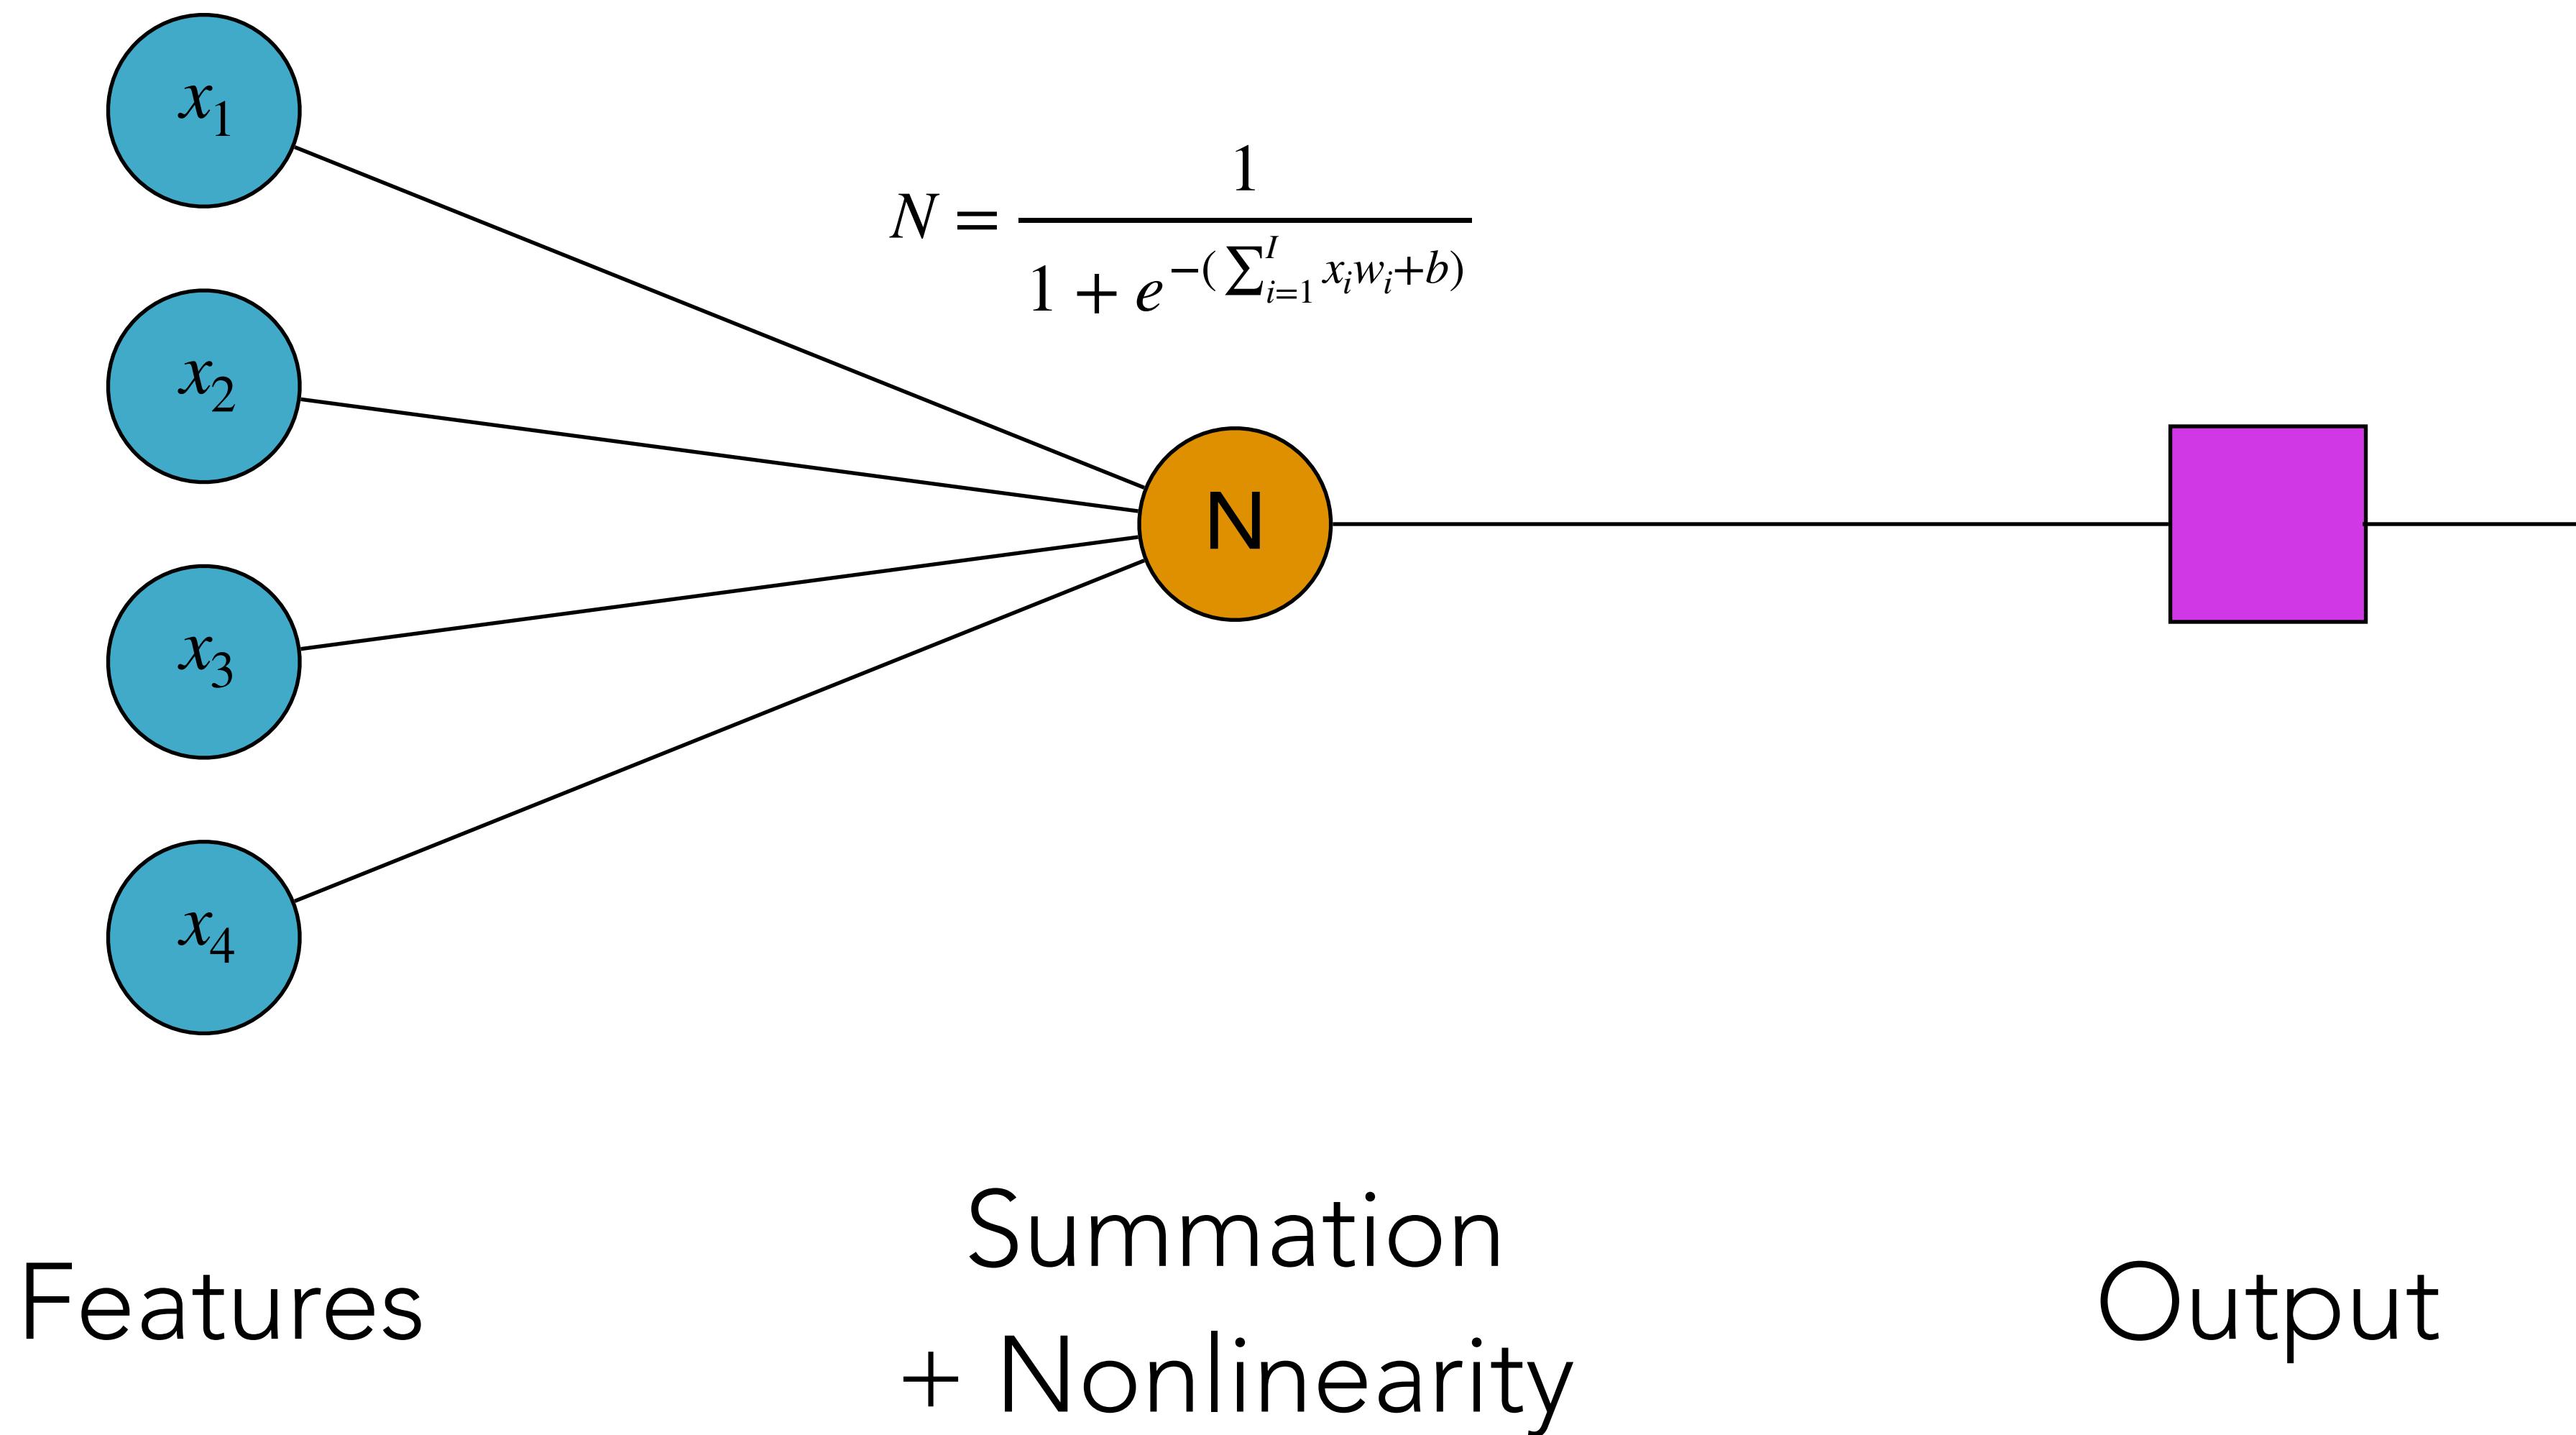


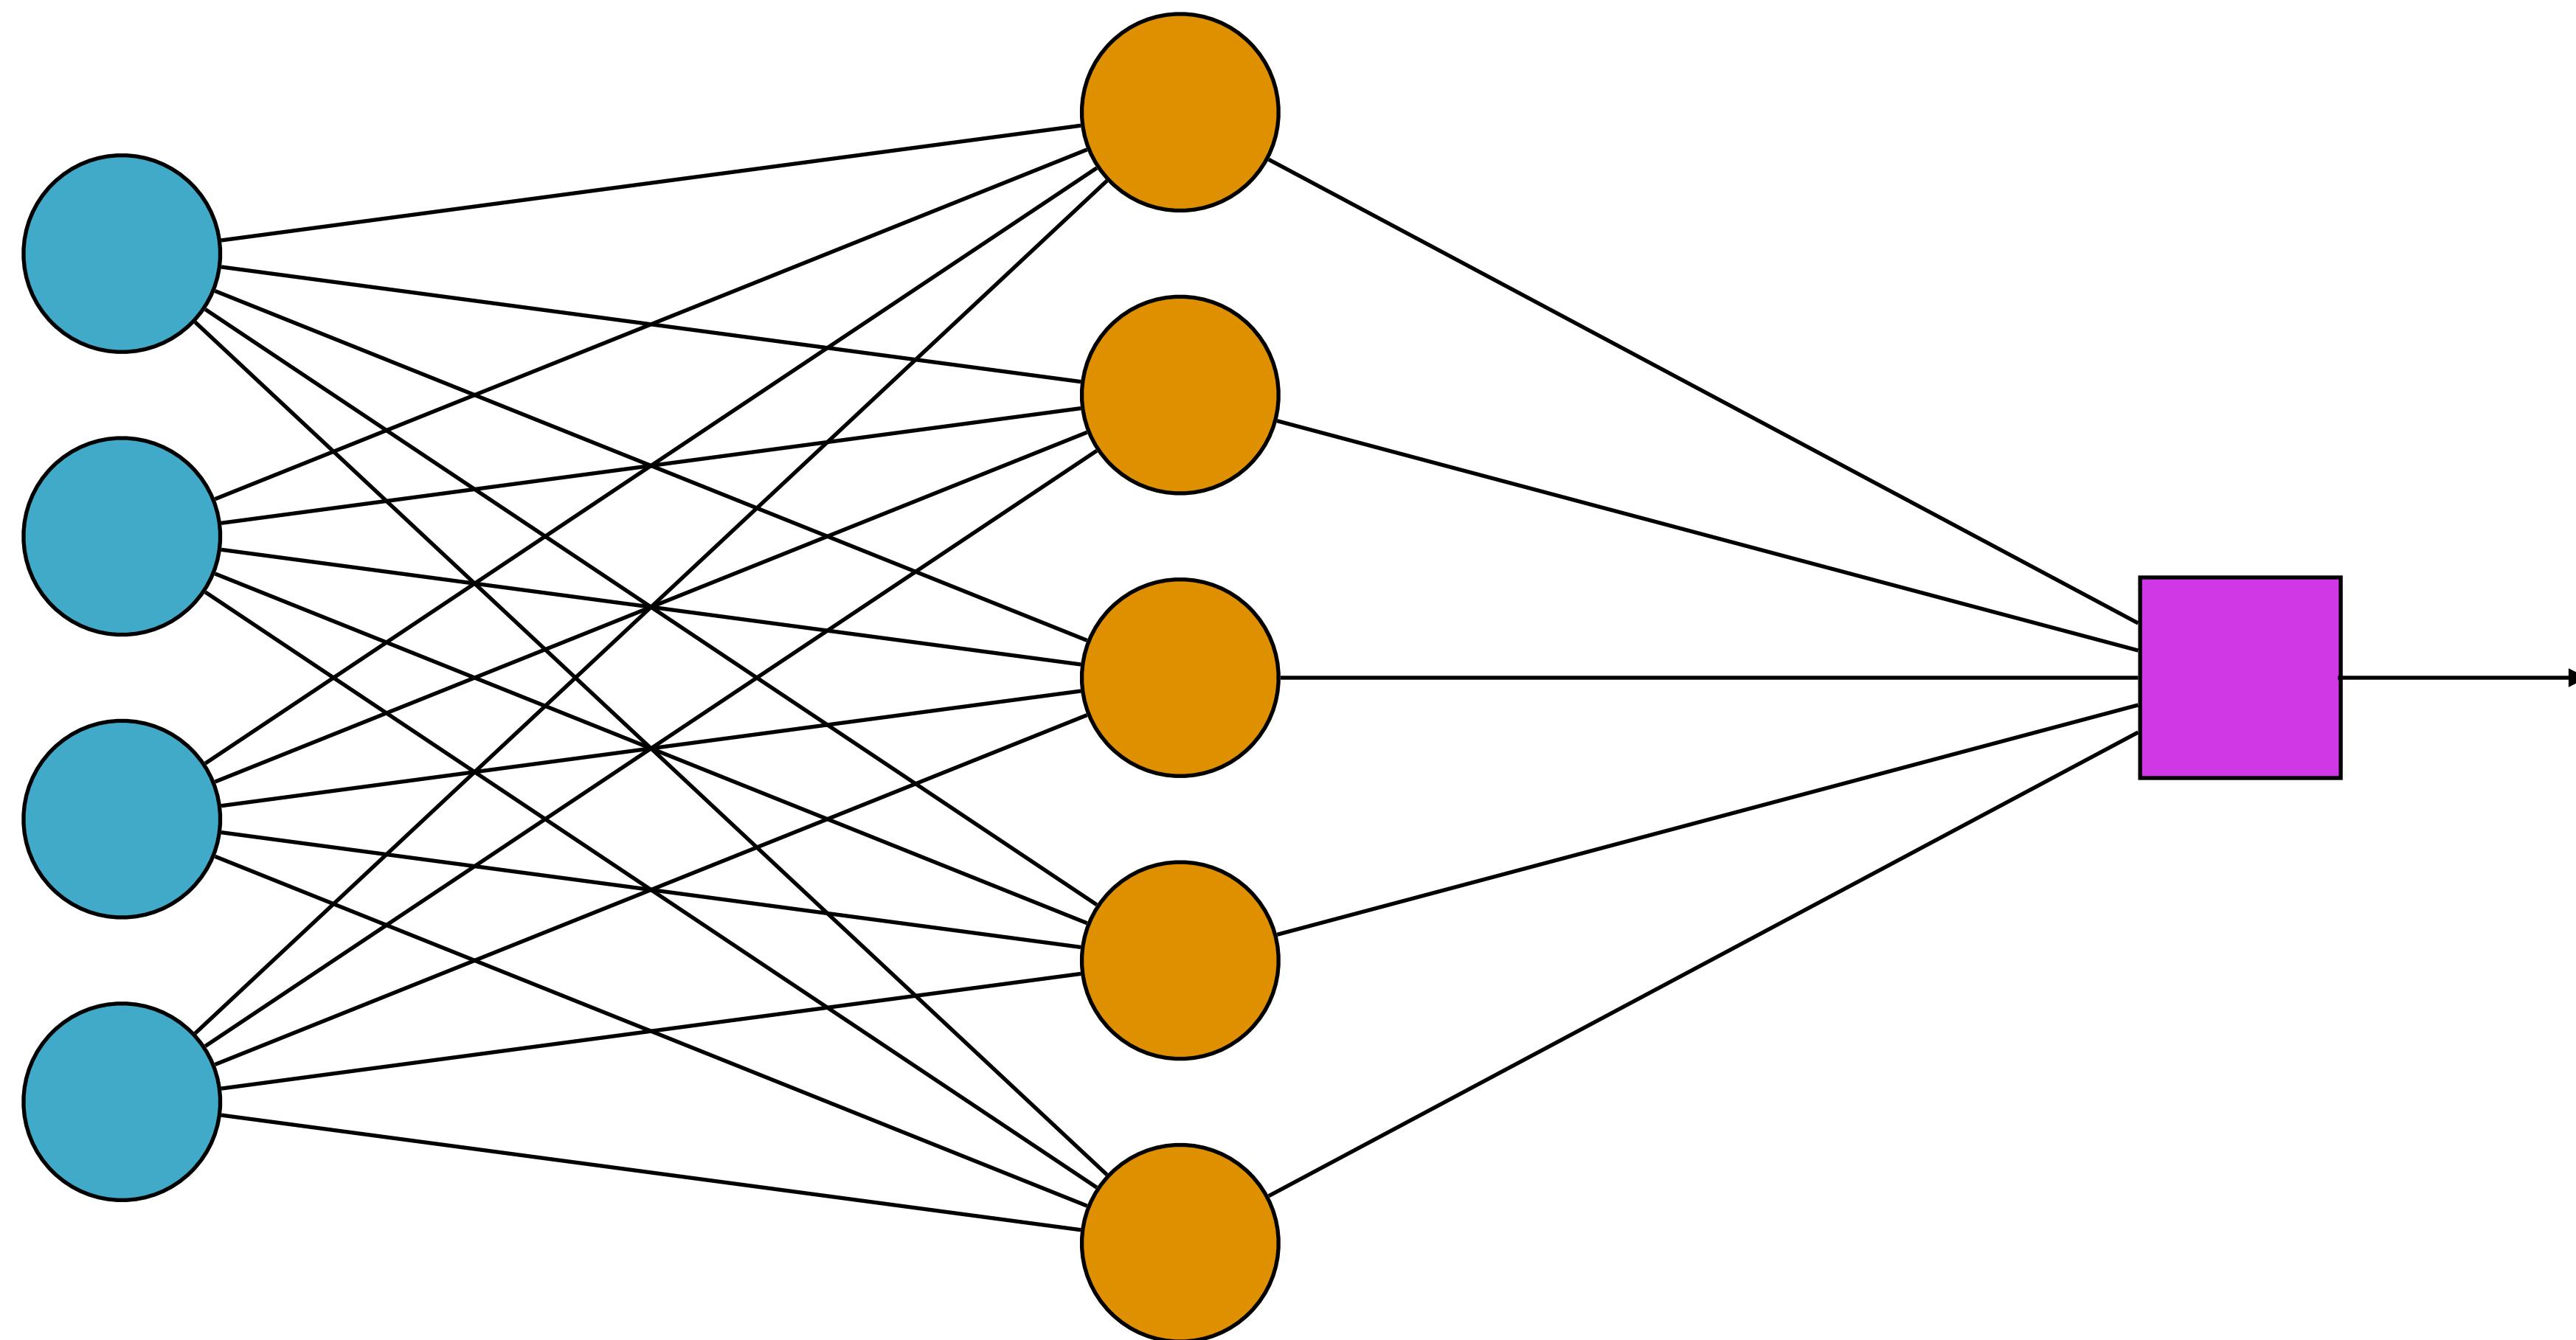
Features

Summation
+ Nonlinearity

Output

CHECK: HOW MANY TRAINABLE PARAMETERS?



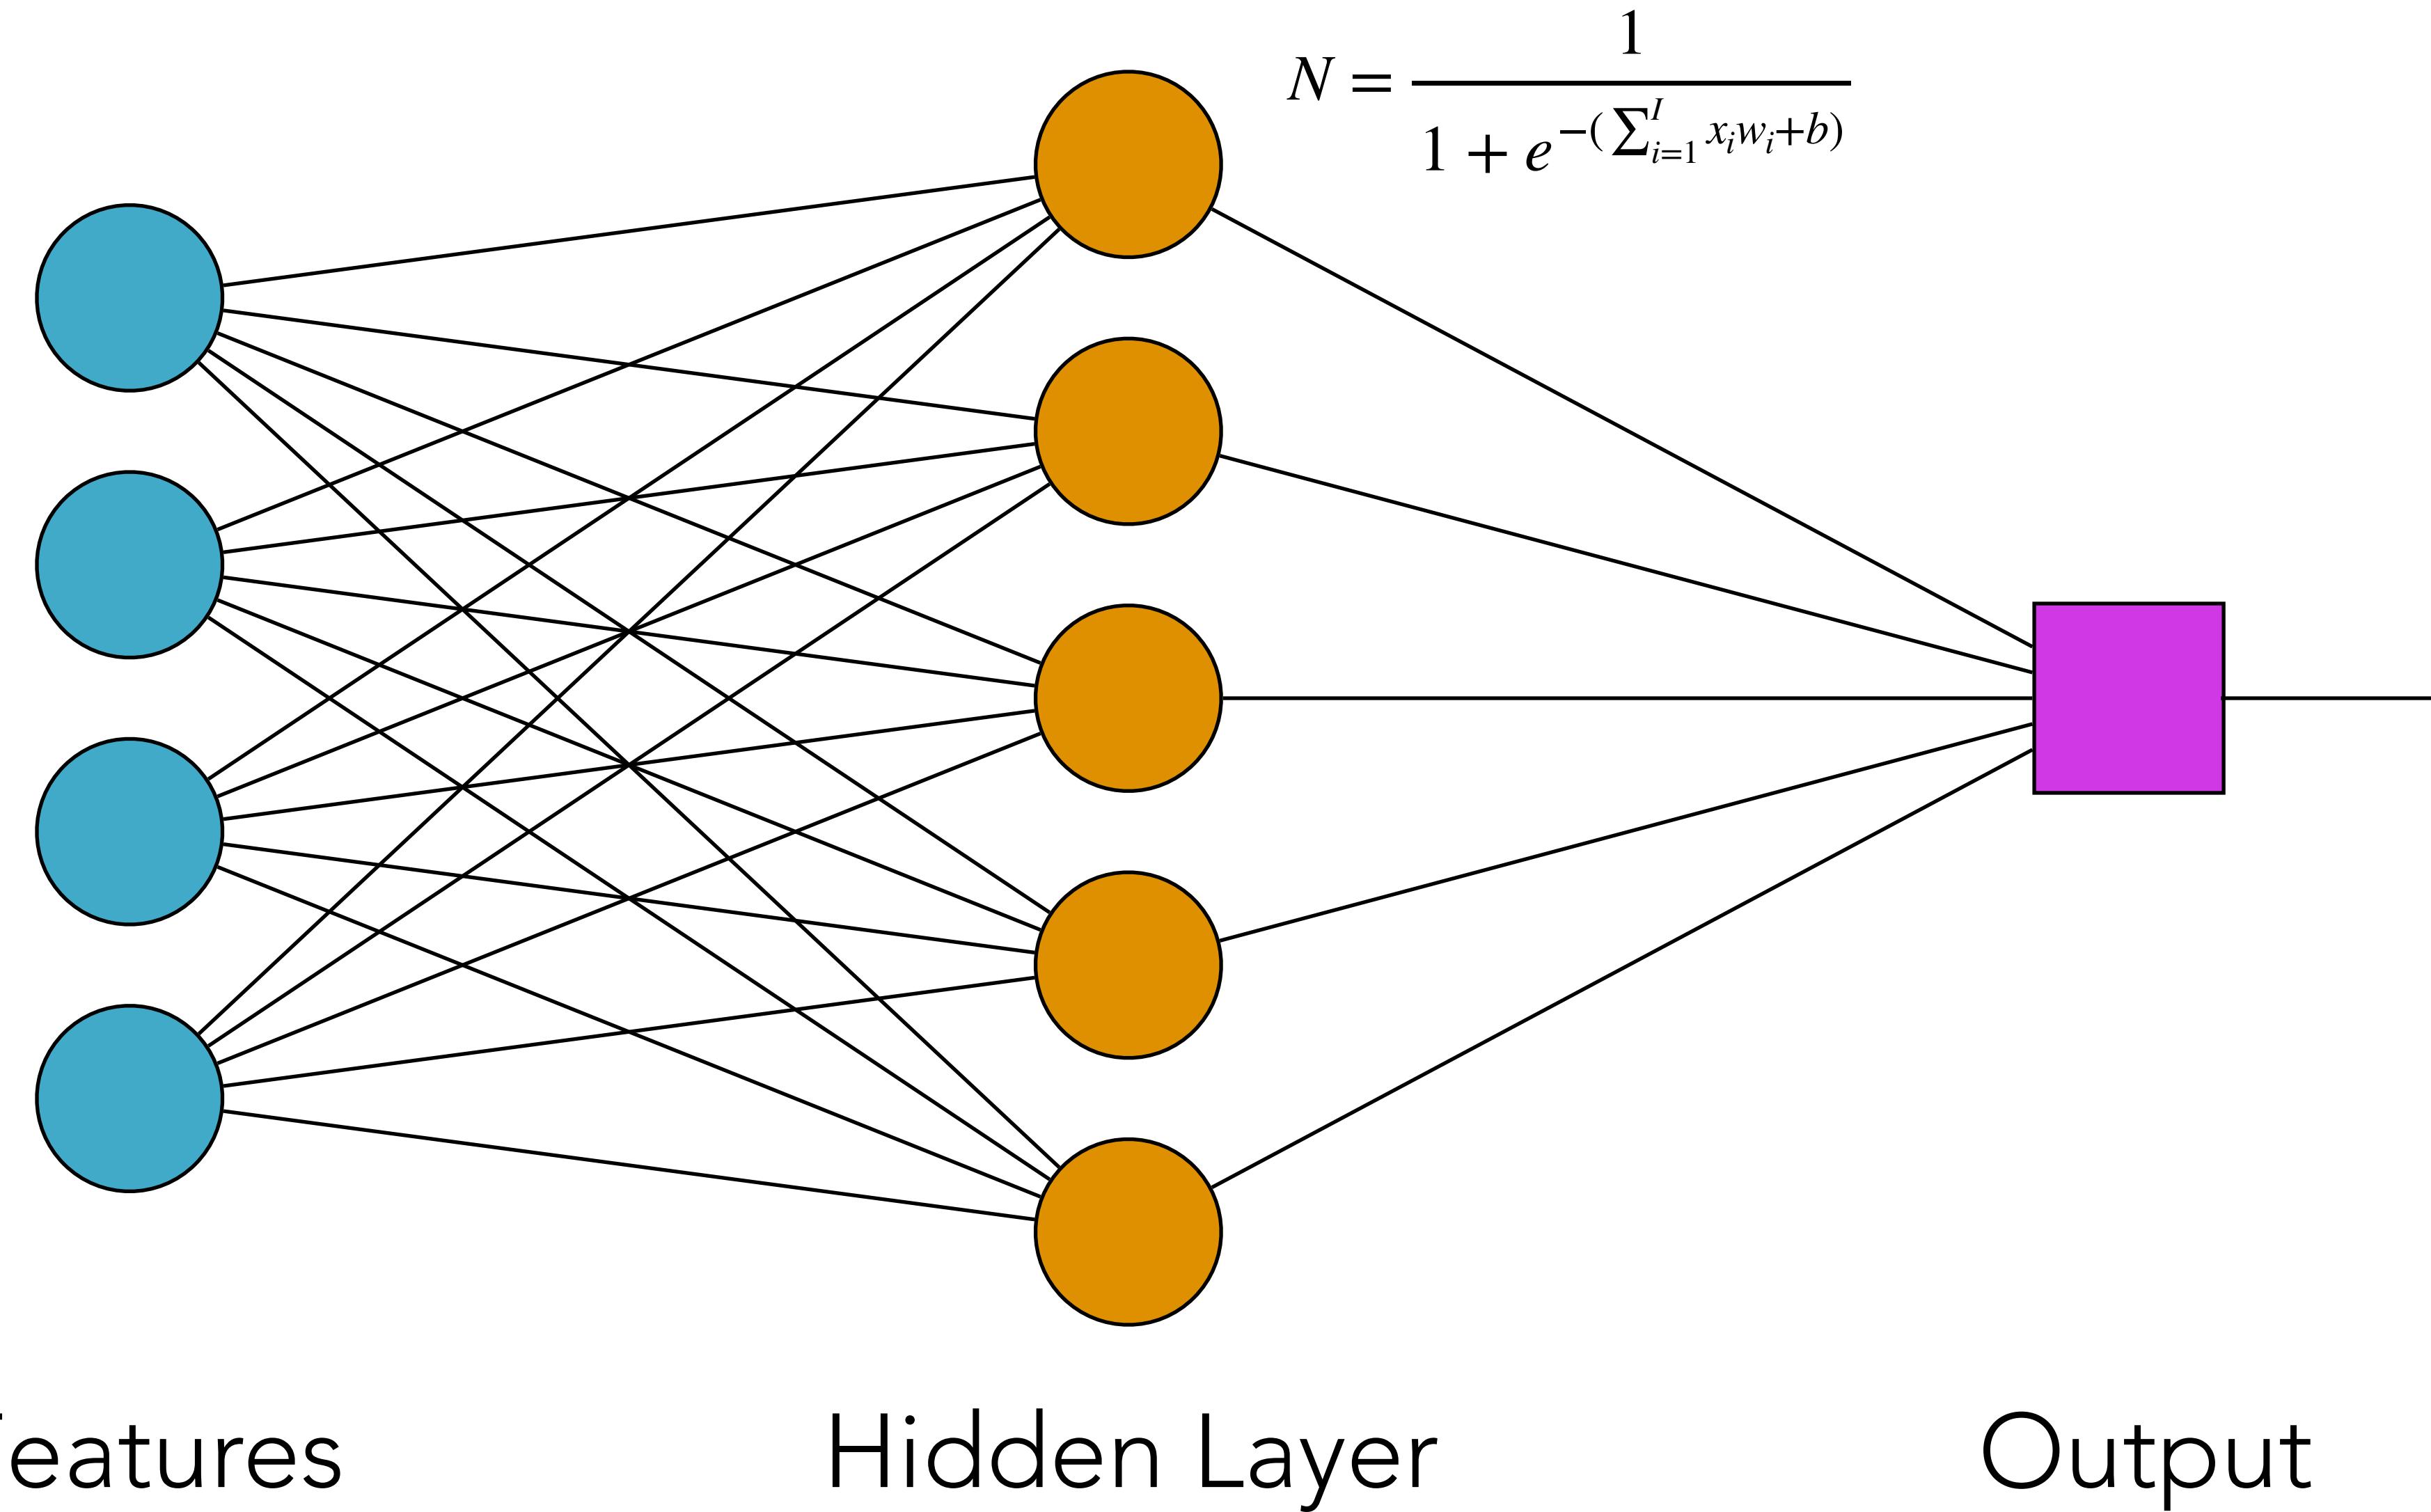


Features

Hidden Layer

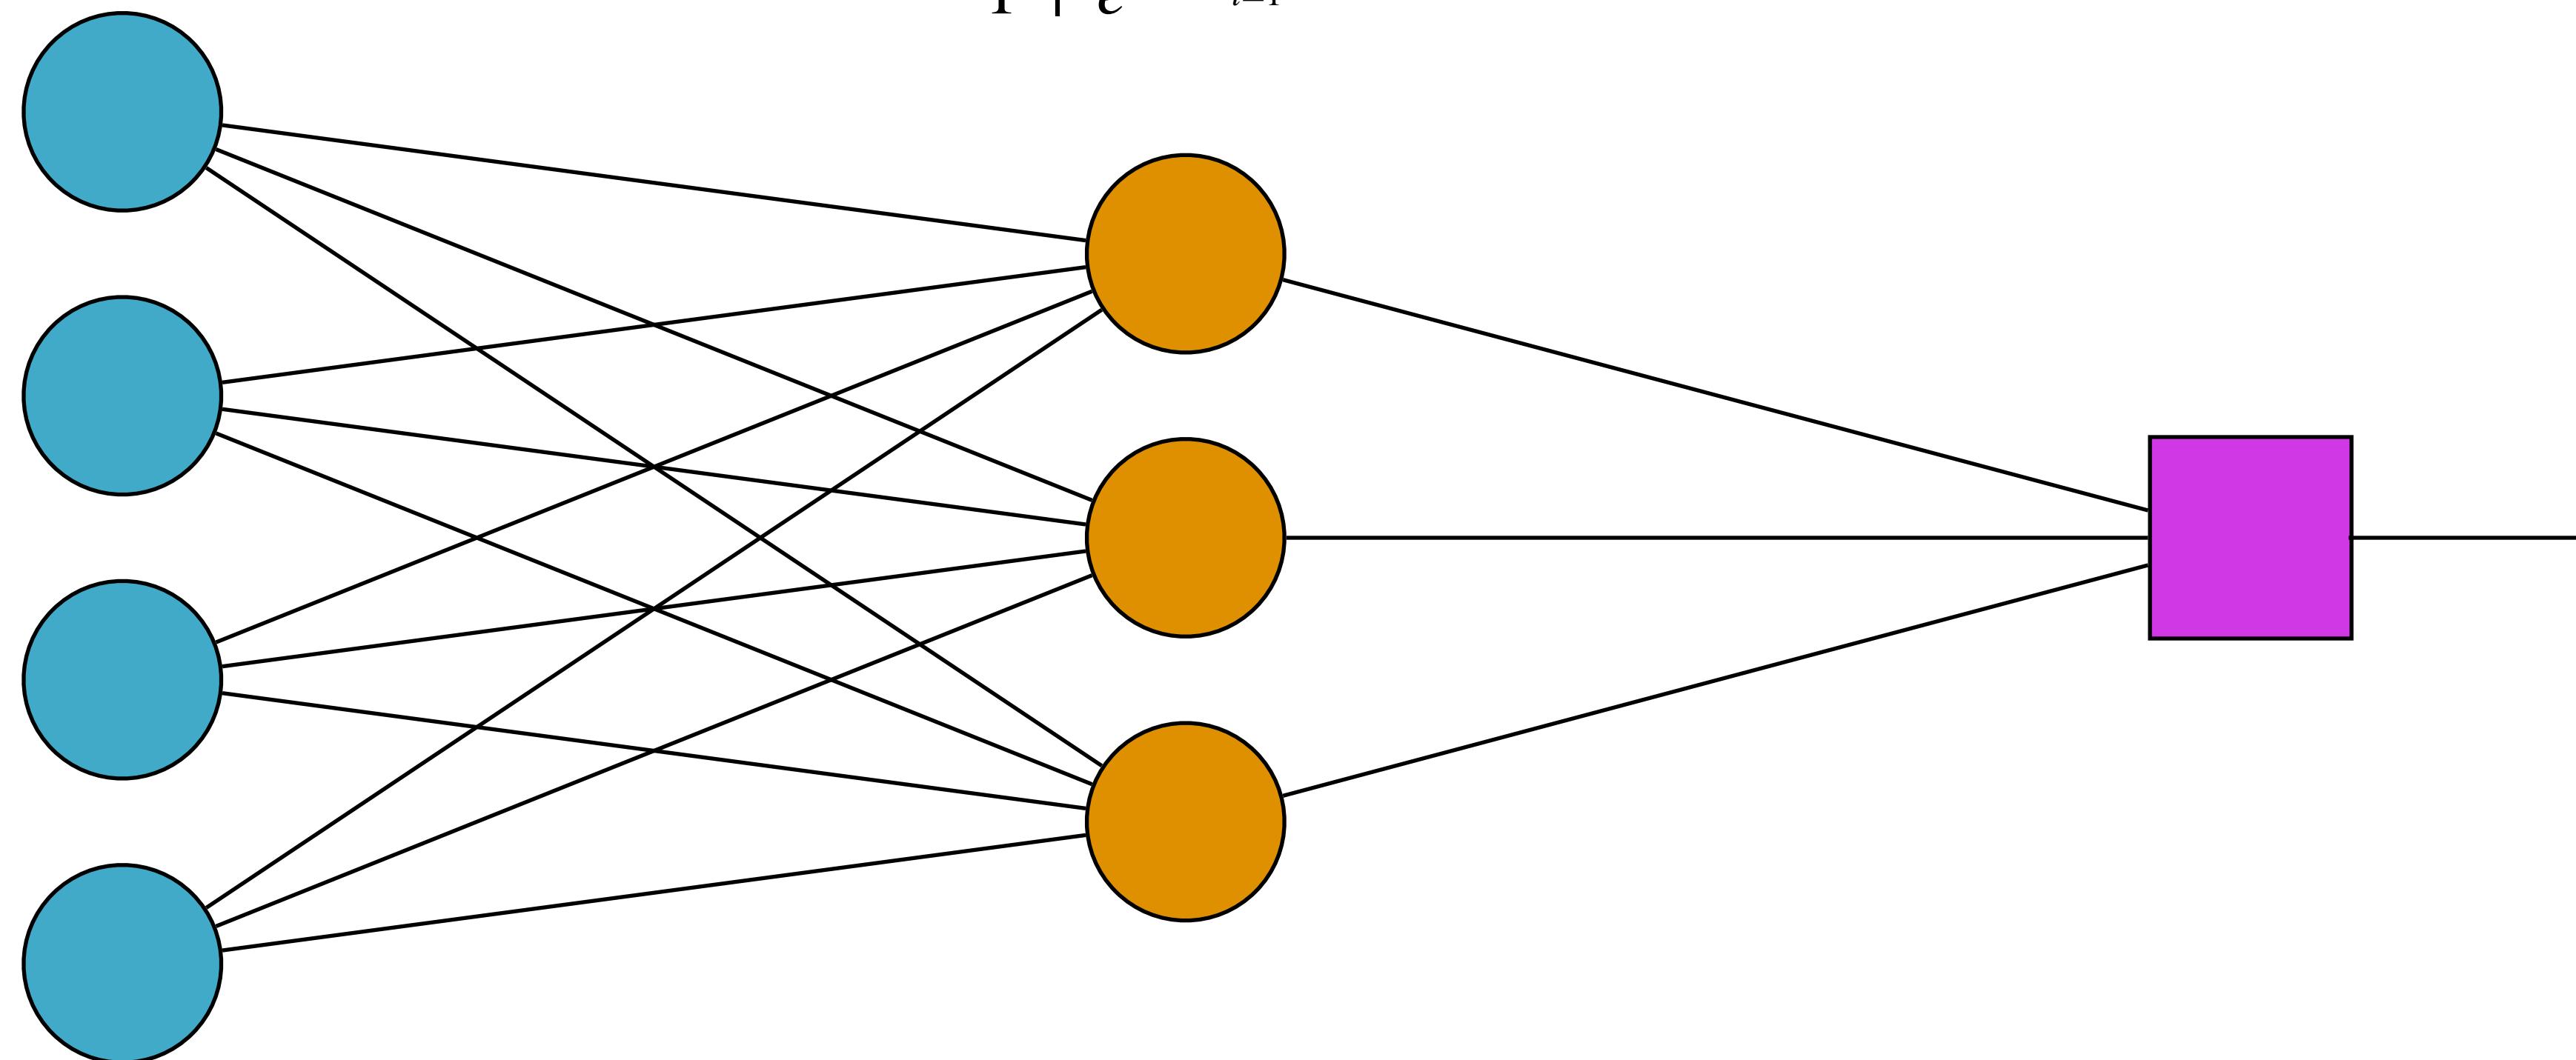
Output

CHECK: HOW MANY TRAINABLE PARAMETERS?



CHECK: HOW MANY TRAINABLE PARAMETERS?

$$N = \frac{1}{1 + e^{-(\sum_{i=1}^I x_i w_i + b)}}$$



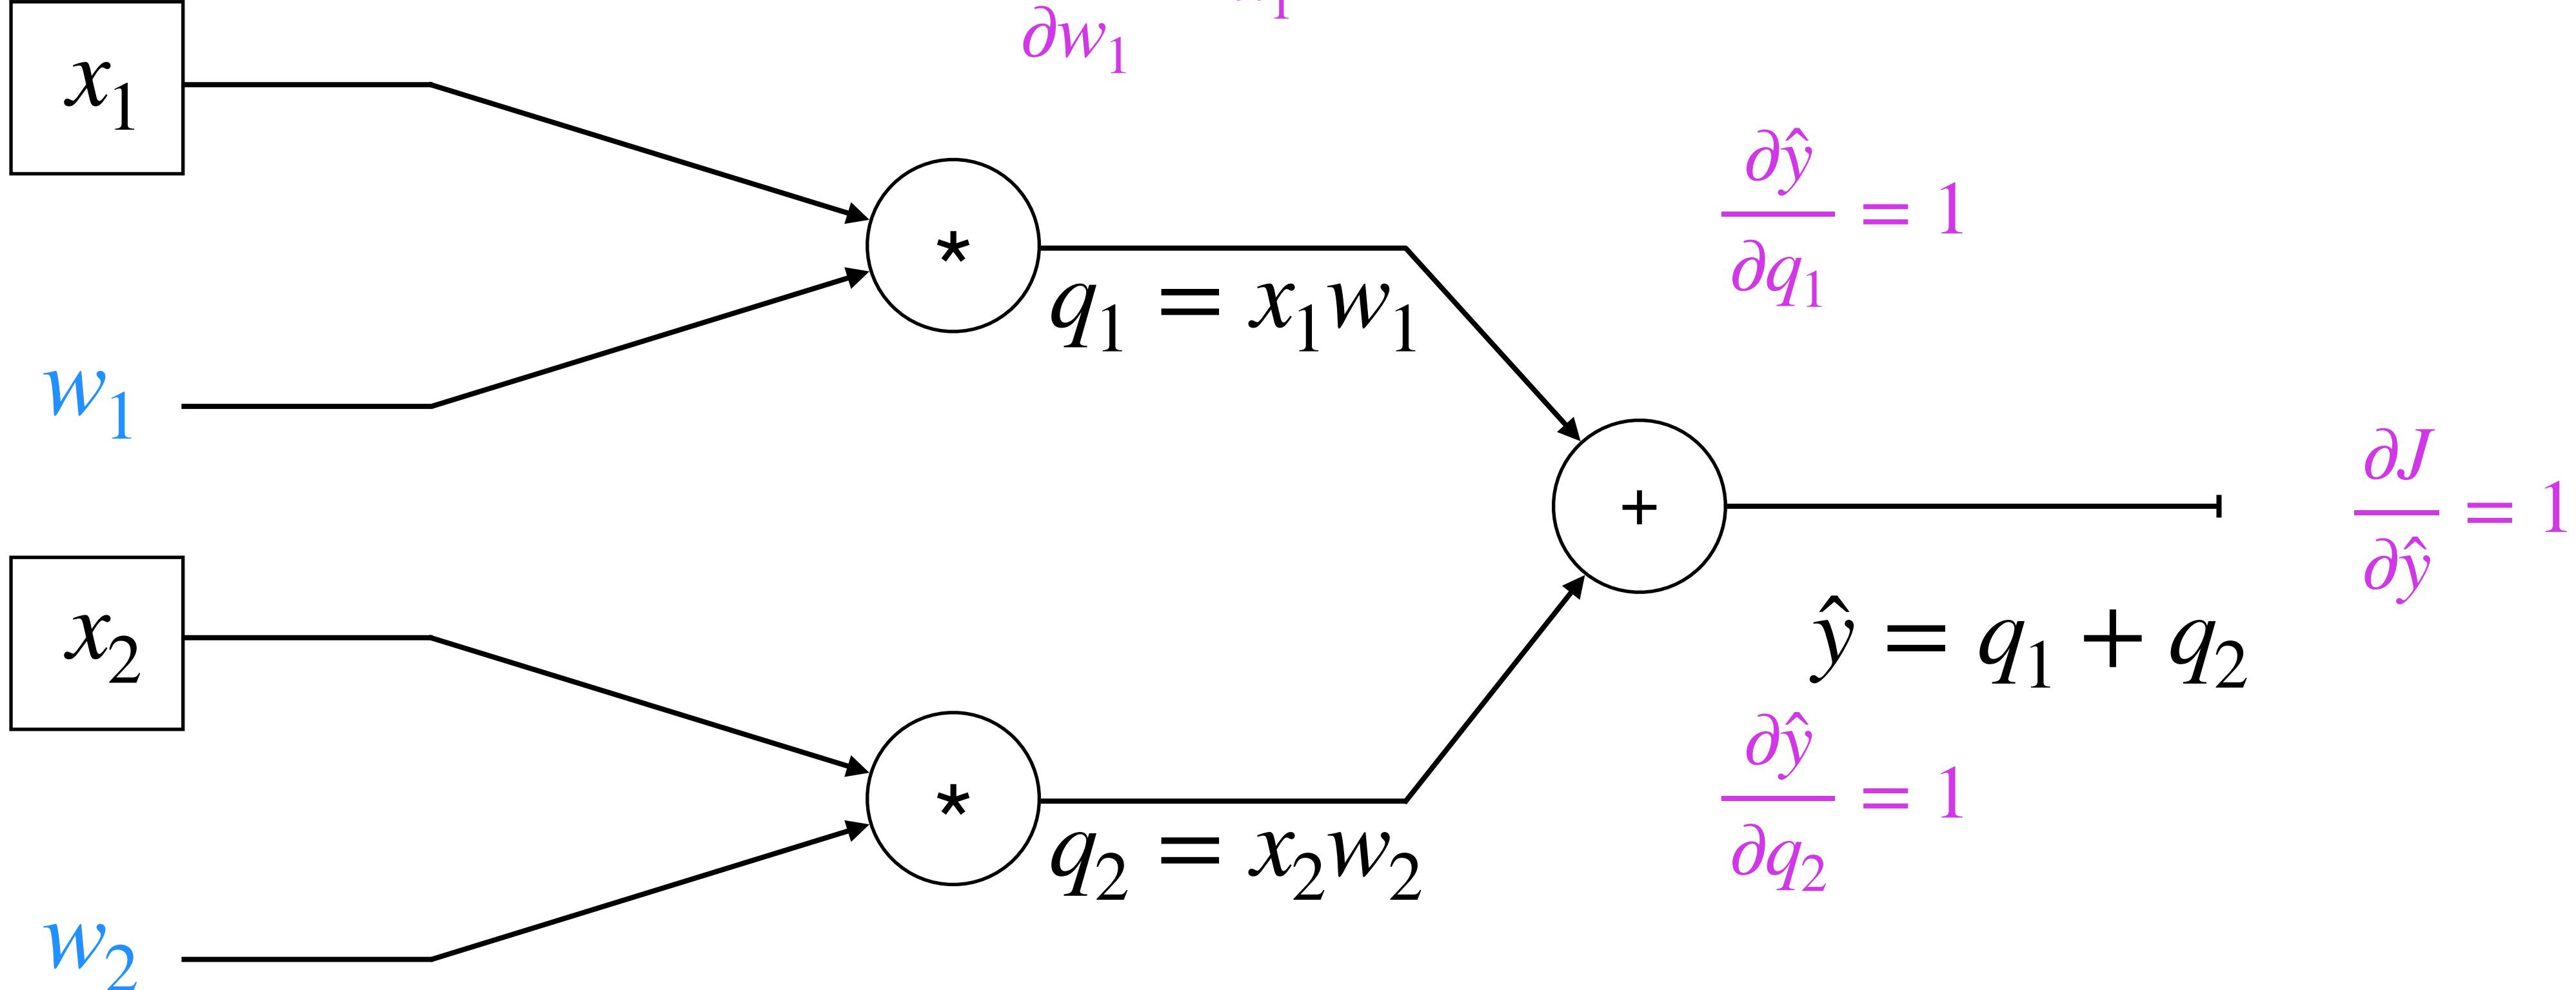
Features

Hidden Layer

Output

BACKPROPAGATION

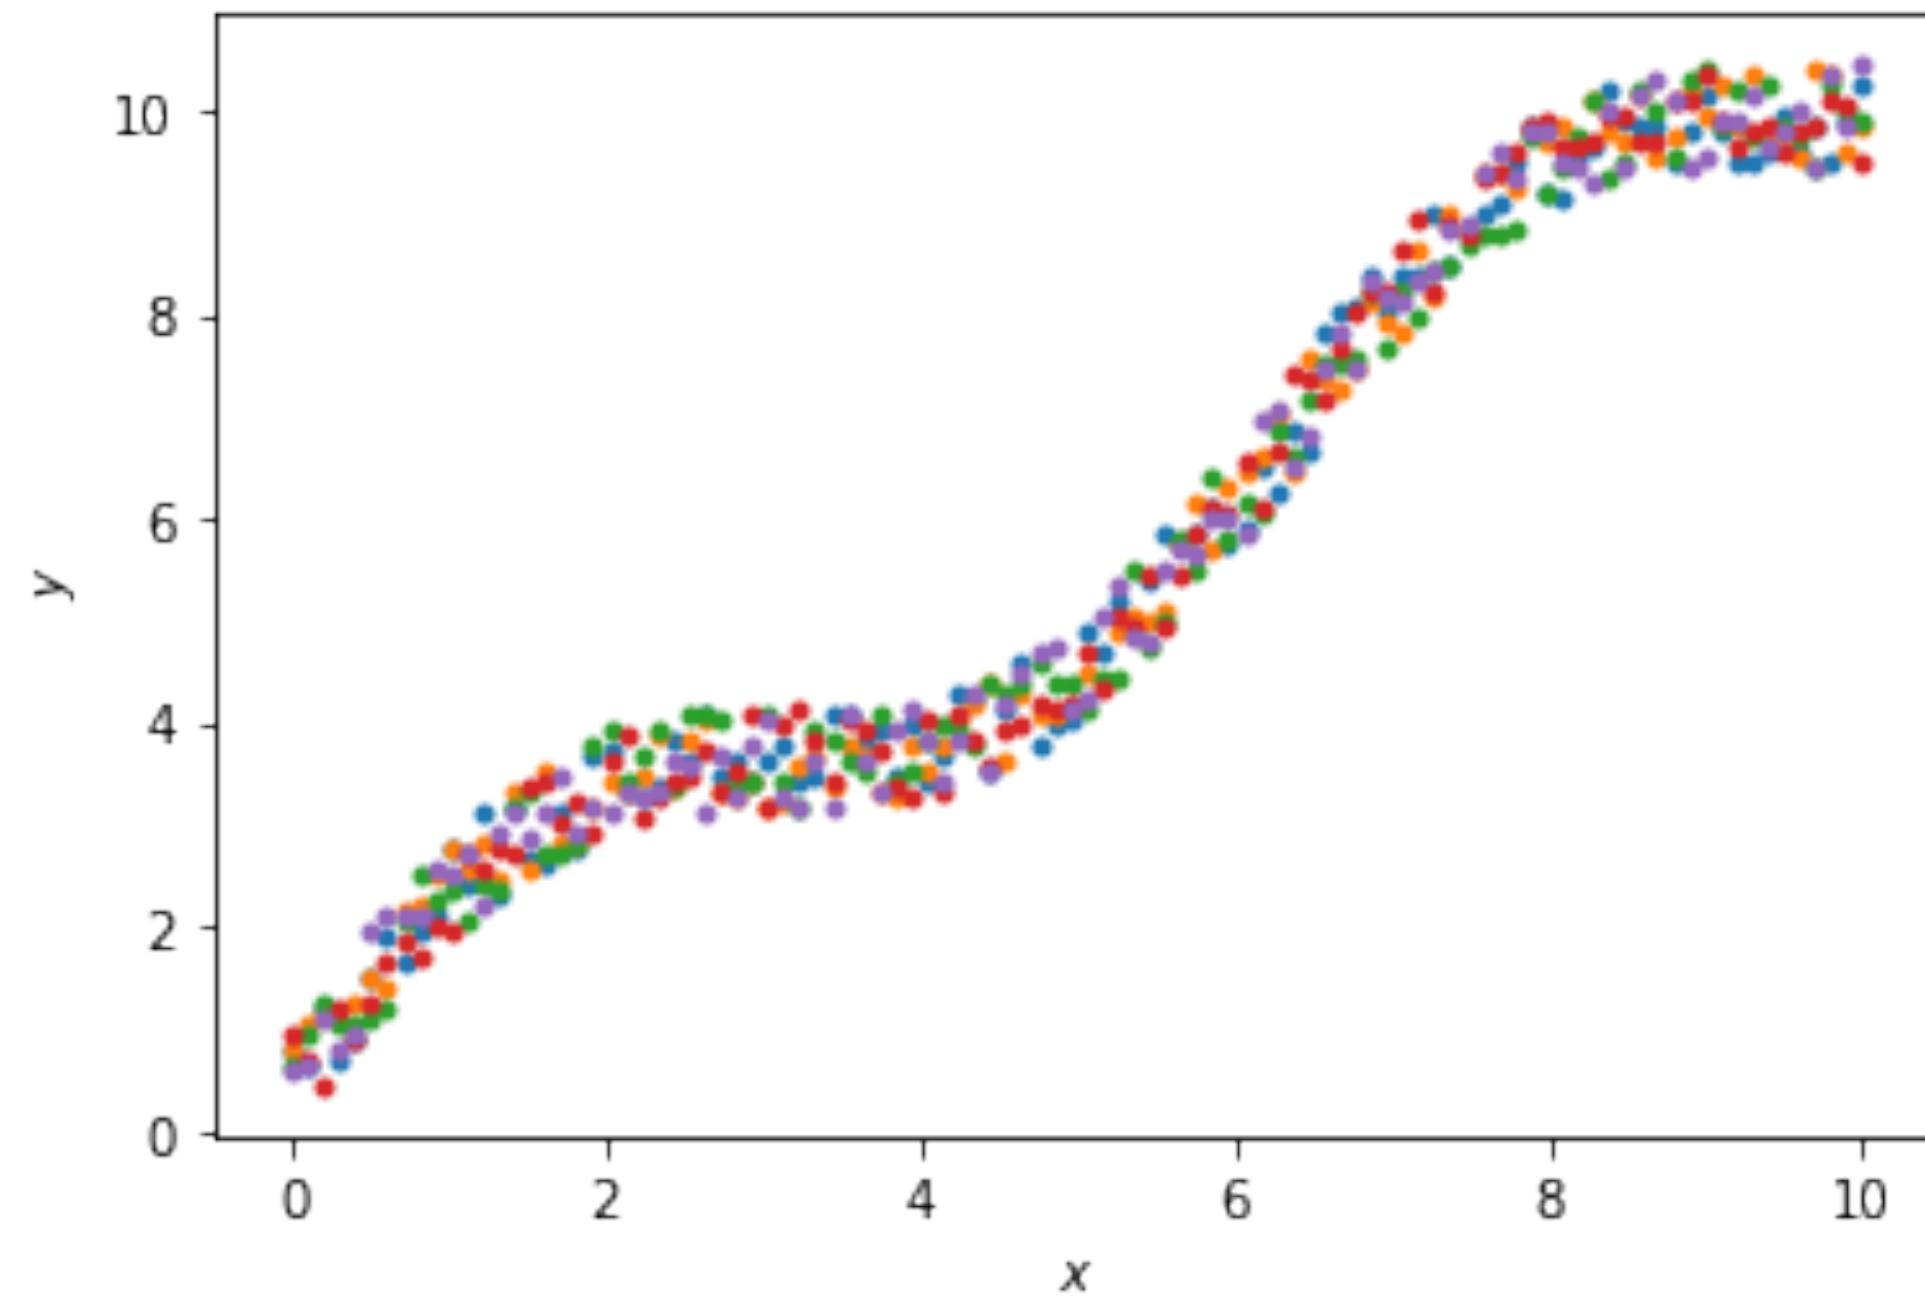
$$w_1 = w_1 - \eta * \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial q_1} \frac{\partial q_1}{\partial w_1}$$



$$w_2 = w_2 - \eta * \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial q_2} \frac{\partial q_2}{\partial w_2}$$

Loss function
 $J(w) = \hat{y} - y$

LOSS FUNCTIONS



Loss function

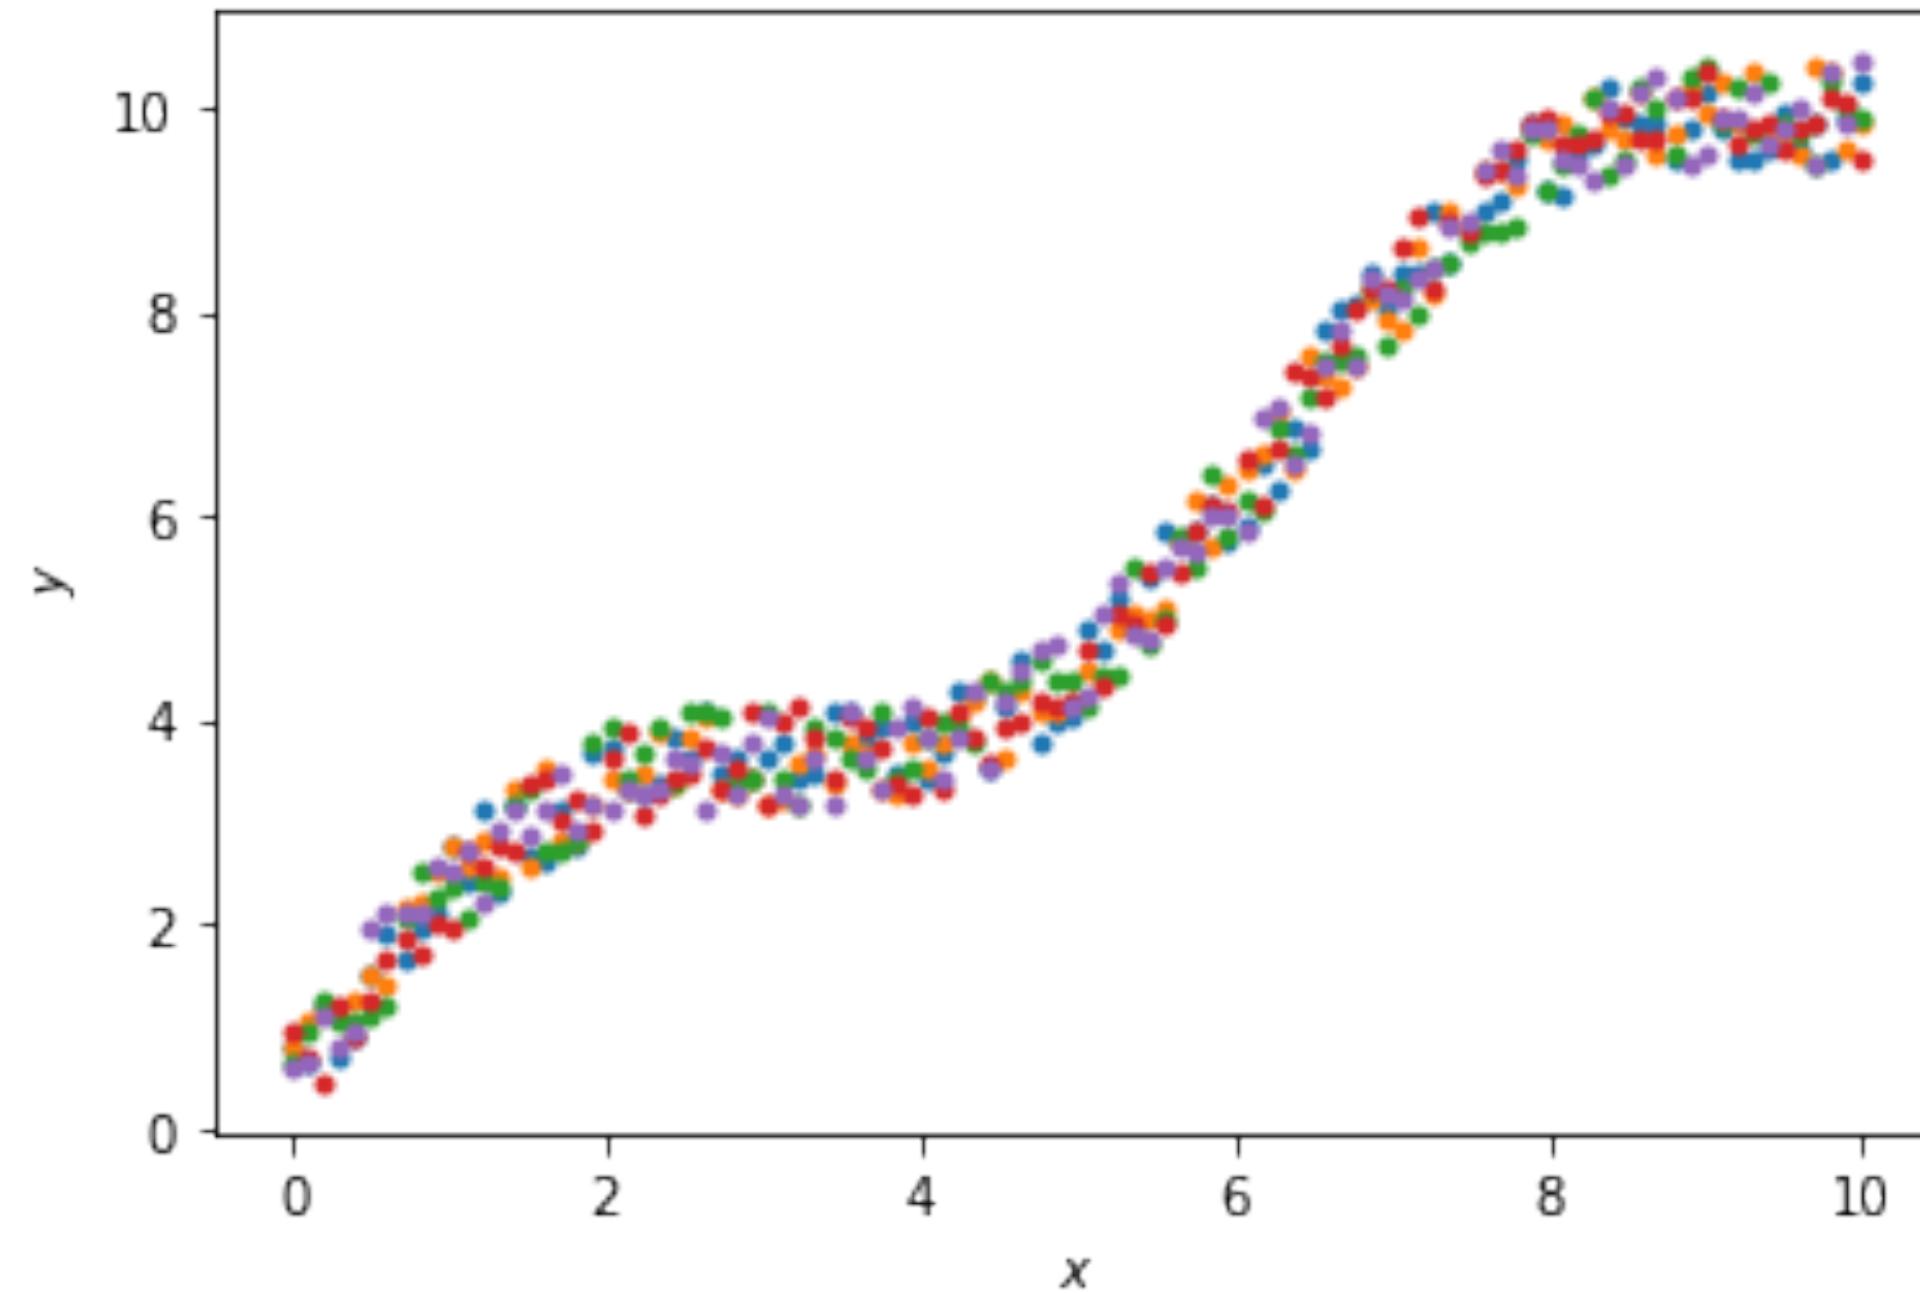
Mean squared error
(MSE)

Mean absolute error
(MAE)

$$J(w) = \frac{1}{N} \sum_{i=0}^N (\hat{y}_i - y_i)^2$$

$$J(w) = \frac{1}{N} \sum_{i=0}^N |\hat{y}_i - y_i|$$

LOSS FUNCTIONS



Loss function

Mean squared error

mean

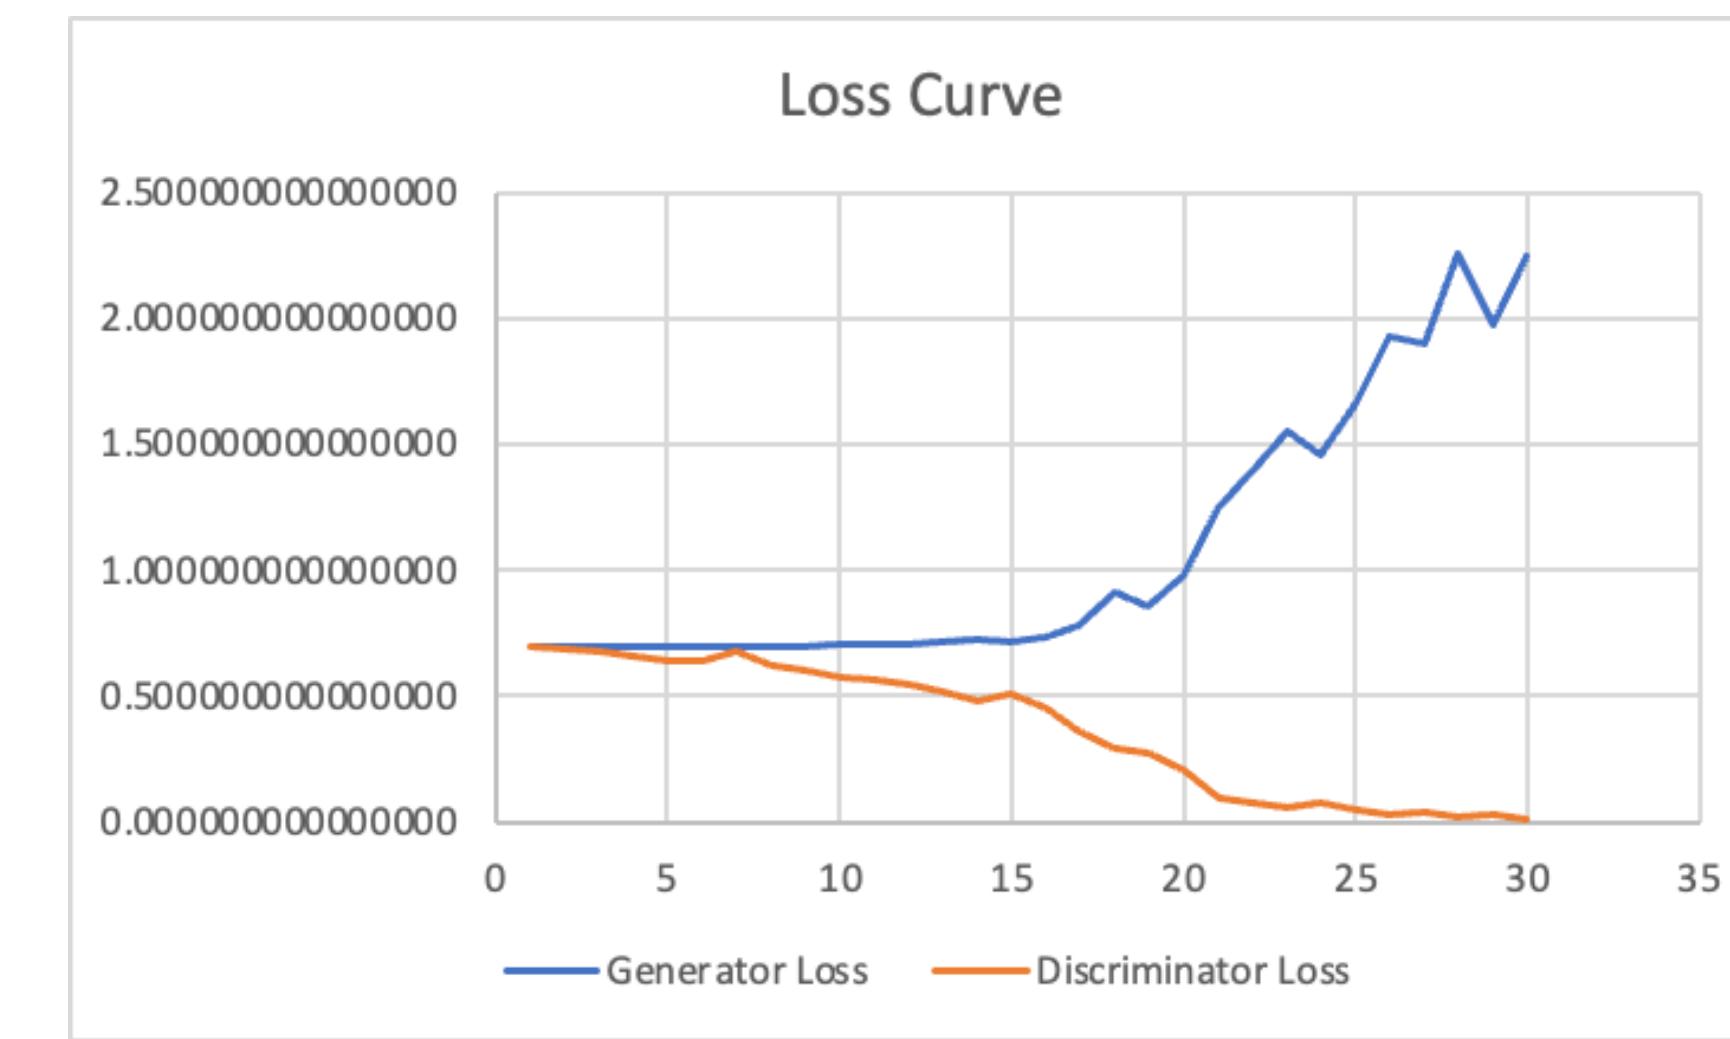
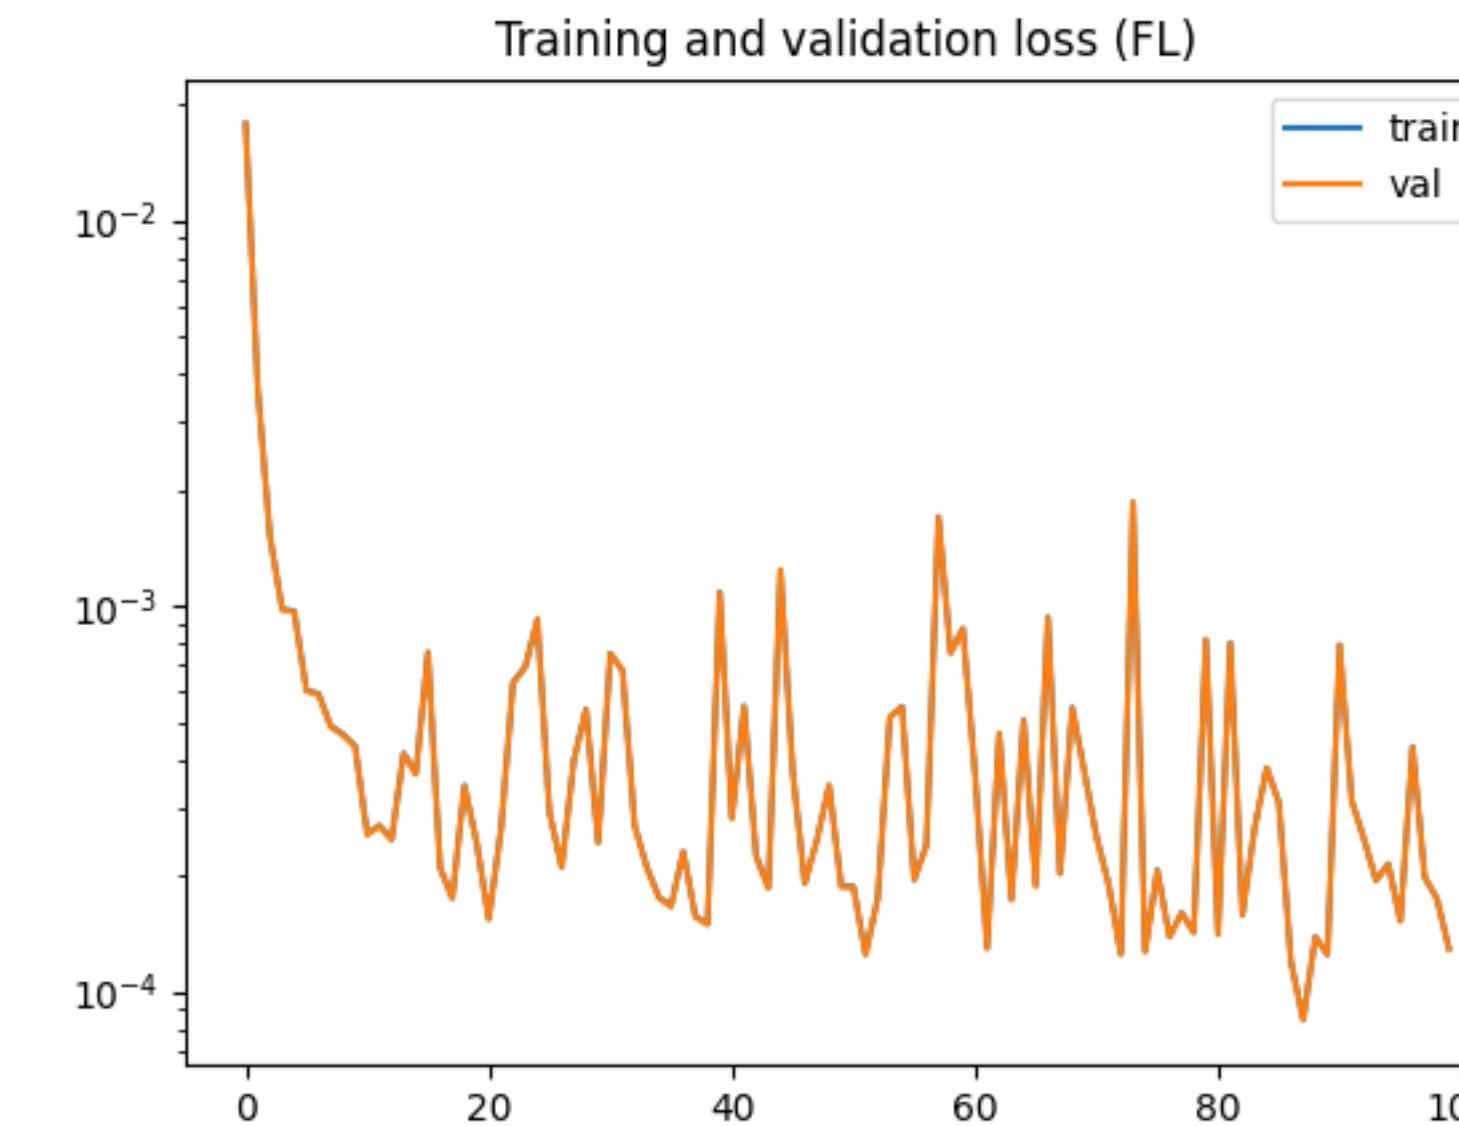
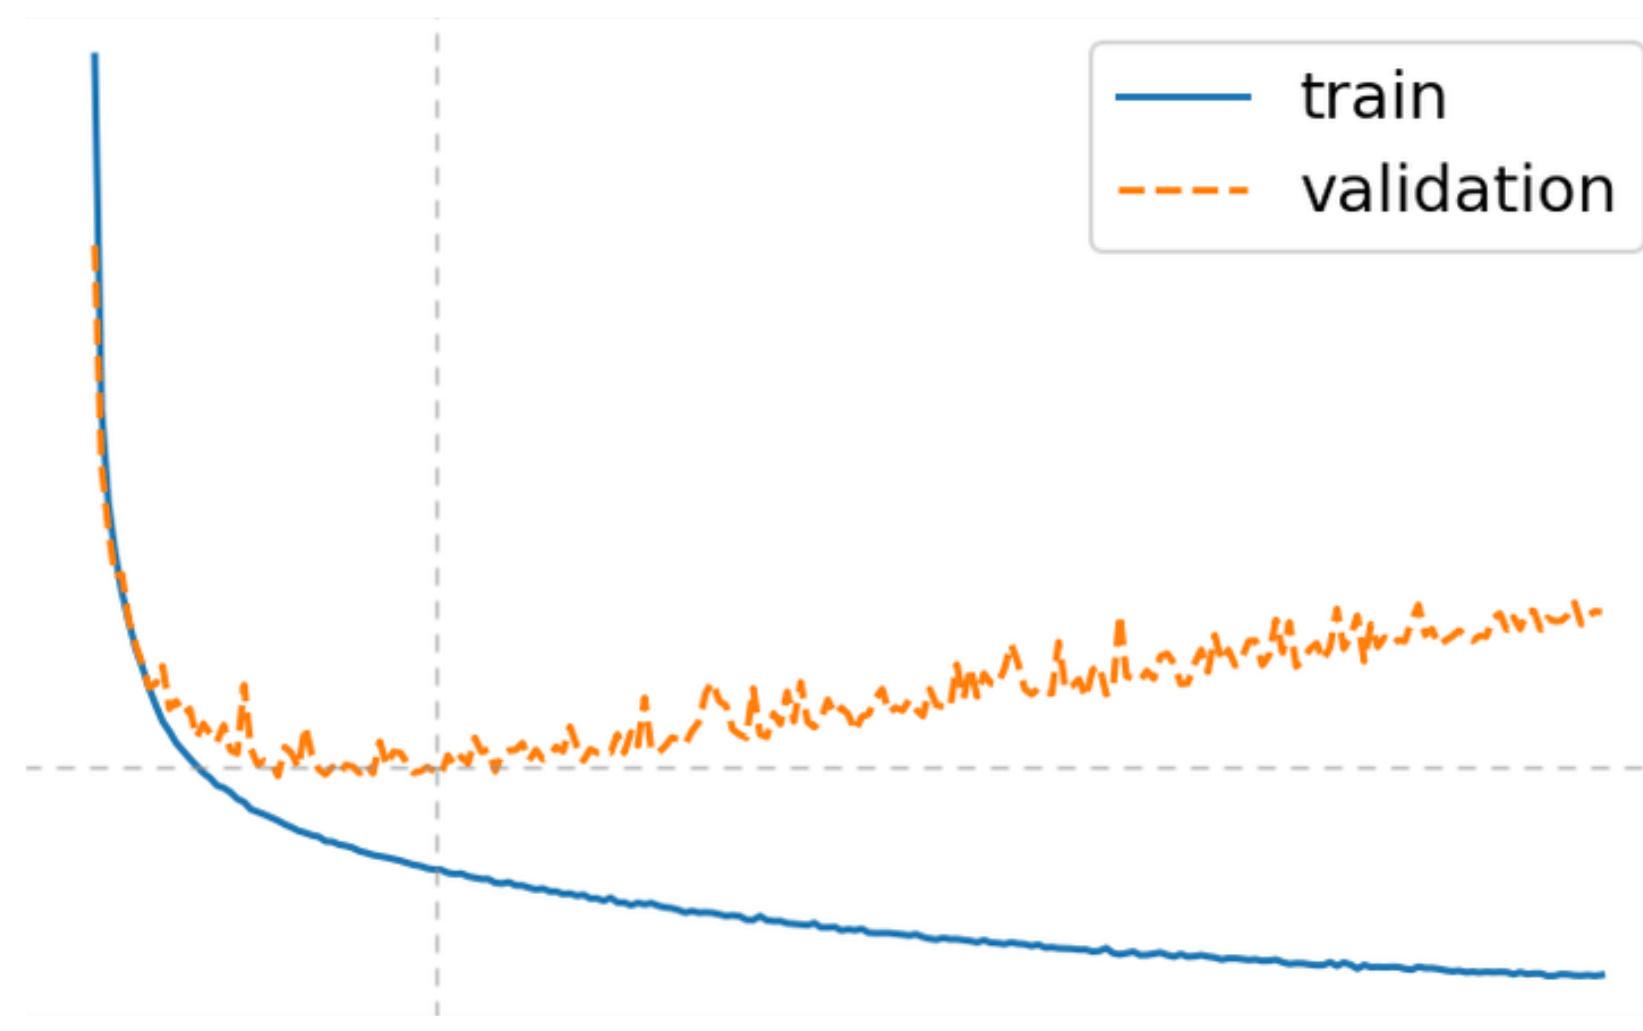
Mean absolute error

median

$$J(w) = \frac{1}{N} \sum_{i=0}^N (\hat{y}_i - y_i)^2$$

$$J(w) = \frac{1}{N} \sum_{i=0}^N |\hat{y}_i - y_i|$$

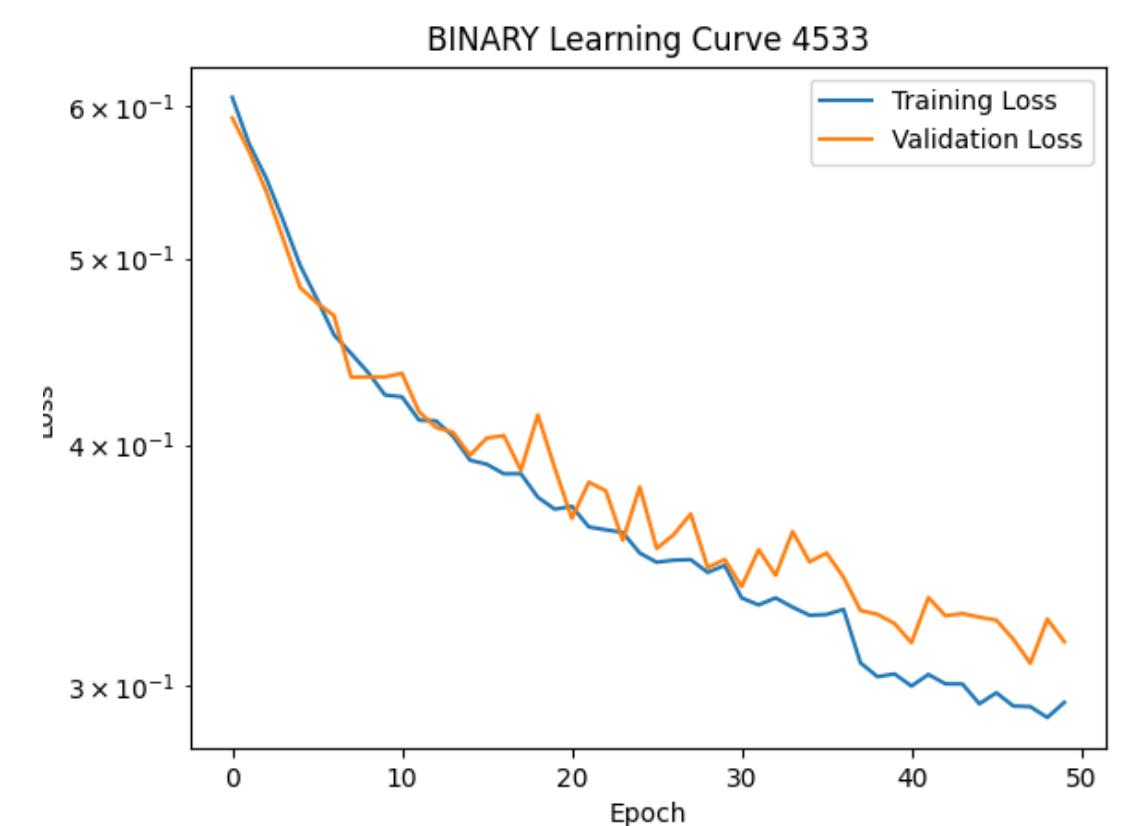
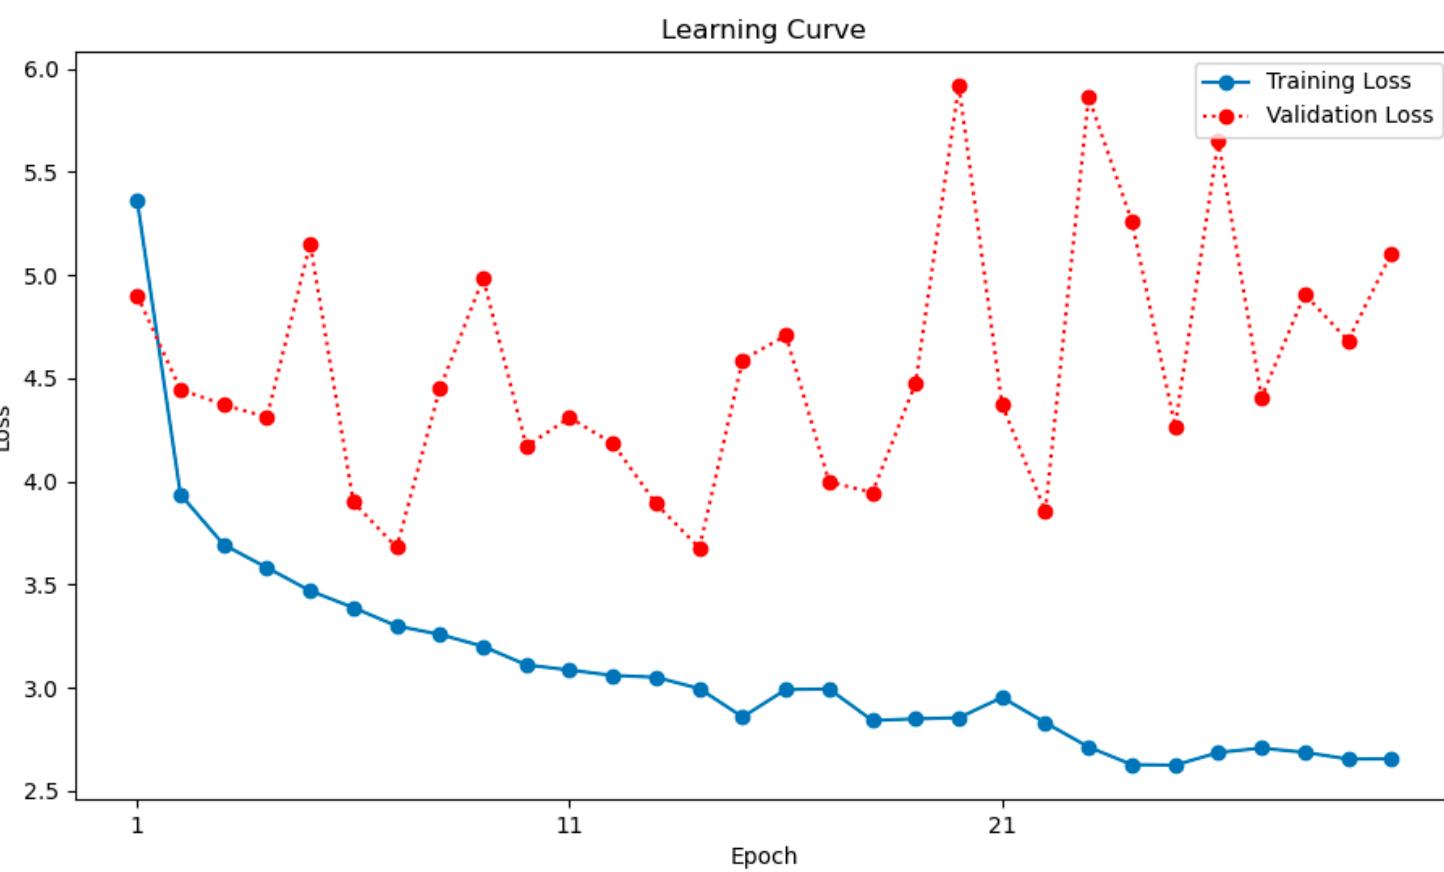
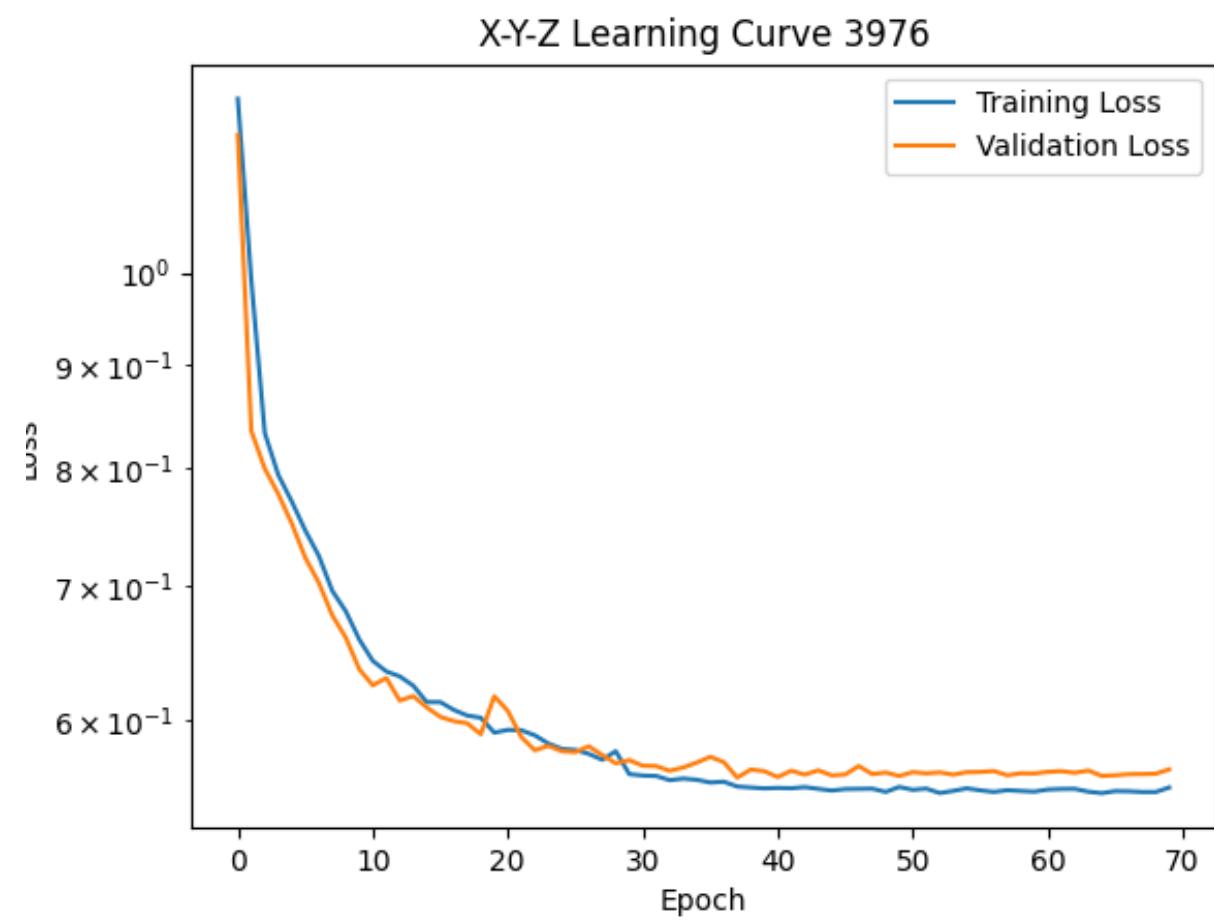
Learning (loss) curves



TRAINING

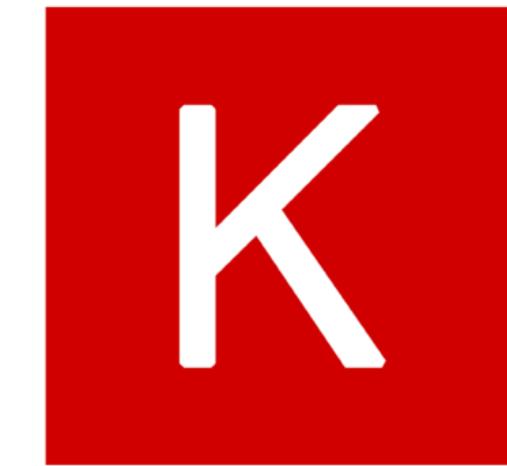
Remember that our goal is NOT to minimize loss on training data!

Learning curves

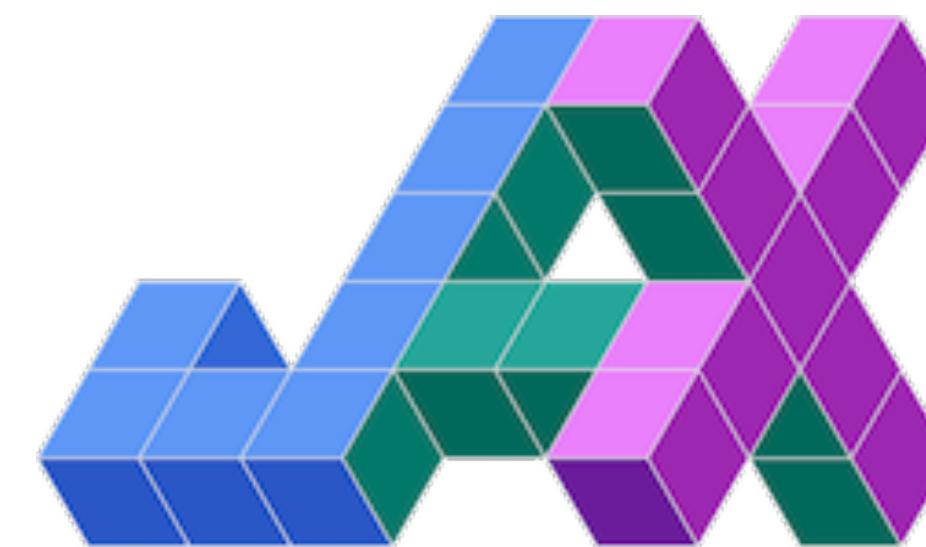


AUTOMATIC DIFFERENTIATION

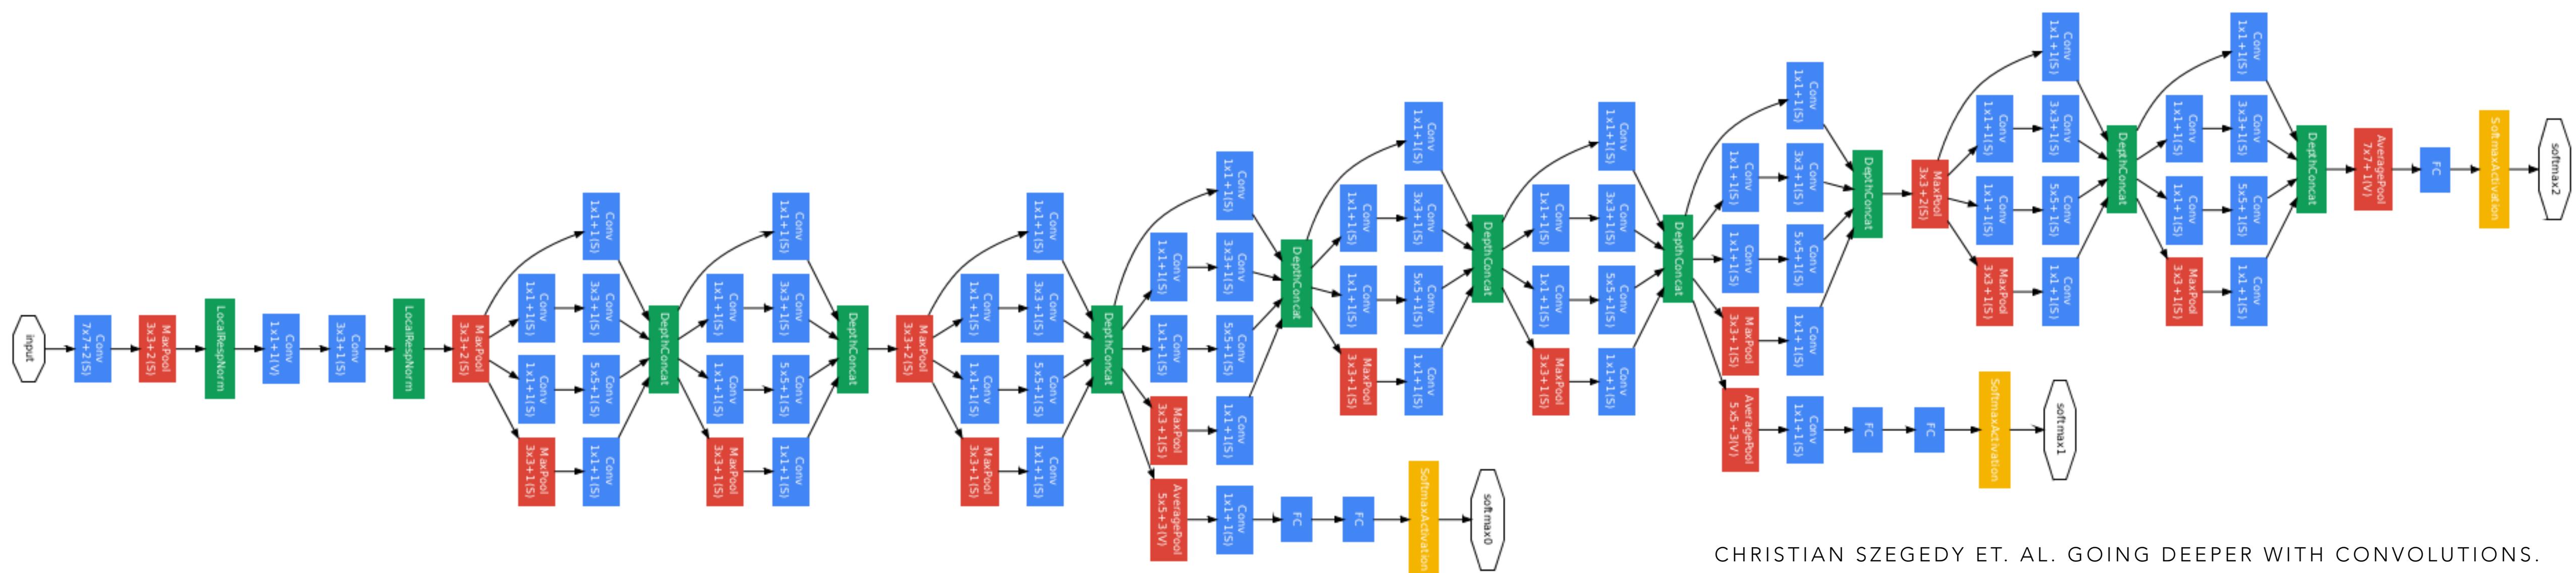
 TensorFlow

 Keras

 PyTorch



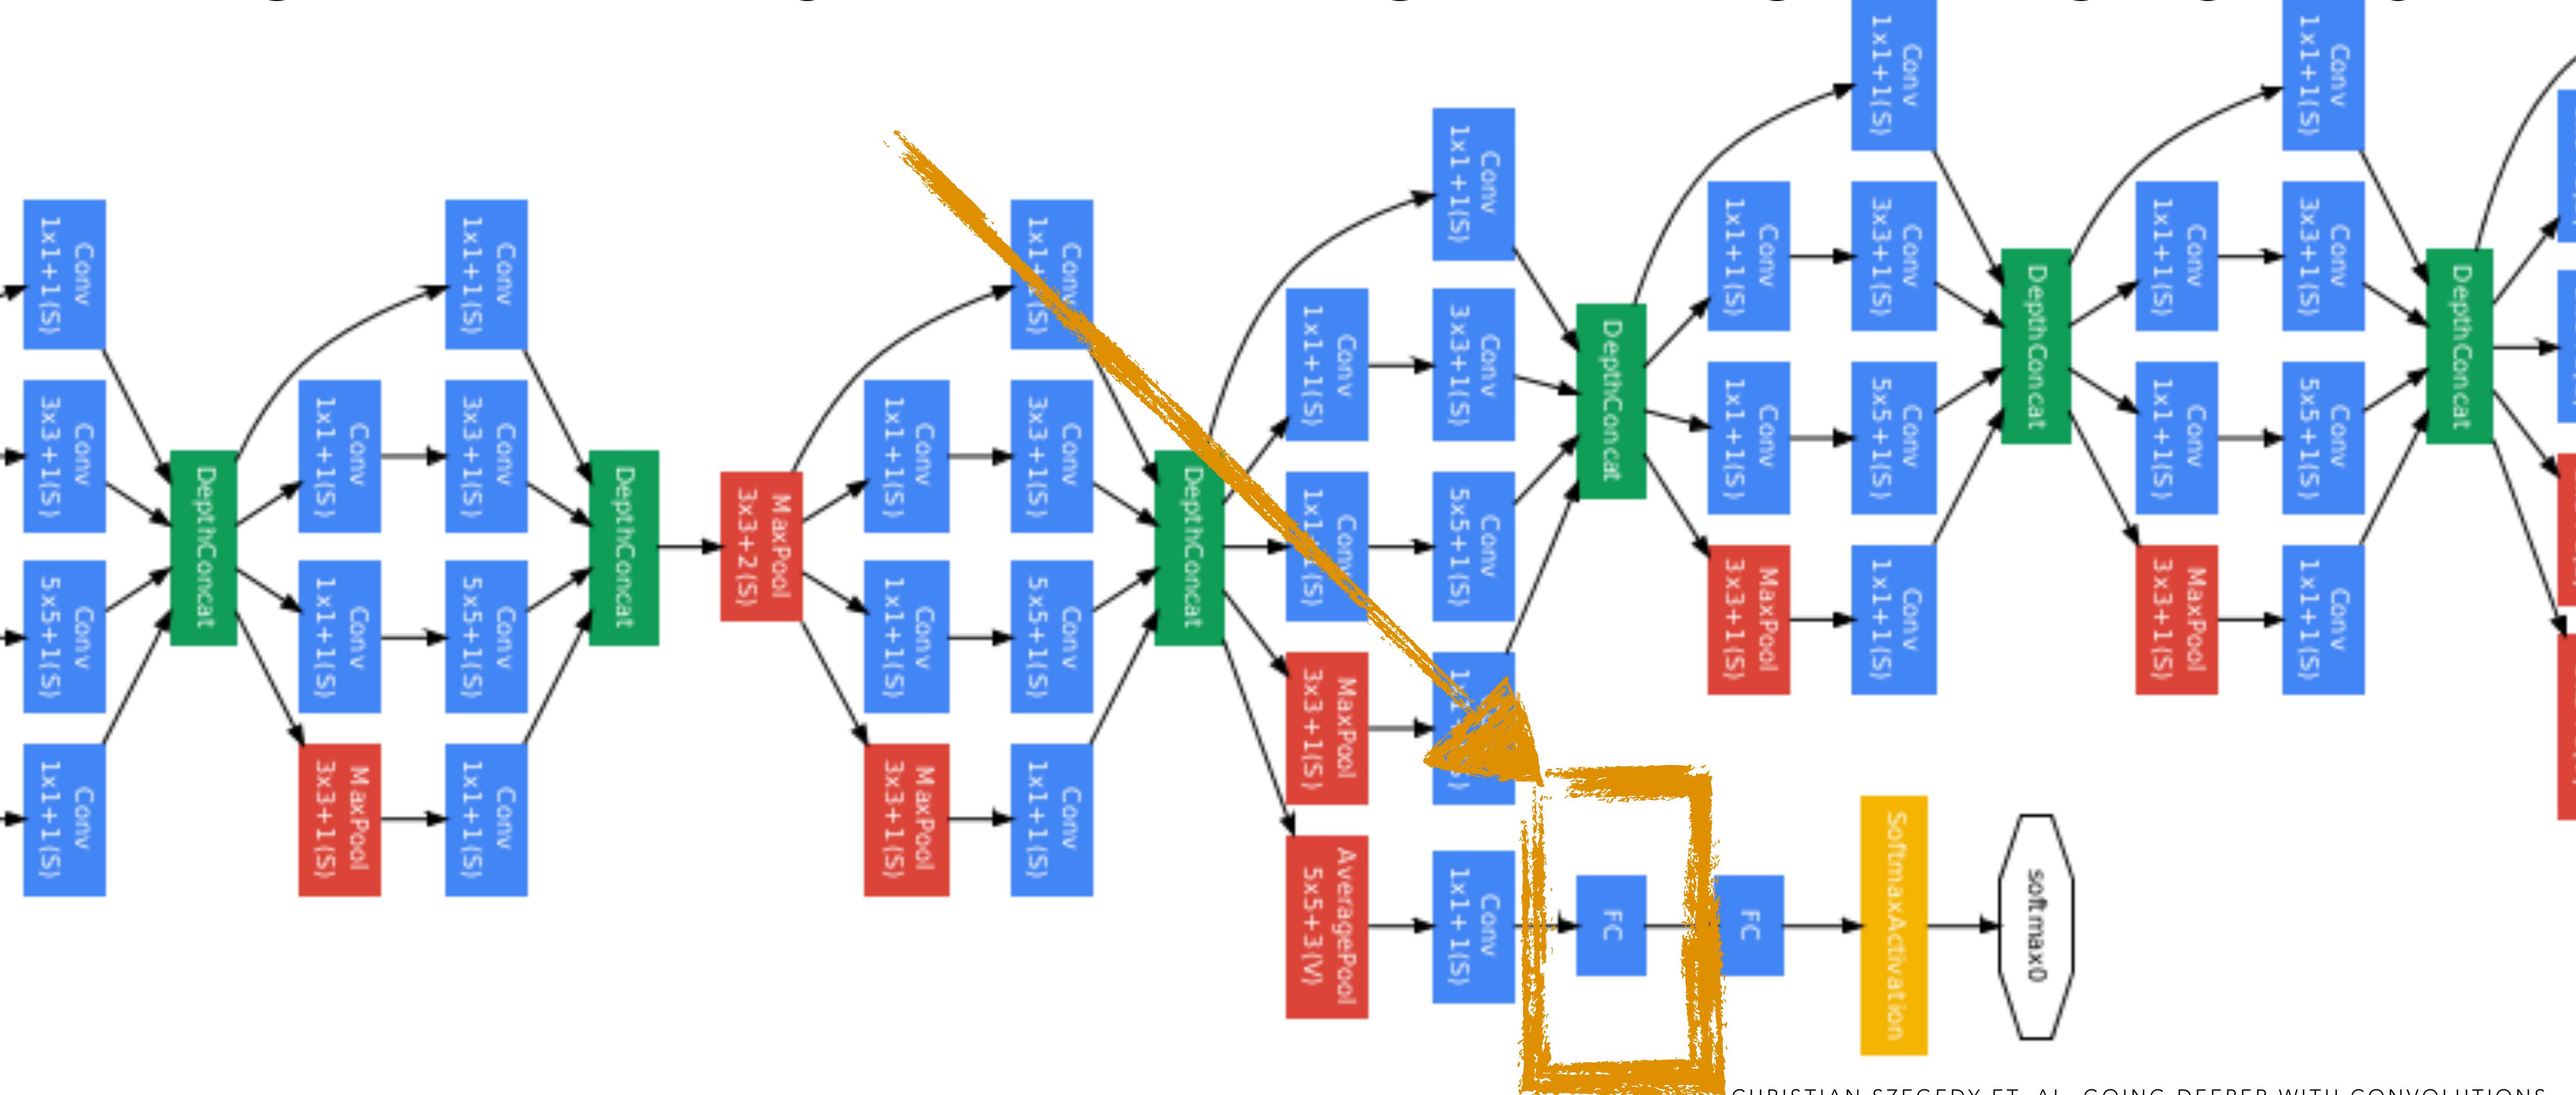
MODERN NEURAL NETWORK ARCHITECTURES



CHRISTIAN SZEGEDY ET. AL. GOING DEEPER WITH CONVOLUTIONS.

"GoogLeNet network with all the bells and whistles"

MODERN NEURAL NETWORK ARCHITECTURES



CHRISTIAN SZEGEDY ET AL. GOING DEEPER WITH CONVOLUTIONS.

“GoogLeNet network with all the bells and whistles”

PRACTICAL TIPS FOR TRAINING MODELS

DATA

	Feature 1	Feature 2	Feature 3	Target
Example 1				
Example 2				
Example 3				
Example 4				

NORMALIZATION

- Puts each feature on same scale
- Allows default hyperparameters to be a good starting point
 - learning rate, initialization of weights, etc.
- Options depend on data distribution
 - Standardization: mean: 0 stdev: 1
 - Min-max: [0,1]

DATA

	Feature 1	Feature 2	Feature 3	Target
Example 1				
Example 2				
Example 3				
Example 4				

ENCODING

- Non-numeric data
- Class-based features:
 - One-hot encoding: $2 \rightarrow [0 \ 1]$
 - When classes do not have sequential meaning:  cars vs dogs vs plants  months

BUILDING AND TRAINING MODELS

TRAINING

- The most challenging part of machine learning is gaining the experience for tuning models well.
- We will work on this skill!

COMMUNITY

- Each of you arrived here with your own backgrounds, specialty, and path in life
- Your experience and expertise are valuable here, no matter what it is
- If the activity is within your background, help others!
- If you are totally (or a little) lost, ask for help!
- It is our shared goal to have **each** of us leave with some new skill/knowledge/understanding